

# Formula 1 telemetry processing using Kafka Streams

Paolo Patierno

Principal Software Engineer

Tom Cooper

Senior Software Engineer



Principal Software Engineer @Red Hat  
Working on Apache Kafka and Strimzi

**@ppatierno**





Senior Software Engineer @Red Hat  
Working on Apache Kafka and Strimzi


**@tomncooper**

# Building an events stream pipeline:

- ▶ How to ingest events reliably
- ▶ How to integrate with different systems for events ingestion (UDP) and providing output
- ▶ How to process events in real time
- ▶ How to show useful insights
- ▶ How to run and deploy the entire pipeline





*influxdb*

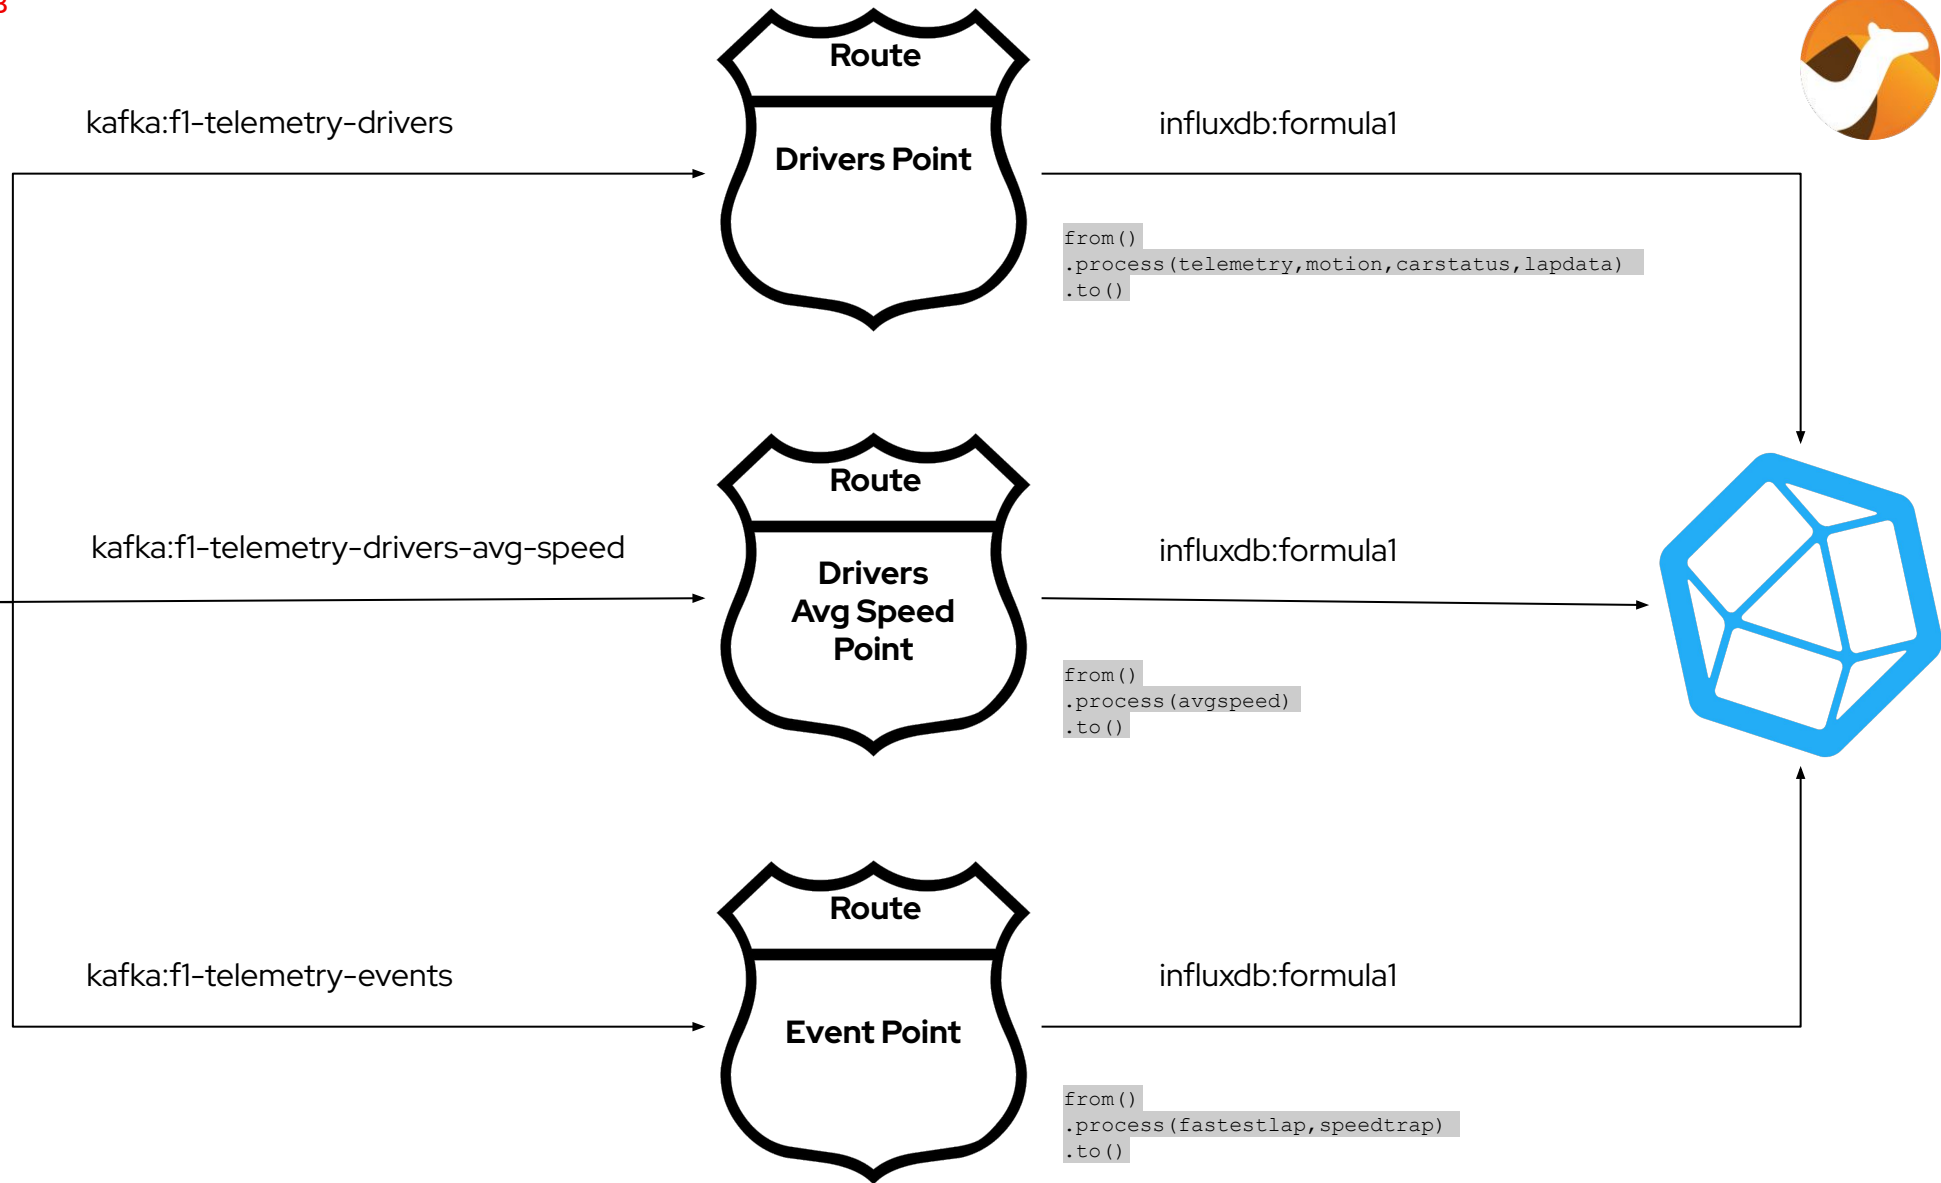




Grafana









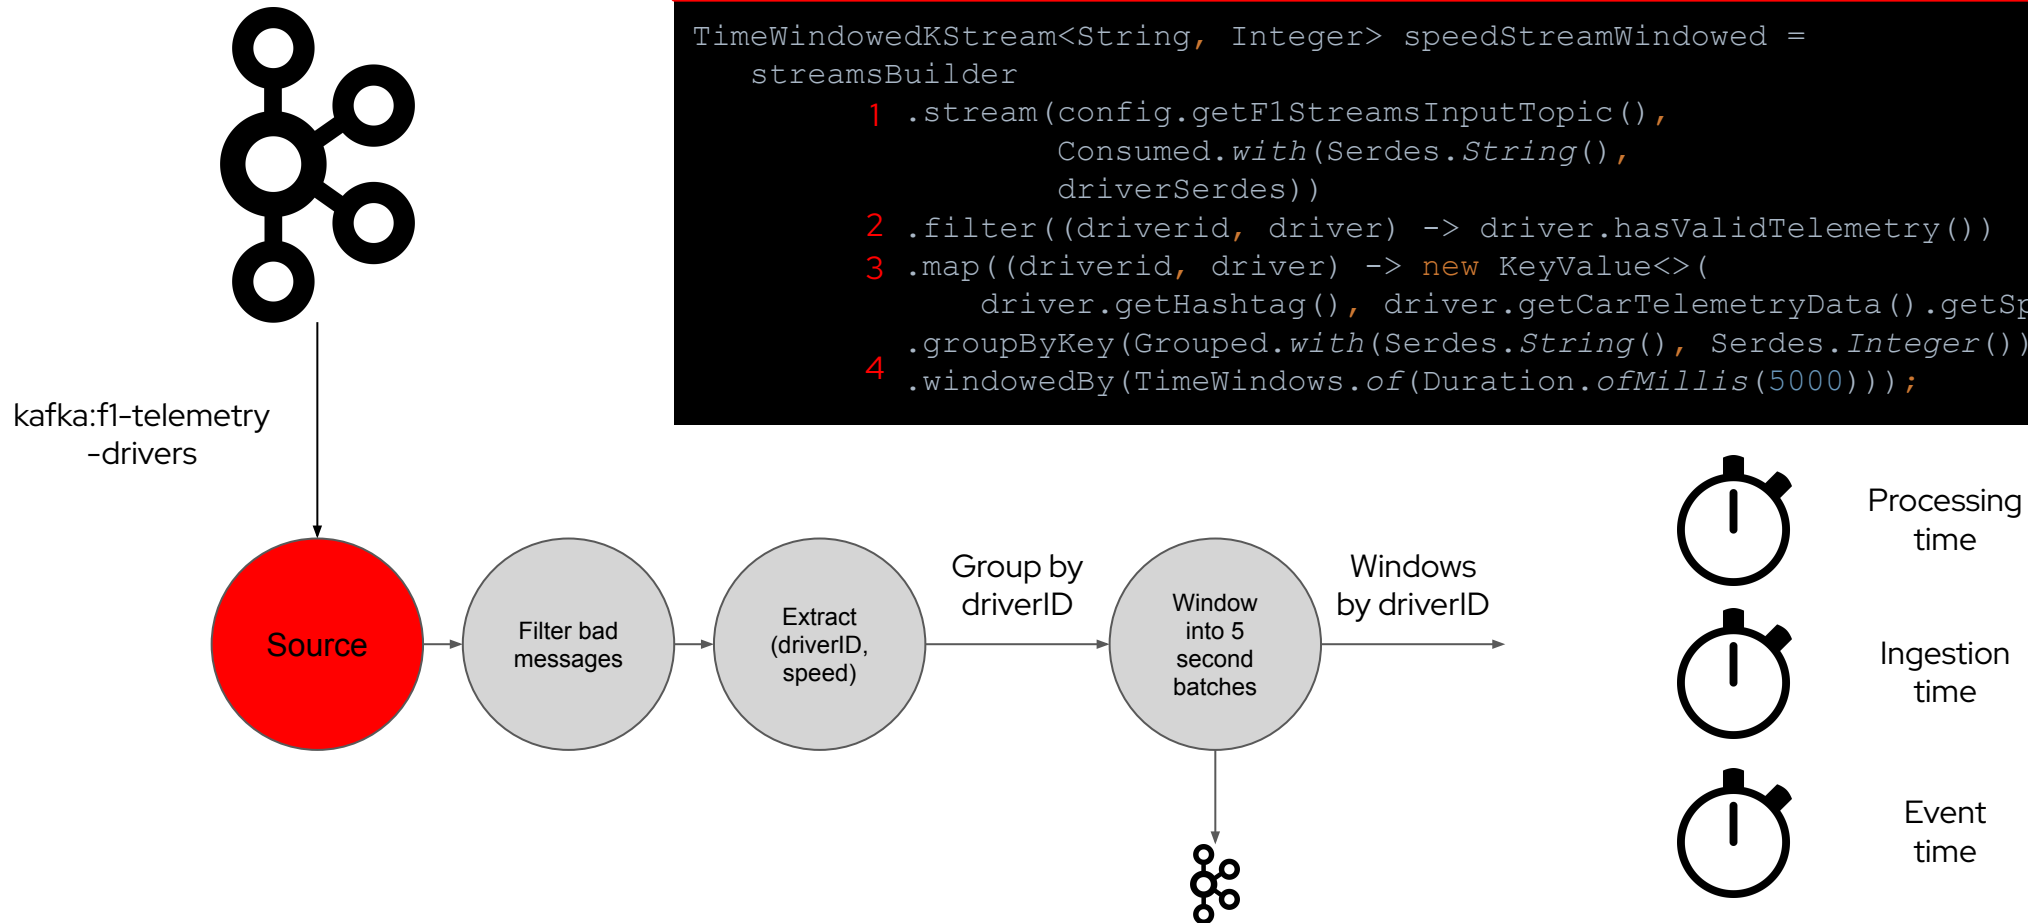




## Setup

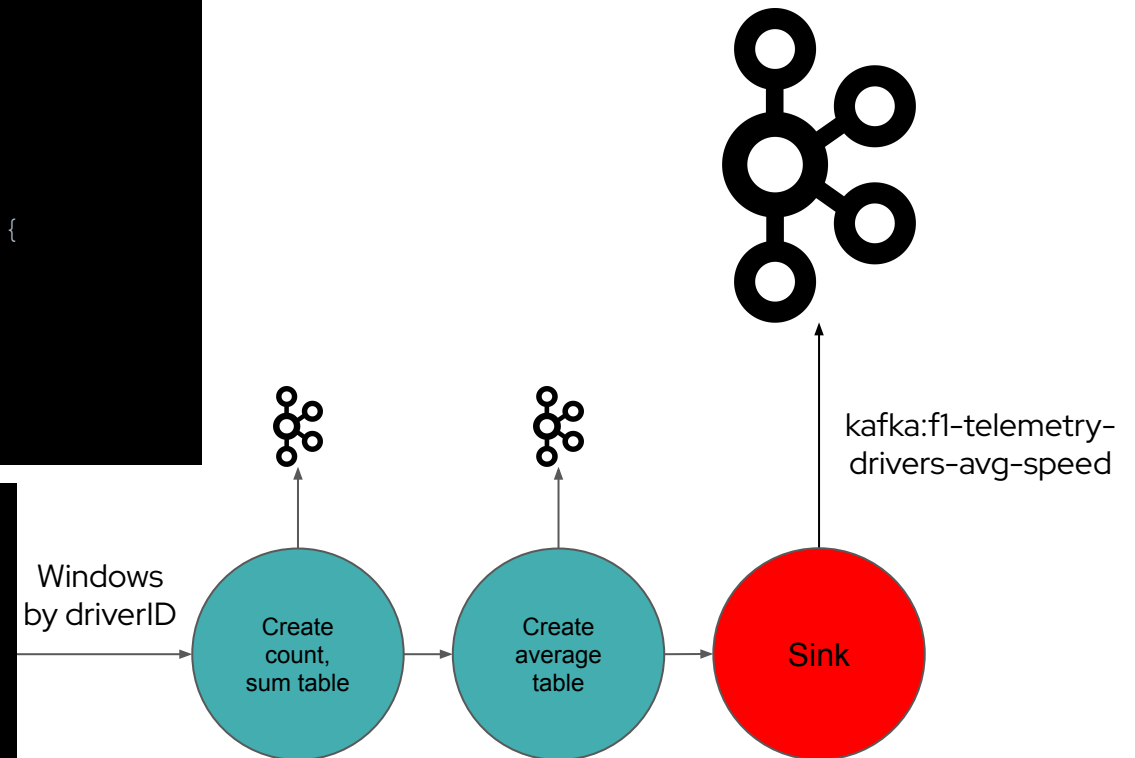
```
StreamsBuilder streamsBuilder = new StreamsBuilder();
Serde<Driver> driverSerdes = Serdes.serdeFrom(
    new DriverSerializer(), new DriverDeserializer());
Serde<SpeedCountAndSum> speedCountAndSumSerde = Serdes.serdeFrom(
    new SpeedCountAndSumSerializer(), new SpeedCountAndSumDeserializer());

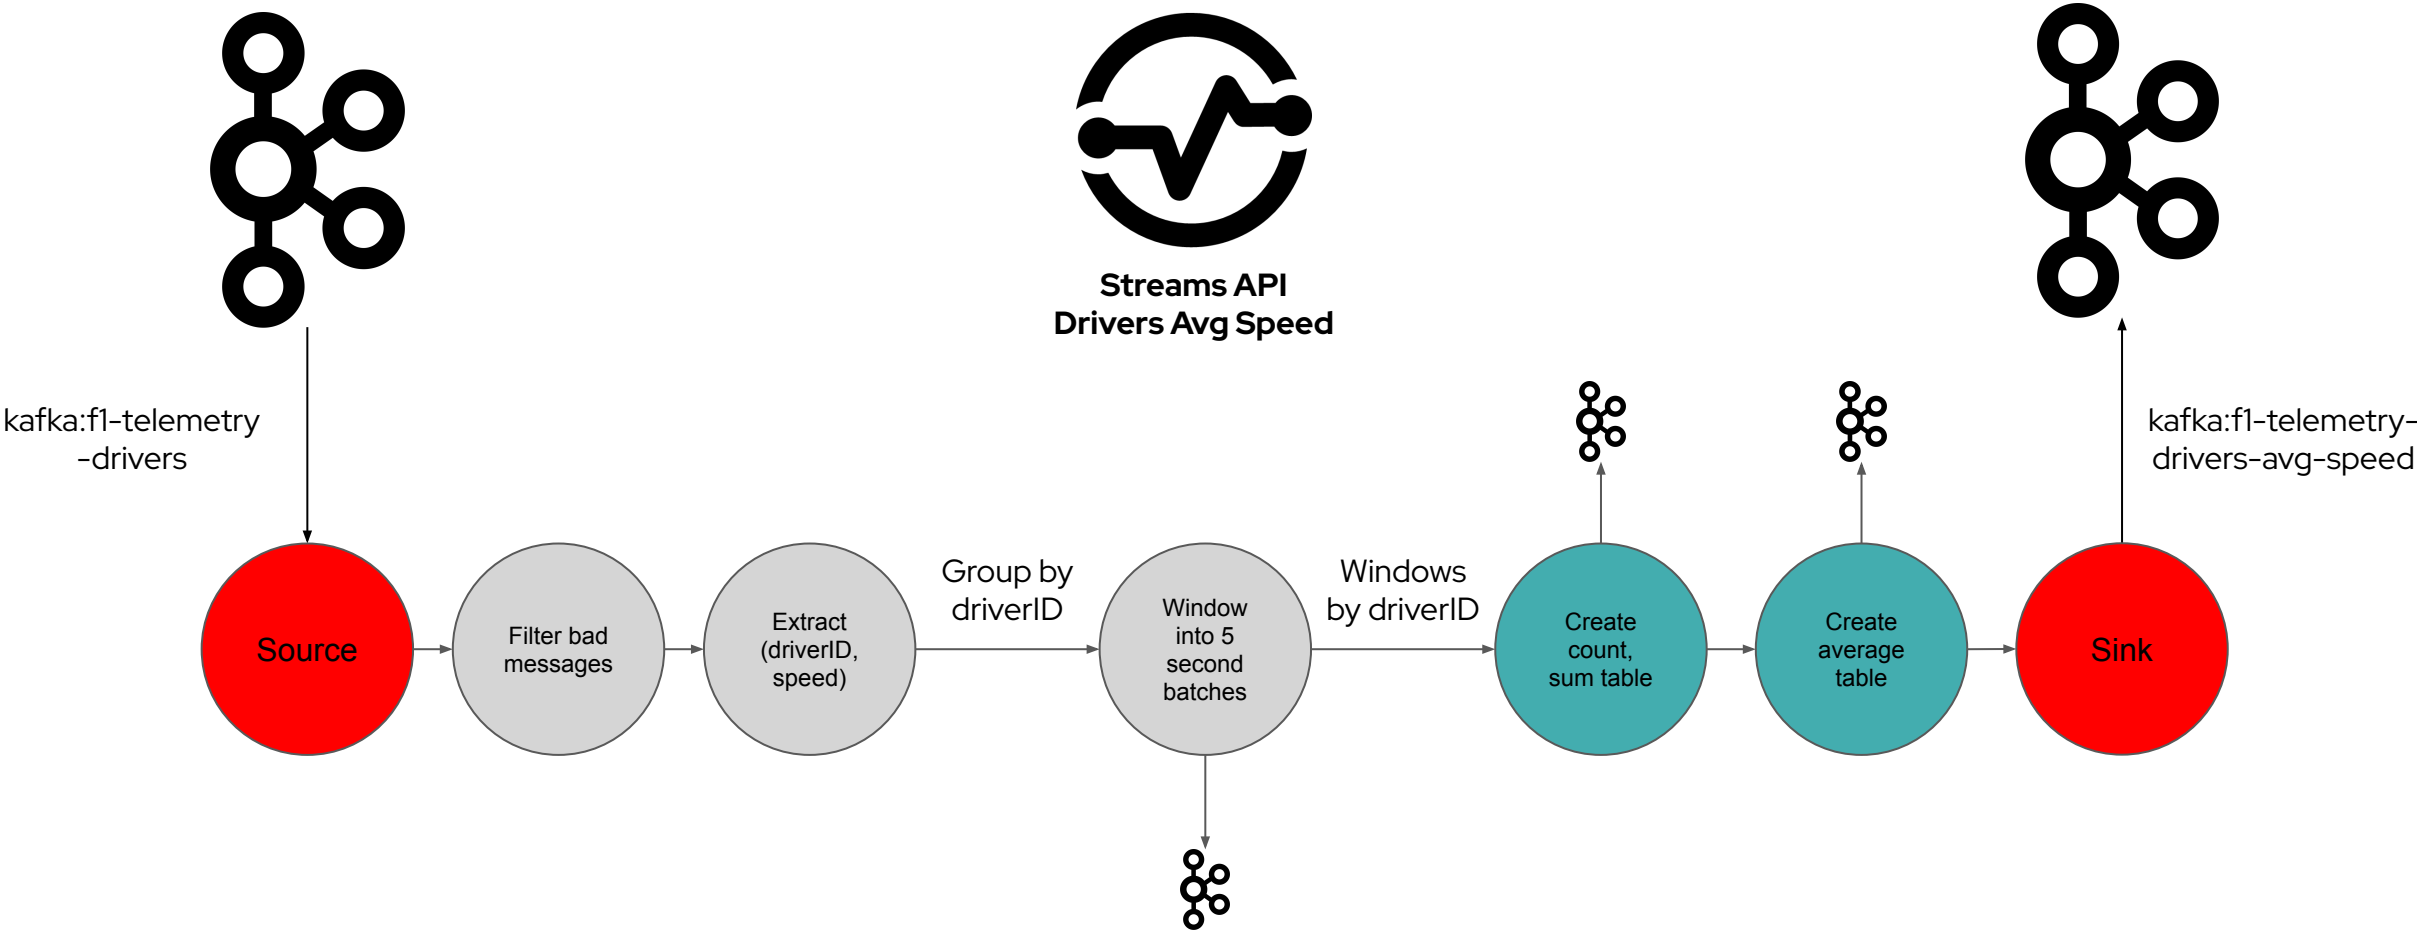
TimeWindowedKStream<String, Integer> speedStreamWindowed =
    streamsBuilder
        1 .stream(config.getF1StreamsInputTopic(),
                Consumed.with(Serdes.String(),
                              driverSerdes))
        2 .filter((driverid, driver) -> driver.isValidTelemetry())
        3 .map((driverid, driver) -> new KeyValue<>(
                driver.getHashtag(), driver.getCarTelemetryData().getSpeed()))
        4 .groupByKey(Grouped.with(Serdes.String(), Serdes.Integer()))
        .windowedBy(TimeWindows.of(Duration.ofMillis(5000)));
```




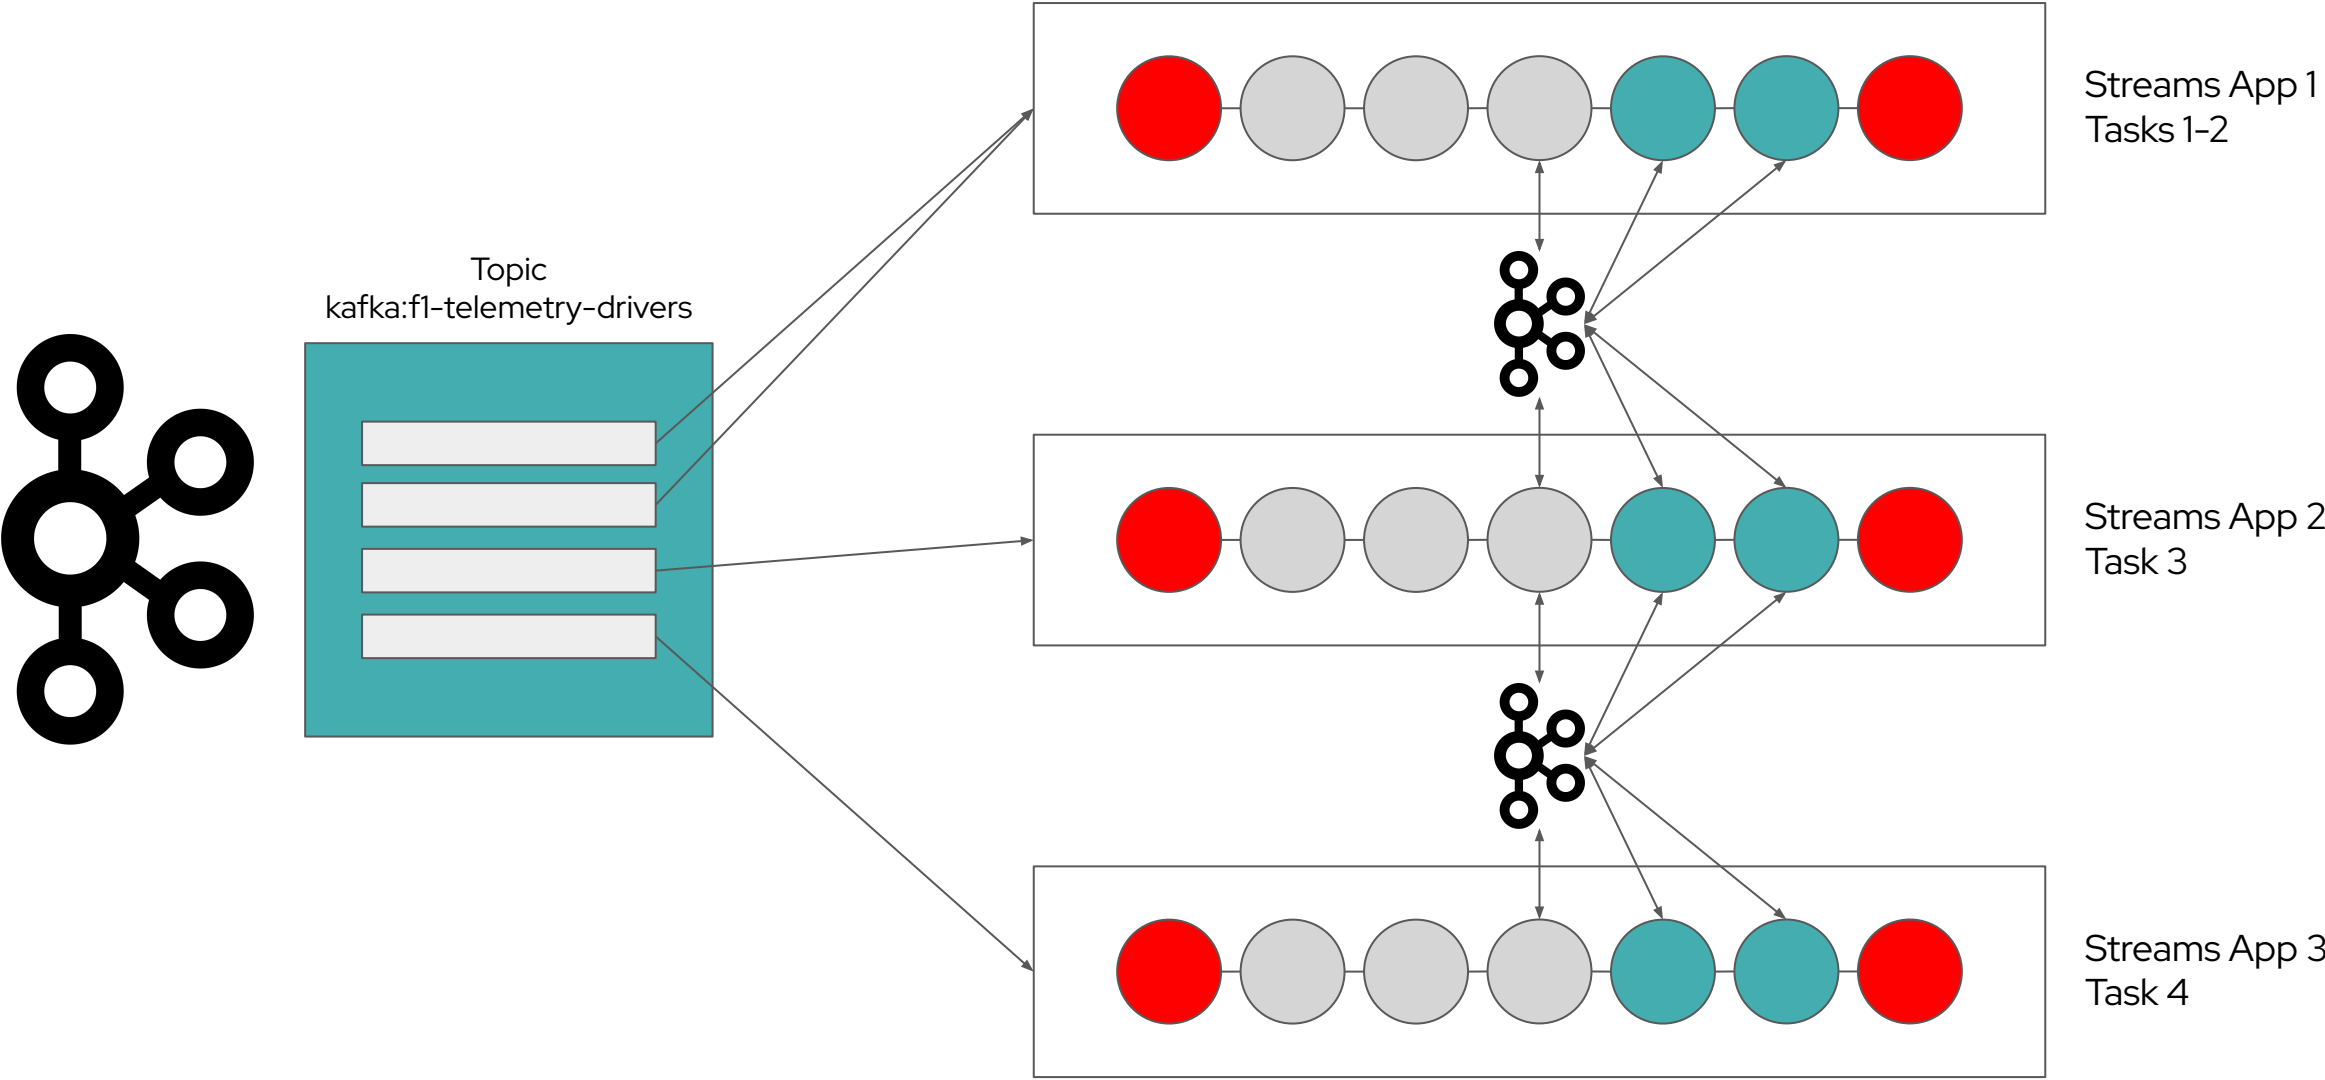
```
KTable<Windowed<String>, SpeedCountAndSum> speedCountAndSum =
    speedStreamWindowed.aggregate(new Initializer<SpeedCountAndSum>() {
        @Override
        public SpeedCountAndSum apply() {
            return new SpeedCountAndSum(0, 0);
        }
    }, new Aggregator<String, Integer, SpeedCountAndSum>() {
        @Override
        public SpeedCountAndSum apply(String key,
            Integer value,
            SpeedCountAndSum aggregate) {
            aggregate.setCount(aggregate.getCount() + 1);
            aggregate.setSum(aggregate.getSum() + value);
            return aggregate;
        }
    }, Materialized.with(Serdes.String(), speedCountAndSumSerde));
```



```
KTable<Windowed<String>, Integer> speedAverage =
    speedCountAndSum.mapValues(
        new ValueMapper<SpeedCountAndSum, Integer>() {
            @Override
            public Integer apply(SpeedCountAndSum speedCountAndSum) {
                return speedCountAndSum.getSum() /
                    speedCountAndSum.getCount();
            }
        }
    );
```












*influxdb*



Grafana



# Demo

# Resources:

- ▶ Blog post: <https://grafana.com/blog/2021/02/02/real-time-monitoring-of-formula-1-telemetry-data-on-kubernetes-with-grafana-apache-kafka-and-strimzi/>
- ▶ F1 decoding library: <https://github.com/ppatierno/formula1-telemetry>
- ▶ F1 Kafka project: <https://github.com/ppatierno/formula1-telemetry-kafka>
- ▶ Video demo: <https://www.youtube.com/watch?v=Re9LOAYZi2A>



- ▶ F1 2020 Codemasters game provides telemetry packets on UDP
  - [Specification is available online](#)
- ▶ [Kubernetes](#) / [OpenShift](#)
  - Deploying the [Apache Kafka](#) cluster through [Strimzi](#) project
  - Running Apache Camel applications, InfluxDB and Grafana
- ▶ [Apache Camel](#)
  - Ingesting telemetry packets to Apache Kafka
  - Gets telemetry data and race events from Apache Kafka; store into InfluxDB
- ▶ [InfluxDB](#)
  - Time-series database to provide data to Grafana
- ▶ [Grafana](#)
  - Showing all telemetry and events on specific dashboards