

Finding a work-load balance

Cruise Control for Kafka on Kubernetes

Kyle Liberti

Software Engineer

Paolo Patierno

Principal Software Engineer



Principal Software Engineer @Red Hat
Working on Apache Kafka and Strimzi

@ppatierno



Software Engineer @Red Hat
Working on Apache Kafka and Strimzi

Kyle Liberti

Overview

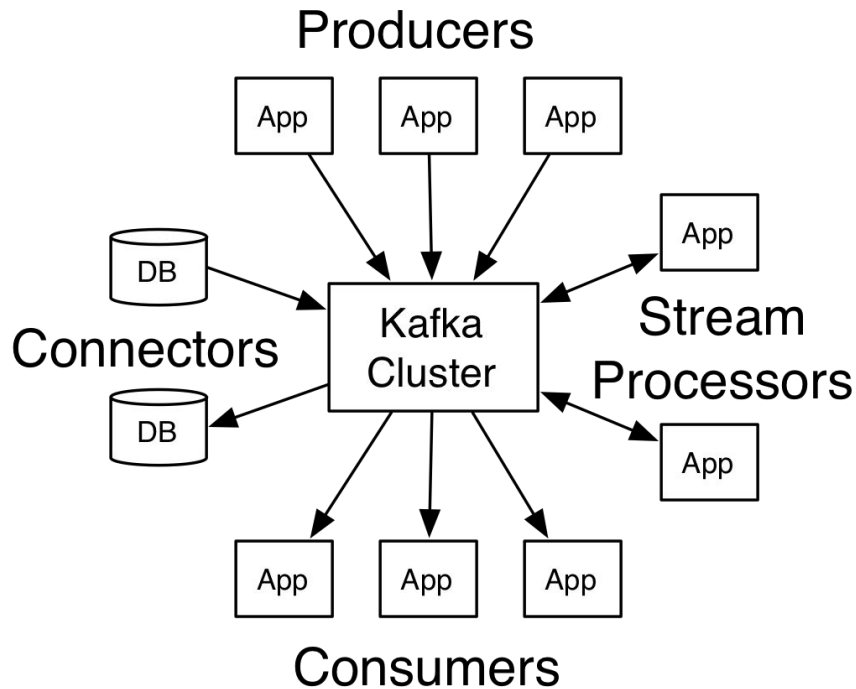
- 1) Intro to Kafka
- 2) Managing Kafka clusters using Strimzi
- 3) Intro to Kafka Cruise Control
- 4) Balancing Kafka clusters using Strimzi

What is Kafka?

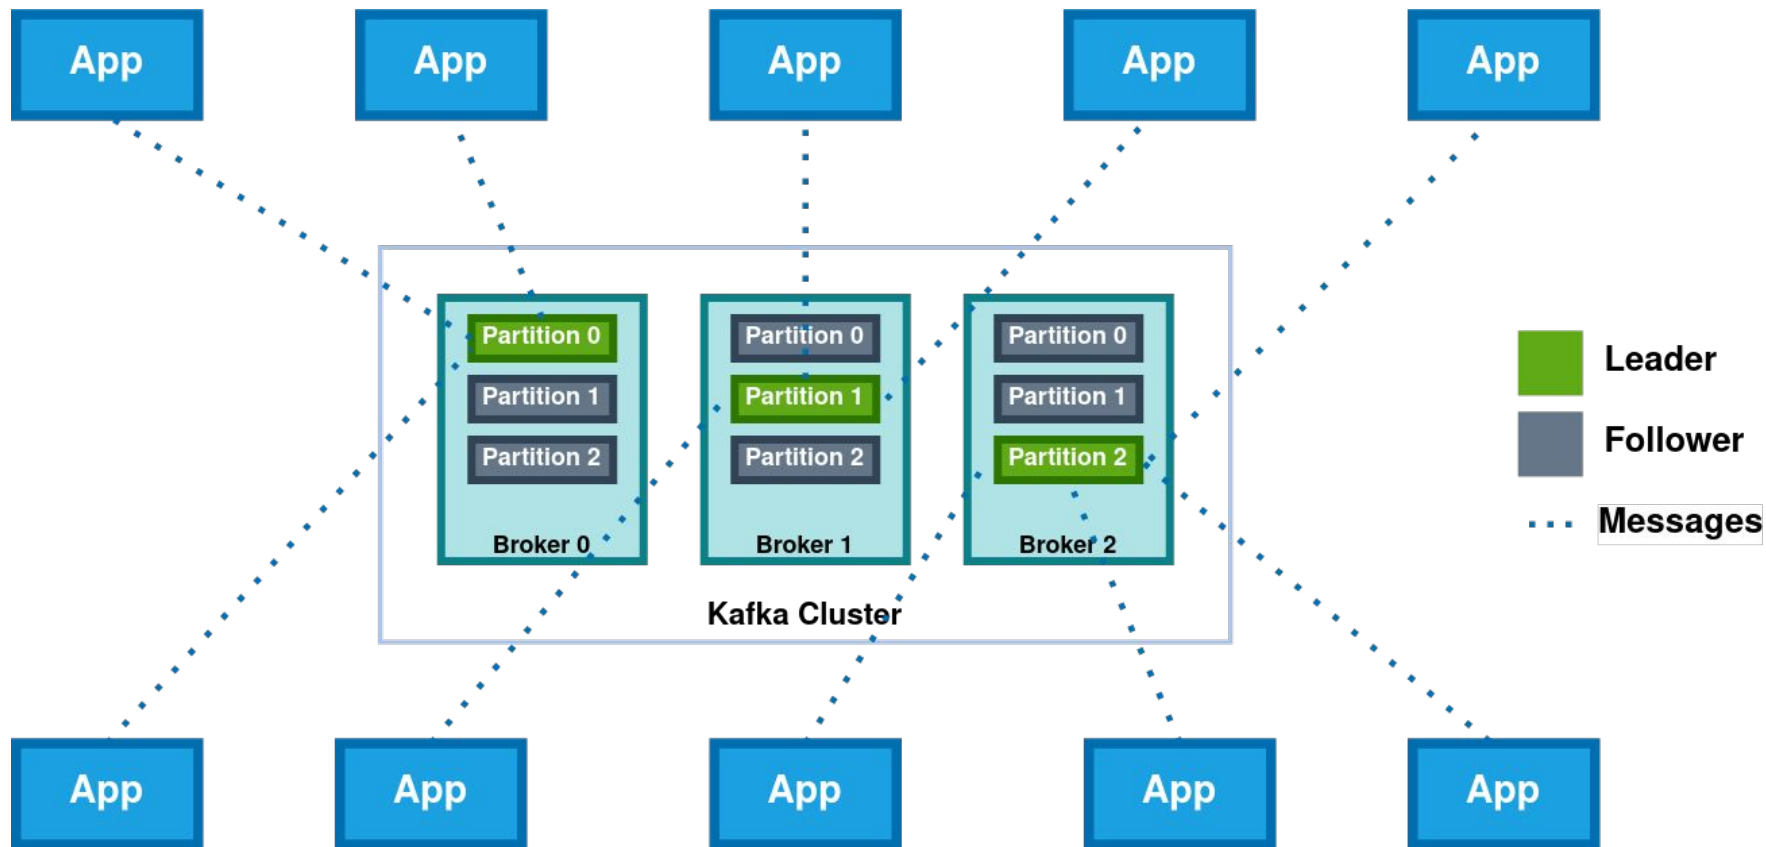


Kafka

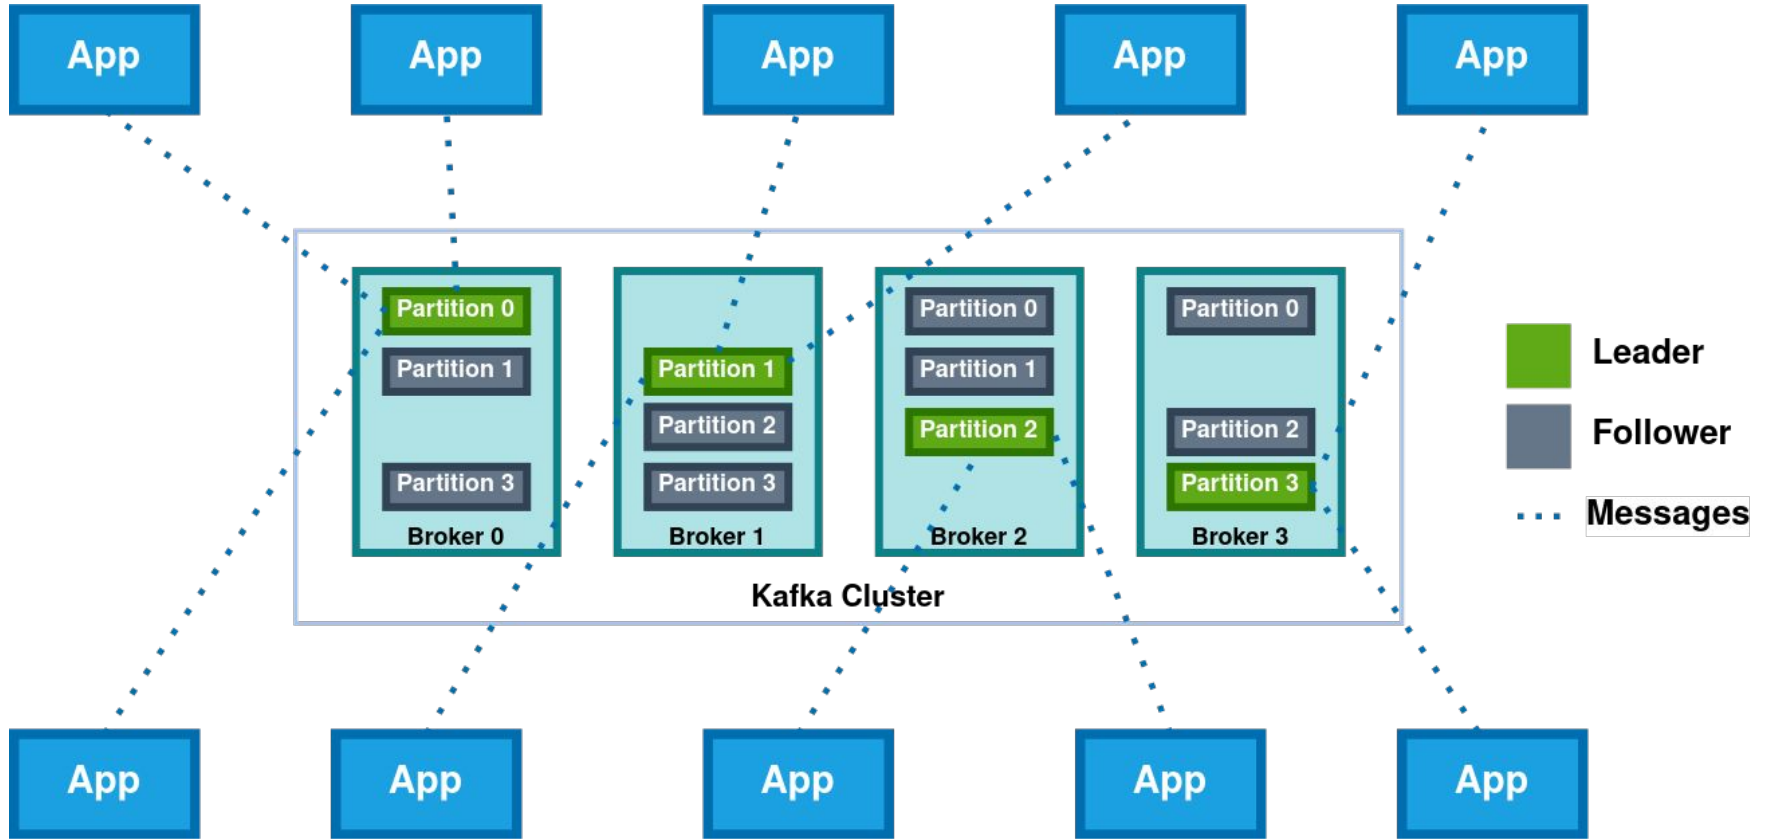
- ▶ Distributed messaging system
- ▶ Scalable and fault tolerant
- ▶ Integrated ecosystem



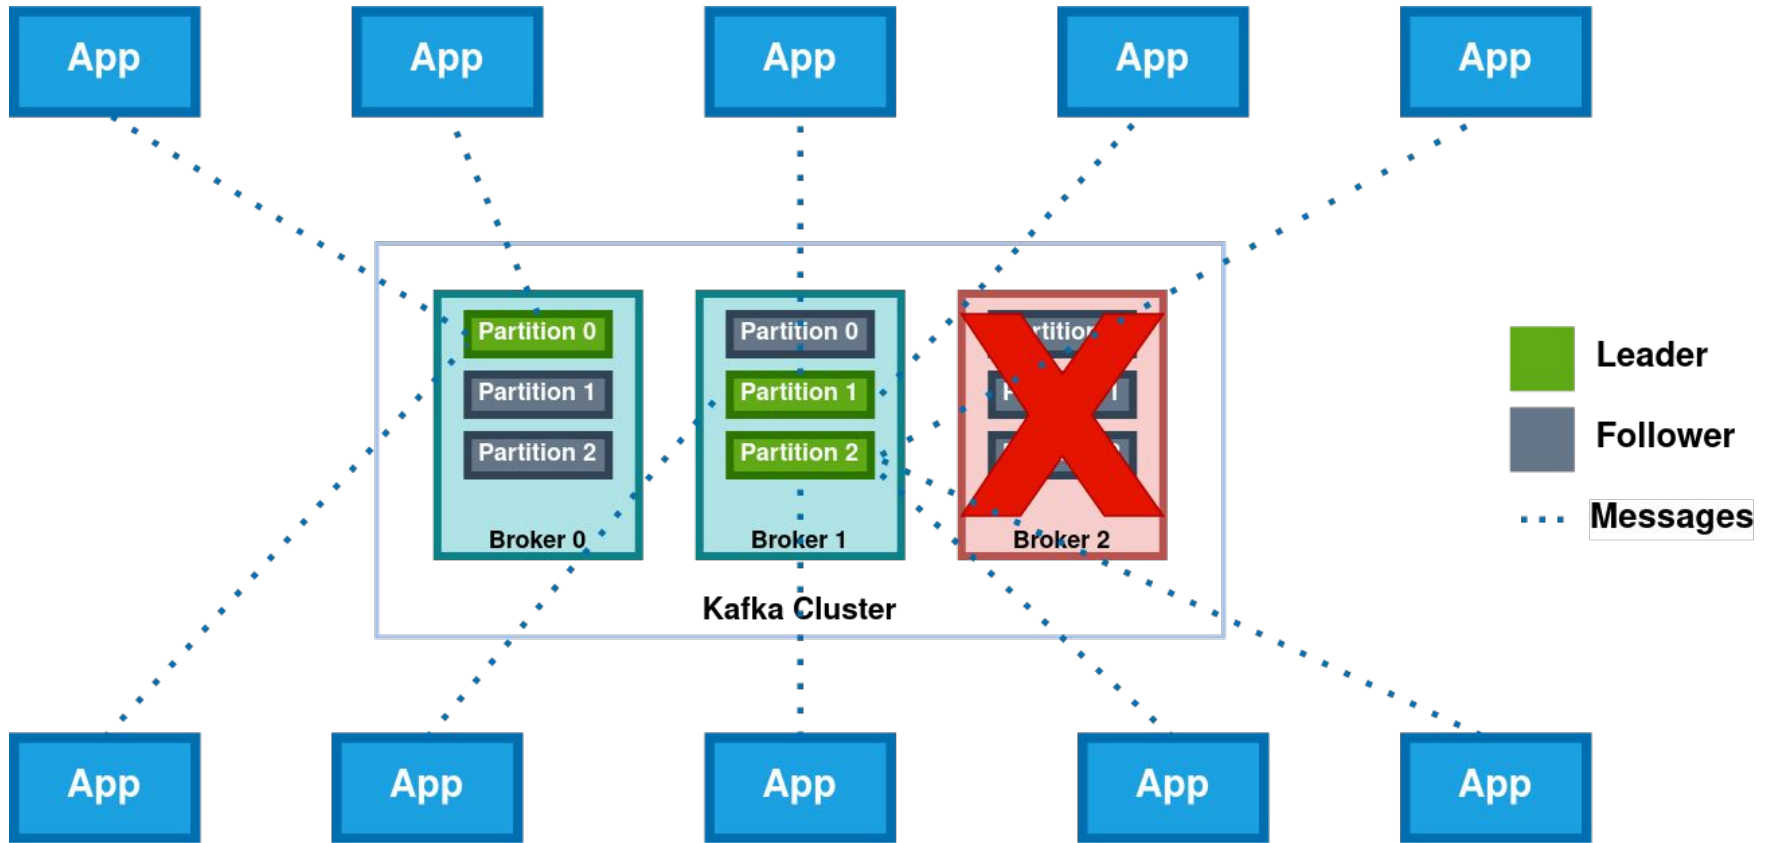
Simplified Kafka Architecture



Kafka scales horizontally



Kafka is fault tolerant



BUT...

Kafka is **HARD**

- ▶ Configuration
- ▶ Operation
- ▶ Development




IS THERE AN EASIER WAY?


Strimzi

Operator for running Kafka on Kubernetes

- ▶ Automated configuration and deployment
- ▶ Built-in security
- ▶ Simple user interface

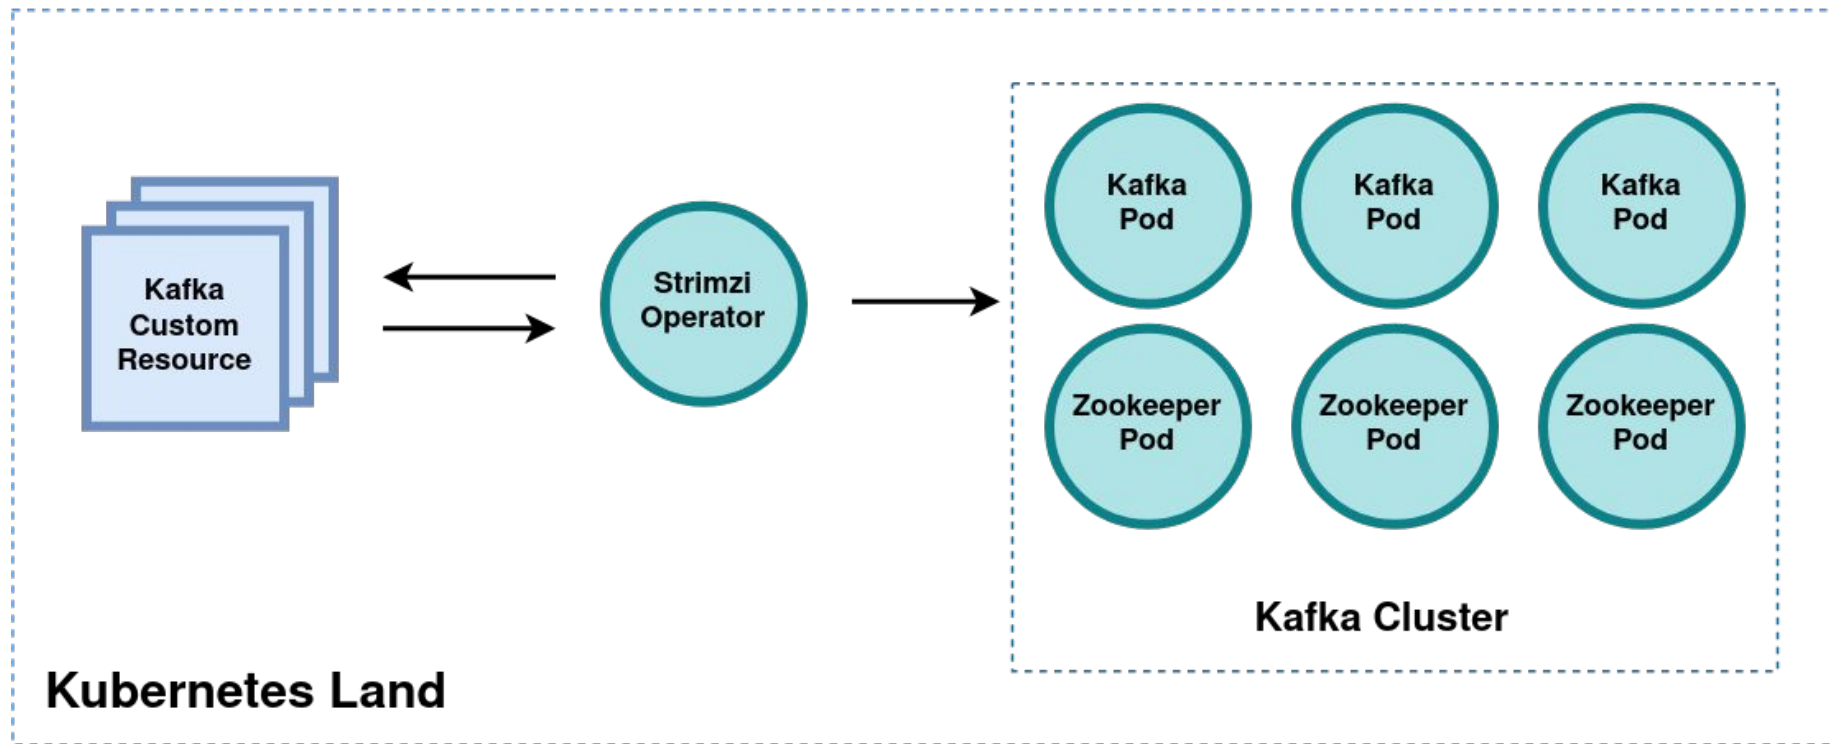


Strimzi Custom Resources



```
apiVersion: kafka.strimzi.io/v1beta2
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    version: 2.8.0
    replicas: 3
    config:
      offsets.topic.replication.factor: 3
      log.message.format.version: "2.8"
      inter.broker.protocol.version: "2.8"
    storage:
      type: ephemeral
  zookeeper:
    replicas: 3
    storage:
      type: ephemeral
```

Simplified Strimzi Architecture



NOT OUT OF THE WOODS YET...

Unbalanced Kafka brokers

- Uneven distribution of load across brokers

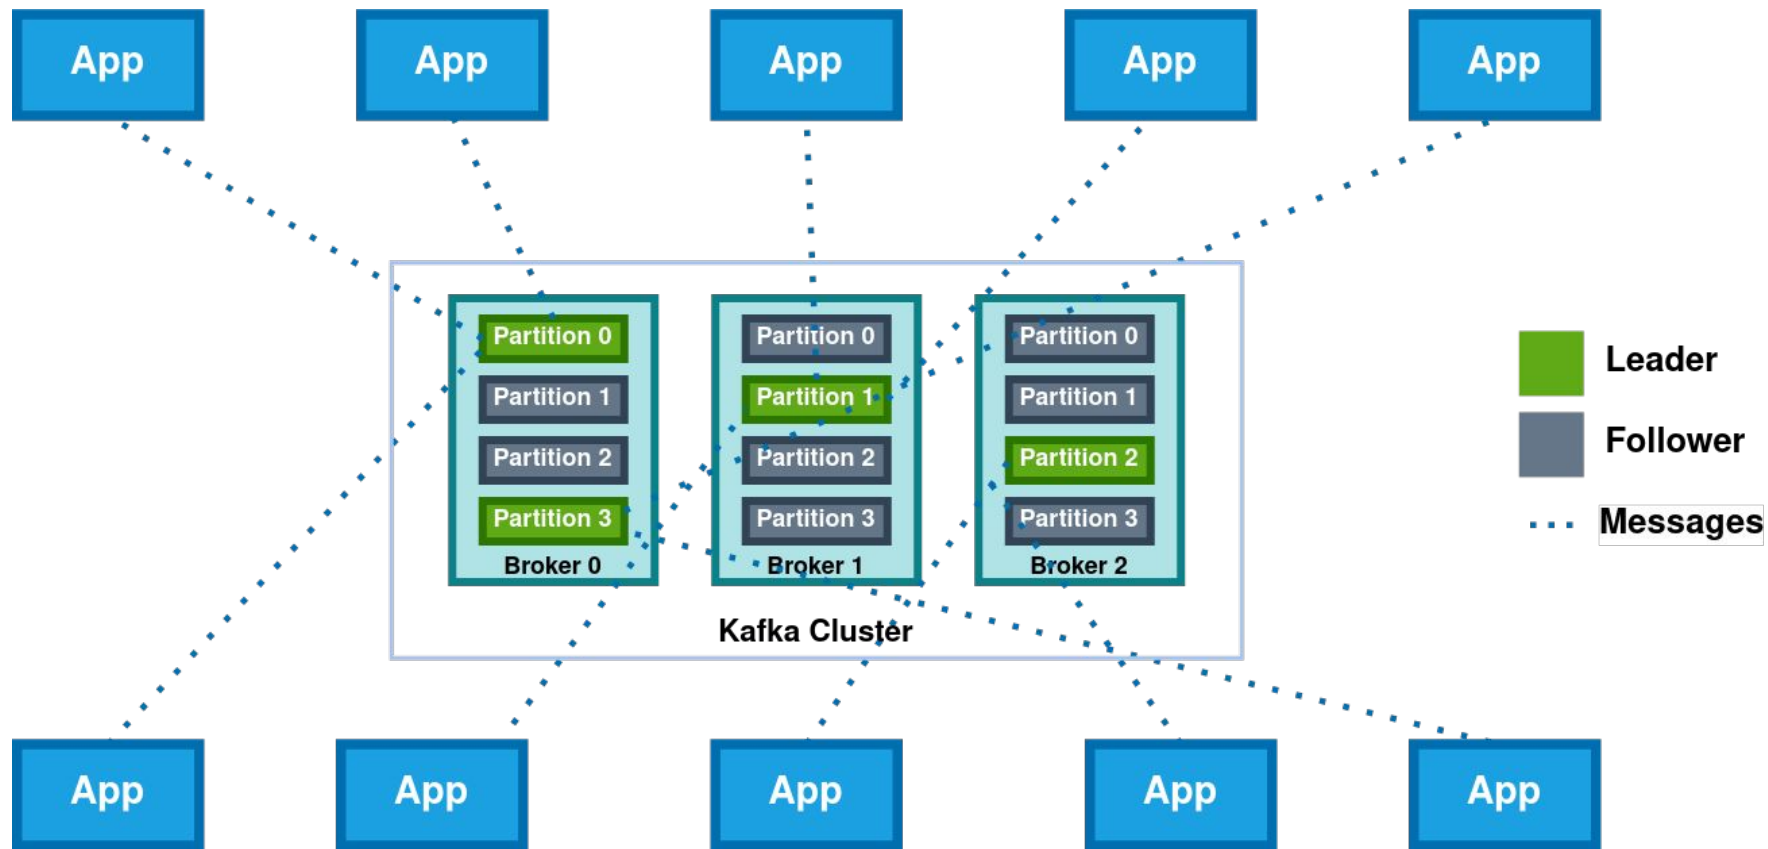
- > **Poor performance**

- > **Storage and Network are badly utilized**

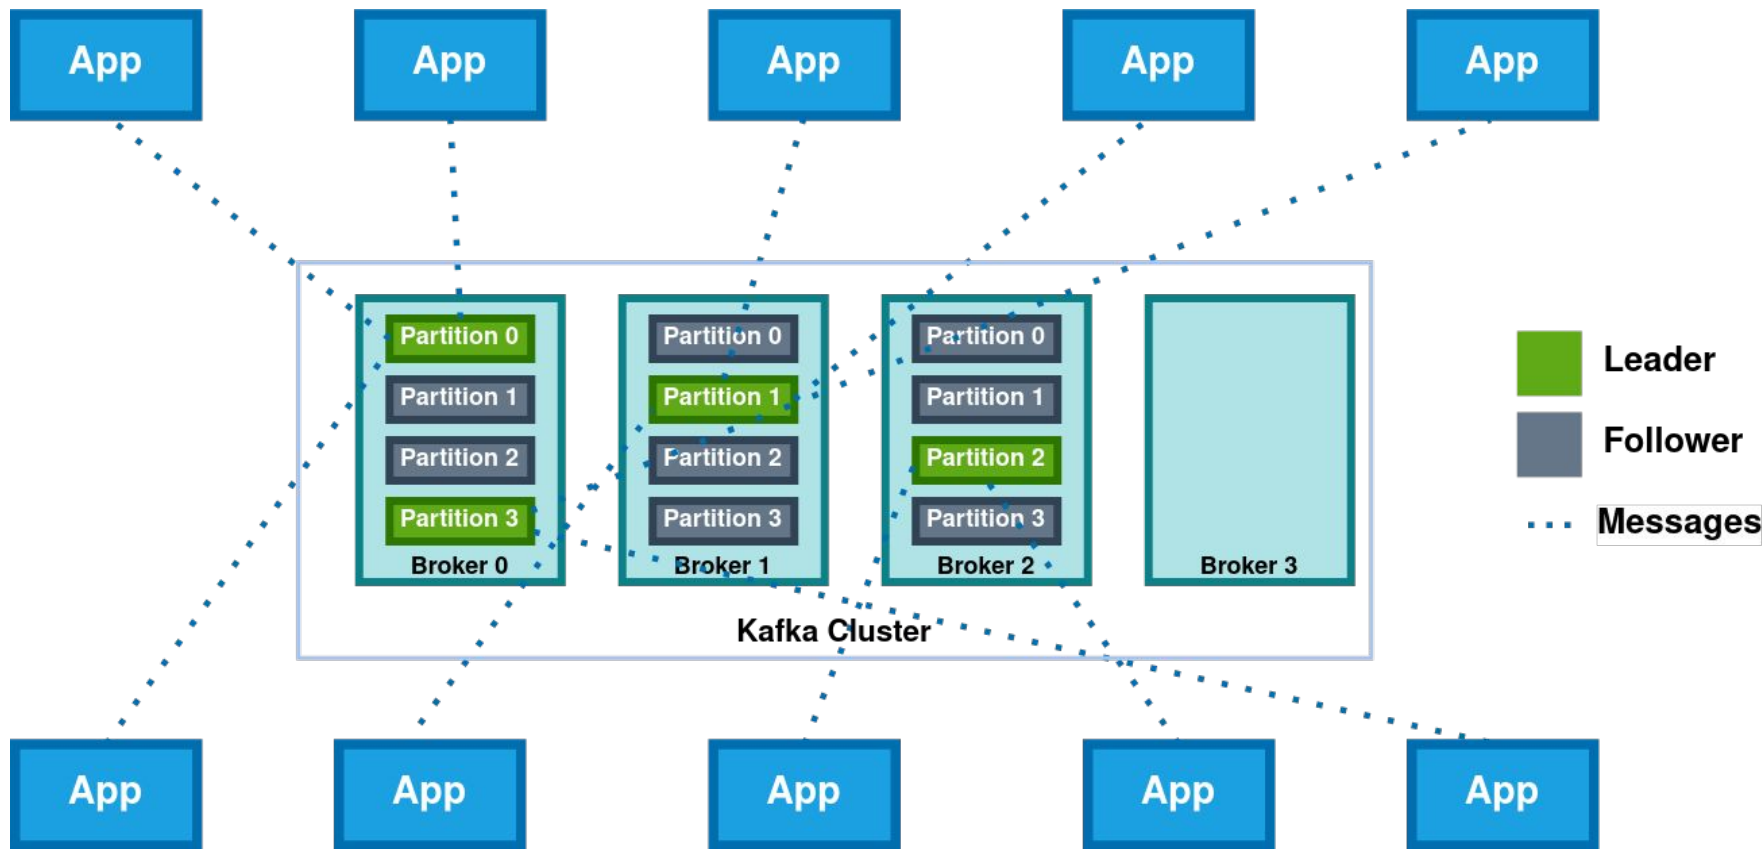
- > **More load on some brokers**



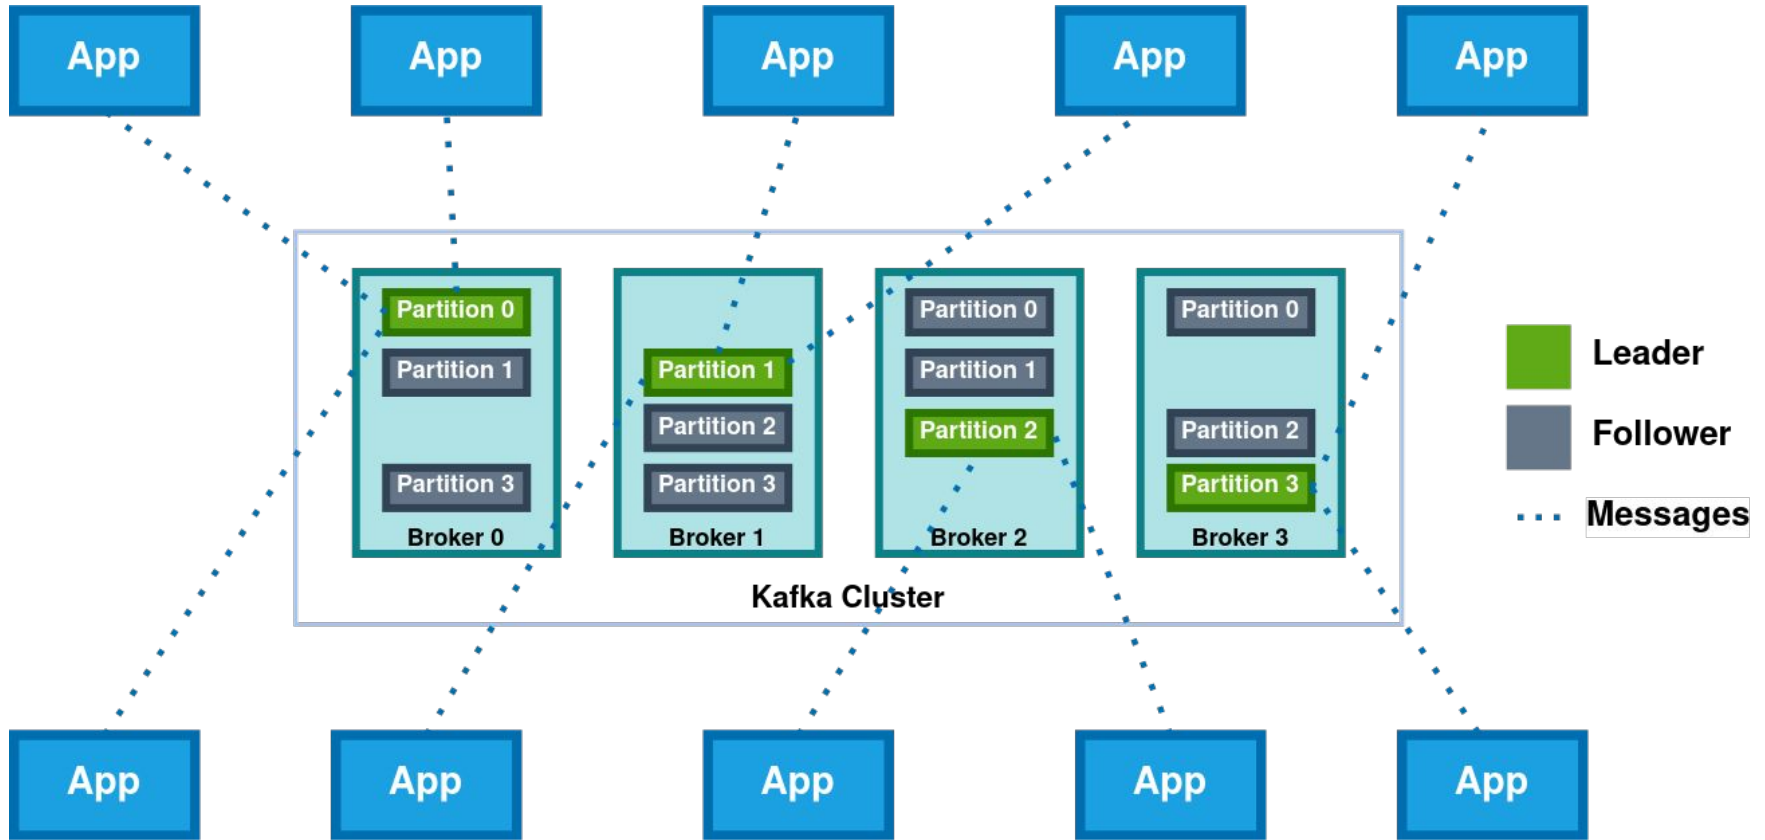
Simplified Kafka Architecture



Unbalanced Kafka Cluster



Balanced Partition Distribution



WHAT CAN WE DO ABOUT IT?


Balancing Partitions is **HARD**

- ▶ Moving target
- ▶ Optimization



Old School Rebalancing

- ▶ kafka-reassign-partitions.sh
- ▶ Manual process
- ▶ One dimensional balancing



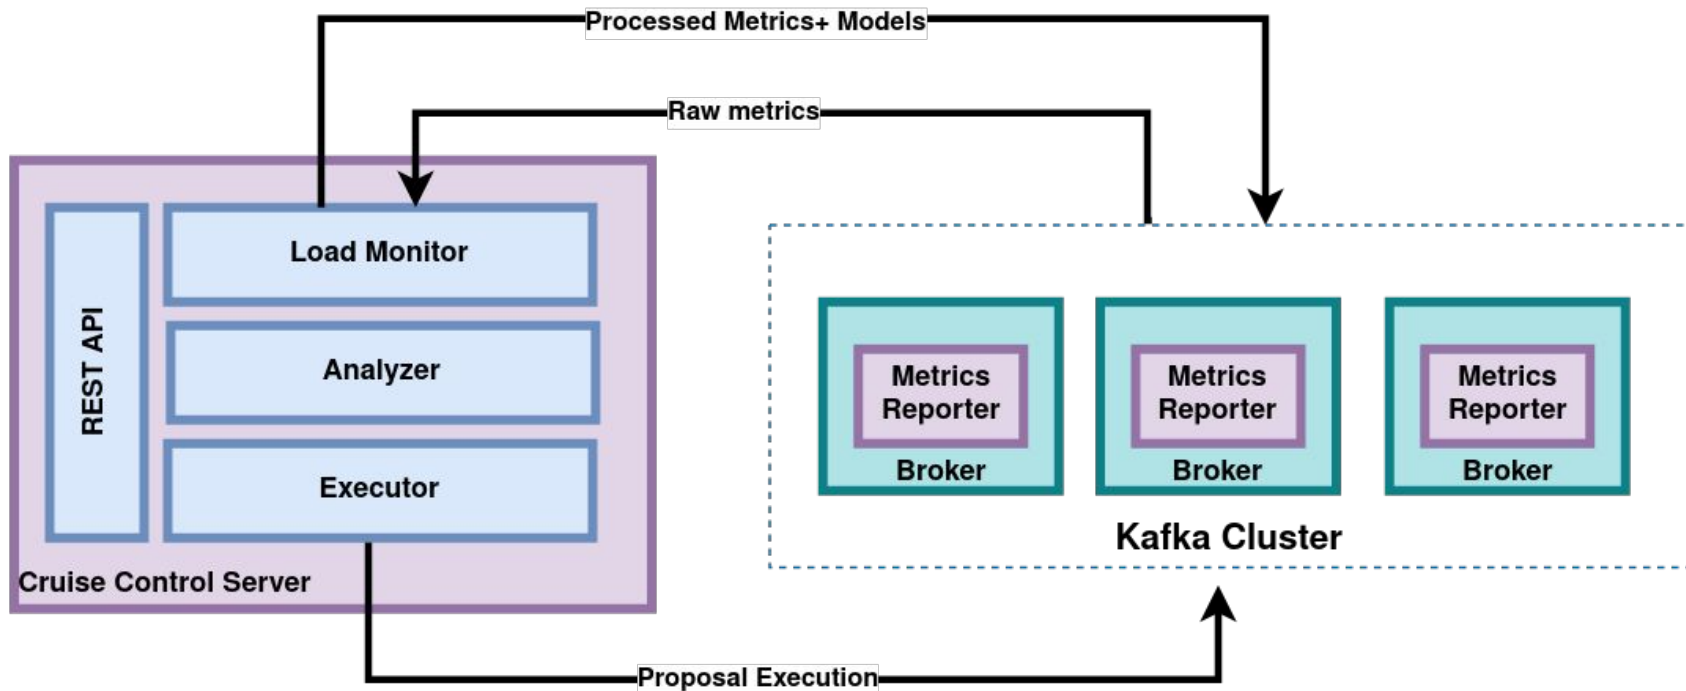
CAN WE DO BETTER?

Cruise Control

- ▶ Fully automated rebalancing
- ▶ Fine-grained resource tracking
- ▶ Multidimensional balancing

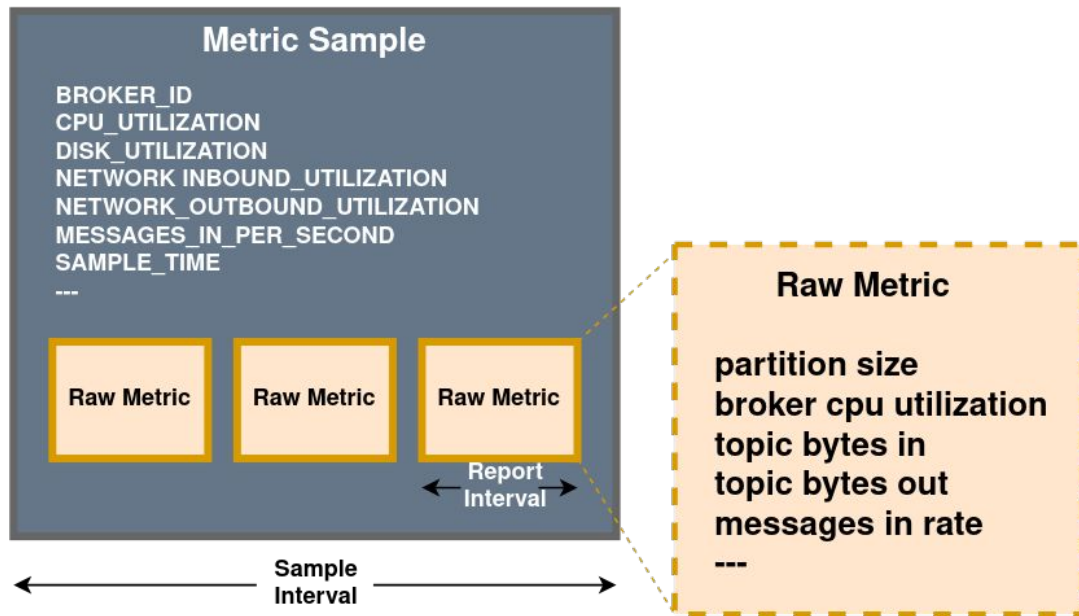


Simplified Cruise Control Architecture

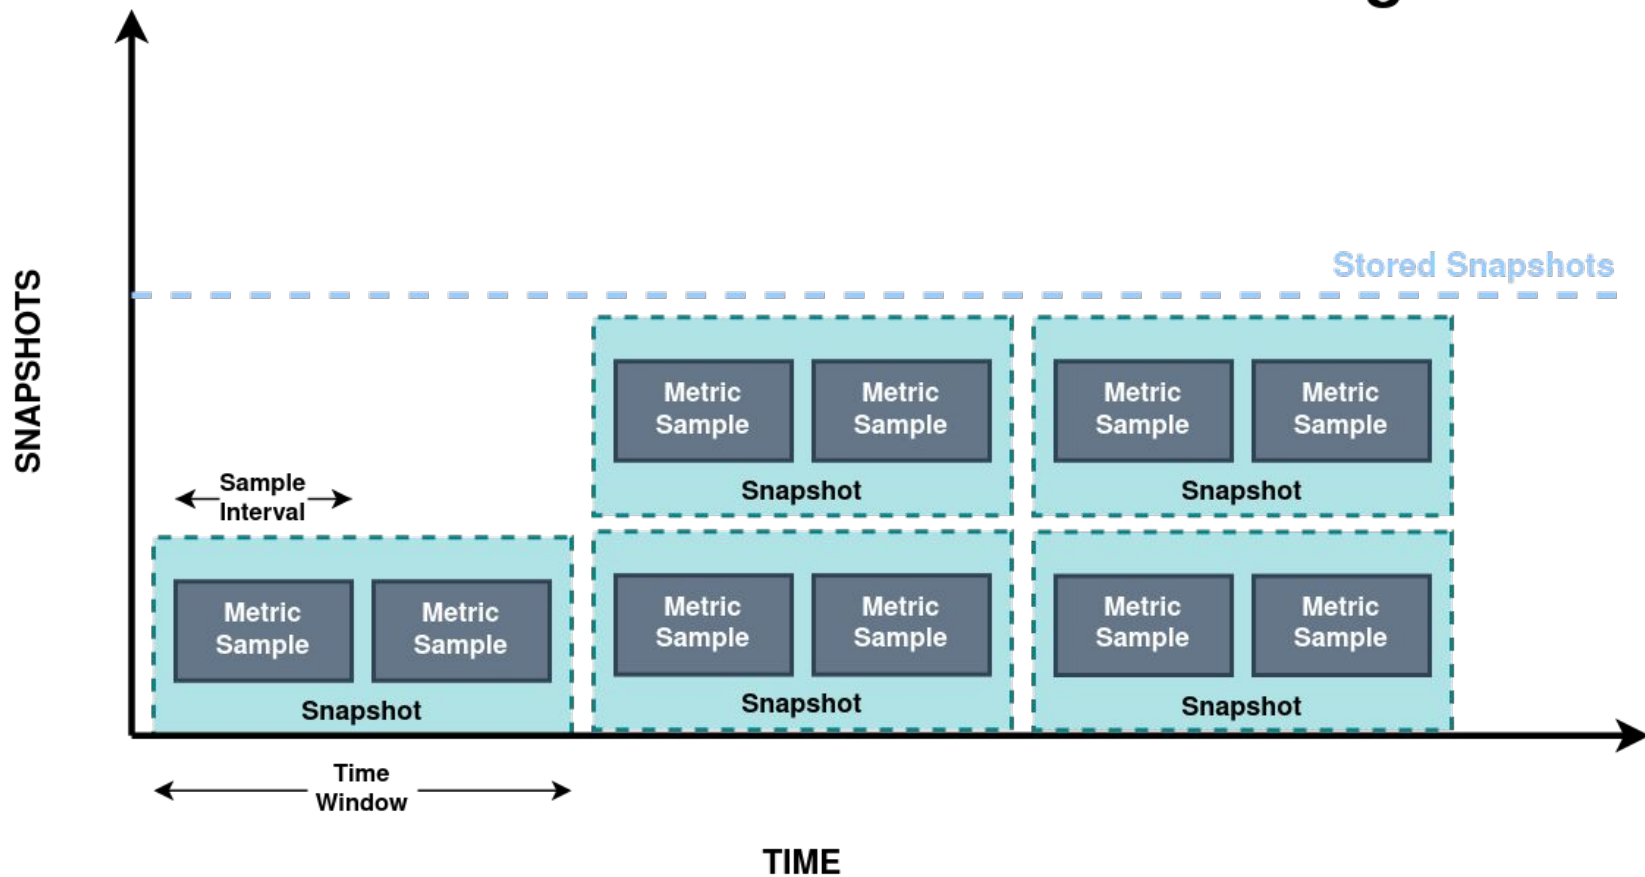


Metric Samples

- ▶ Processed Kafka broker metrics
- ▶ Configurable sampling intervals
- ▶ Stored back into Kafka



Load Monitor Metrics Processing

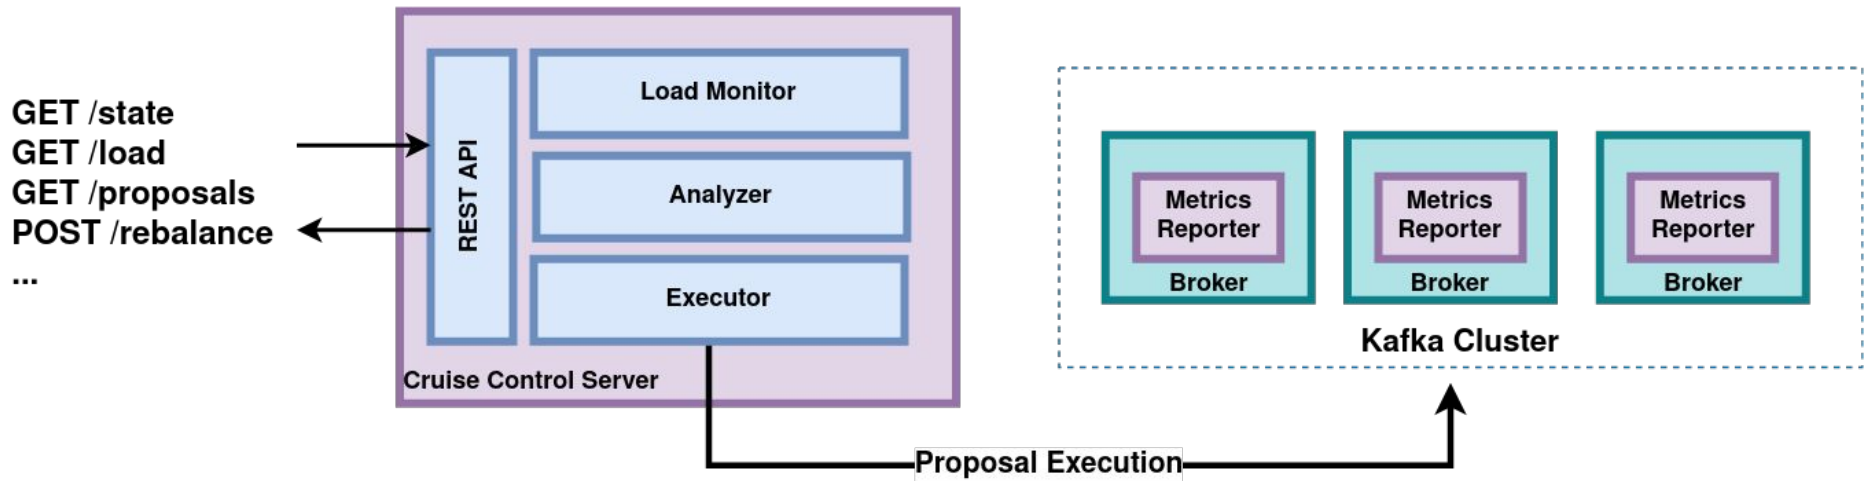


Cluster Workload Model

- ▶ Workload data of cluster resources
- ▶ Used to simulate partition movements

Broker	Topic/Partition		Disk/CPU		Network Rate	
Host	#Replicas	#Leaders	Disk	CPU	IN	OUT
my-cluster-broker-0	24	24	100 GB	8.05%	5 MBps	8 MBps
my-cluster-broker-1	22	22	56 GB	9.60%	6 MBps	12 MBps
my-cluster-broker-2	19	19	76 GB	8.52%	2 MBps	9 MBps

Cruise Control REST API



Cruise Control

Integration with Strimzi

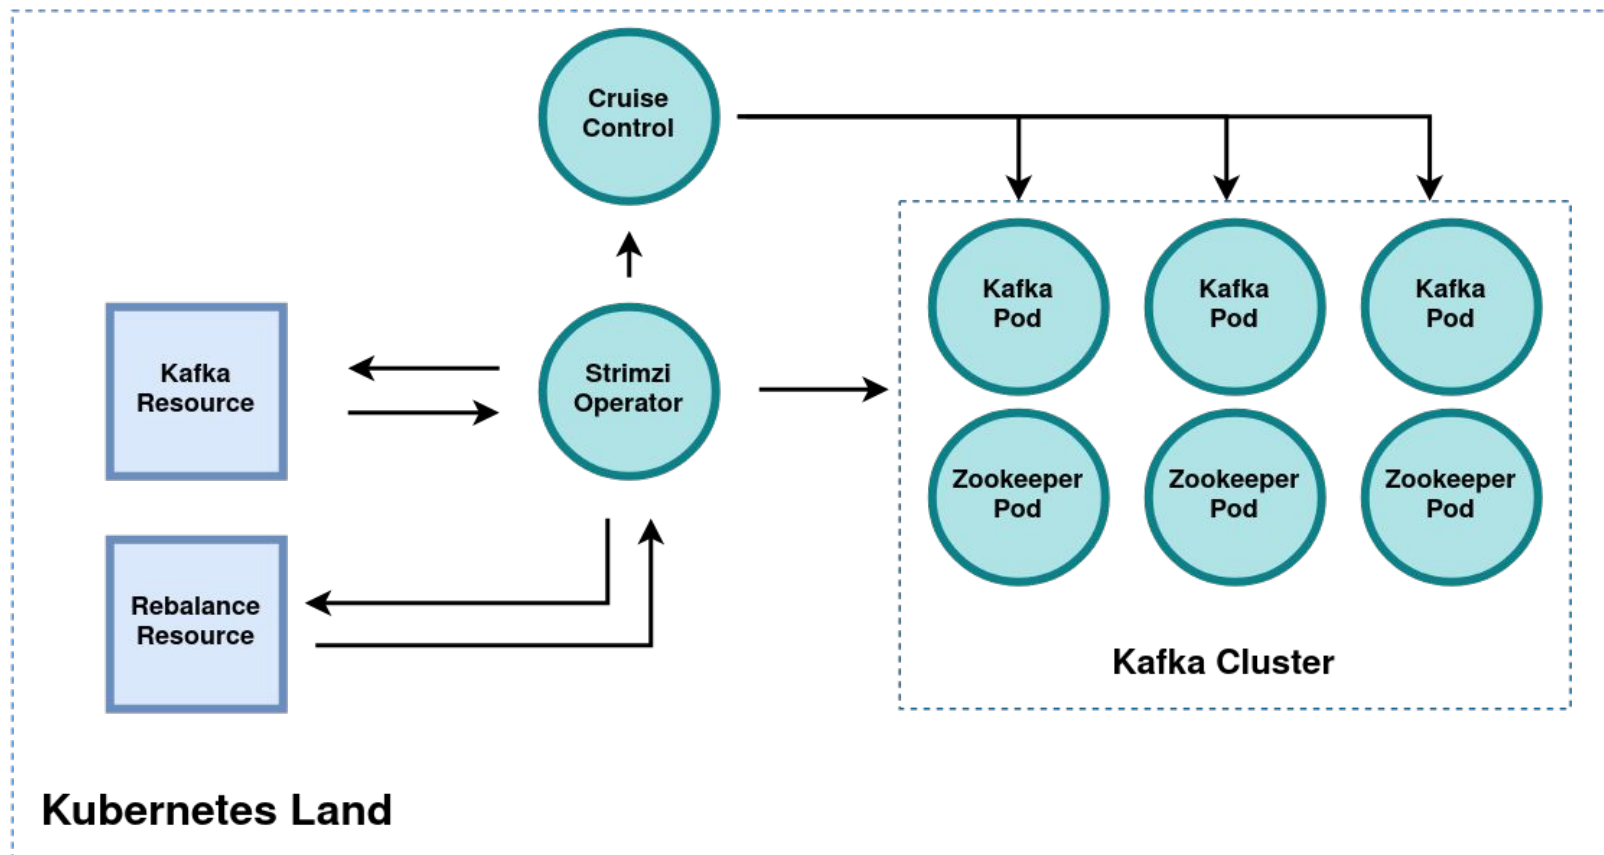
- ▶ Automated deployment
- ▶ Simple user interface



STRIMZI



Strimzi Architecture with Cruise Control



Kafka Resource with Cruise Control

Kafka
Custom
Resource



```
apiVersion: kafka.strimzi.io/v1beta2
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    ...
  cruiseControl:
    config:
      default.goals: >
        <path>.RackAwareGoal,
        <path>.ReplicaCapacityGoal
        ...
      cpu.balance.threshold: 1.1
      metadata.max.age.ms: 300000
      send.buffer.bytes: 131072
    brokerCapacity:
      disk: 100Gi
      cpuUtilization: 100
      inboundNetwork: 10000KiB/s
      outboundNetwork: 10000KiB/s
```


Cruise Control Goals


- ▶ Analyzer proposes load optimizations using selected goals
- ▶ Goals cover different dimensions of a rebalance
 - Rack awareness
 - Replica capacity
 - Resource capacity/utilization
- ▶ The combination and priorities of goals is customizable



Cruise Control Capacity Configurations

Configurable capacity limits for:

- Disk usage
- CPU utilization
- Network throughput



Kafka Rebalance Resource



apiVersion: kafka.strimzi.io/v1beta2

kind: KafkaRebalance

metadata:

name: my-rebalance

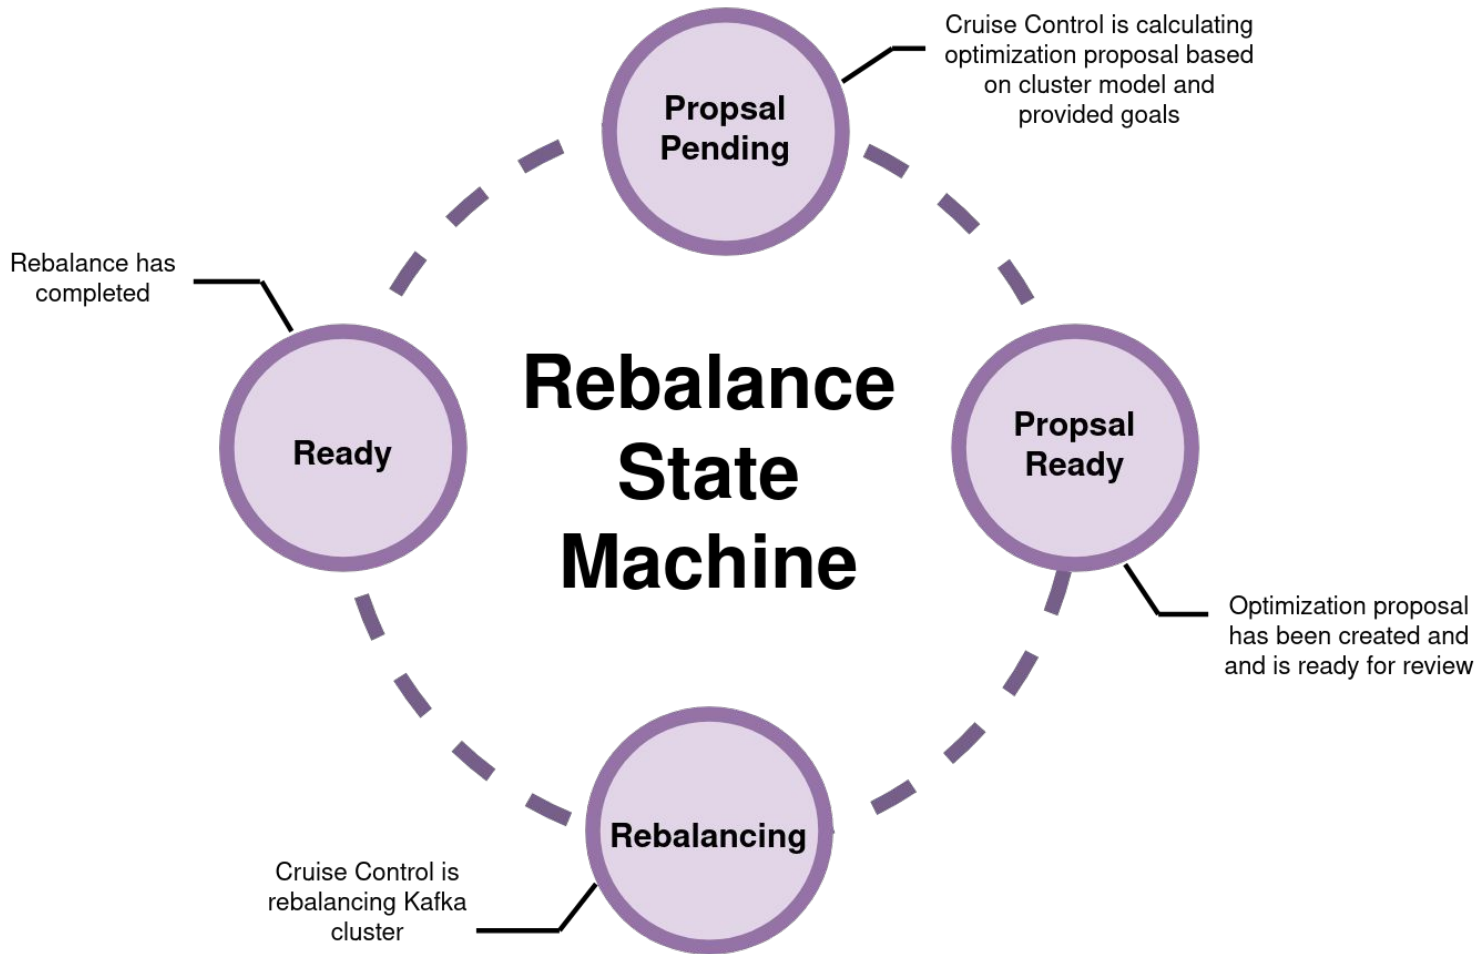
labels:

strimzi.io/cluster: my-cluster

spec:

goals:

- CpuCapacityGoal
- NetworkInboundCapacityGoal
- DiskCapacityGoal
- RackAwareGoal
- MinTopicLeadersPerBrokerGoal
- NetworkOutboundCapacityGoal
- ReplicaCapacityGoal



Kafka Rebalance Status



Status:

Conditions:

Last Transition Time: 2020-05-19T13:50:12.533Z

Status: ProposalReady

Type: State

Observed Generation: 1

Optimization Result:

Data To Move MB: 0

Excluded Brokers For Leadership:

Excluded Brokers For Replica Move:

Excluded Topics:

Intra Broker Data To Move MB: 0

Monitored Partitions Percentage: 100

Num Intra Broker Replica Movements: 0

Num Leader Movements: 0


Num Replica Movements: 26

On Demand Balancedness Score After: 81.8666802863978

On Demand Balancedness Score Before: 78.01176356230222

Recent Windows: 1

Session Id: 05539377-ca7b-45ef-b359-e13564f1458c



DEMO

What's Next?

Cruise Control Integration Roadmap

- ▶ API authentication + authorization
- ▶ Intra broker data balancing
- ▶ Changing topic replication factor



Questions?

Thank You!

 linkedin.com/company/red-hat

 youtube.com/user/RedHatVideos

 facebook.com/redhatinc

 twitter.com/RedHat