




TOPIC

Operate Kubernetes Workloads:

Extend the platform with the operator pattern!

Paolo Patierno, Principal Software Engineer @RedHat





#CodeGen2021

@cloudgen_verona

Who am I?



apiVersion: v1

kind: PrincipalSoftwareEngineer

metadata:

name: Paolo Patierno

namespace: Red Hat, Messaging & Data Streaming

labels:

cncf/maintainer: Strimzi

eclipse/committer: Vert.x, Hono & Paho

microsoft/mvp: Azure

annotations:

family: dad of two, husband of one

sports: running, swimming, motogp, vr46, formula1, ferrari, ssc napoli

community: cncf napoli, devday

spec:

replicas: 1

containers:

- **image:** patiernohub.io/paolo:latest



@ppatierno



“ A system for ...”

“ ... automating deployment ...”

“ ... scaling ...”

“ ... management ...”

“ ... of containerized applications ...”

“ It's like a Linux kernel ... but for distributed systems”



Container
scheduling

Self healing

Secret &
configuration
management

Service
discovery

Horizontal
scaling

Load balancing

Storage
orchestration

Automated
rollout/rollback

Batch
execution





“

How does Kubernetes handle scaling, rollout, batch execution and so on?



- Don't use Pod(s) ... let's use something more sophisticated!
- ReplicaSet
 - Guarantees a specific number of running replicas
 - Spins pods, based on a template, if there are not enough
 - Deletes pods if too many match the selector
- Deployment
 - It's based on ReplicaSet (used to run replicas)
 - Adds extra layer for rollout and rollback
- ... and more with StatefulSet, Job, DaemonSet ...



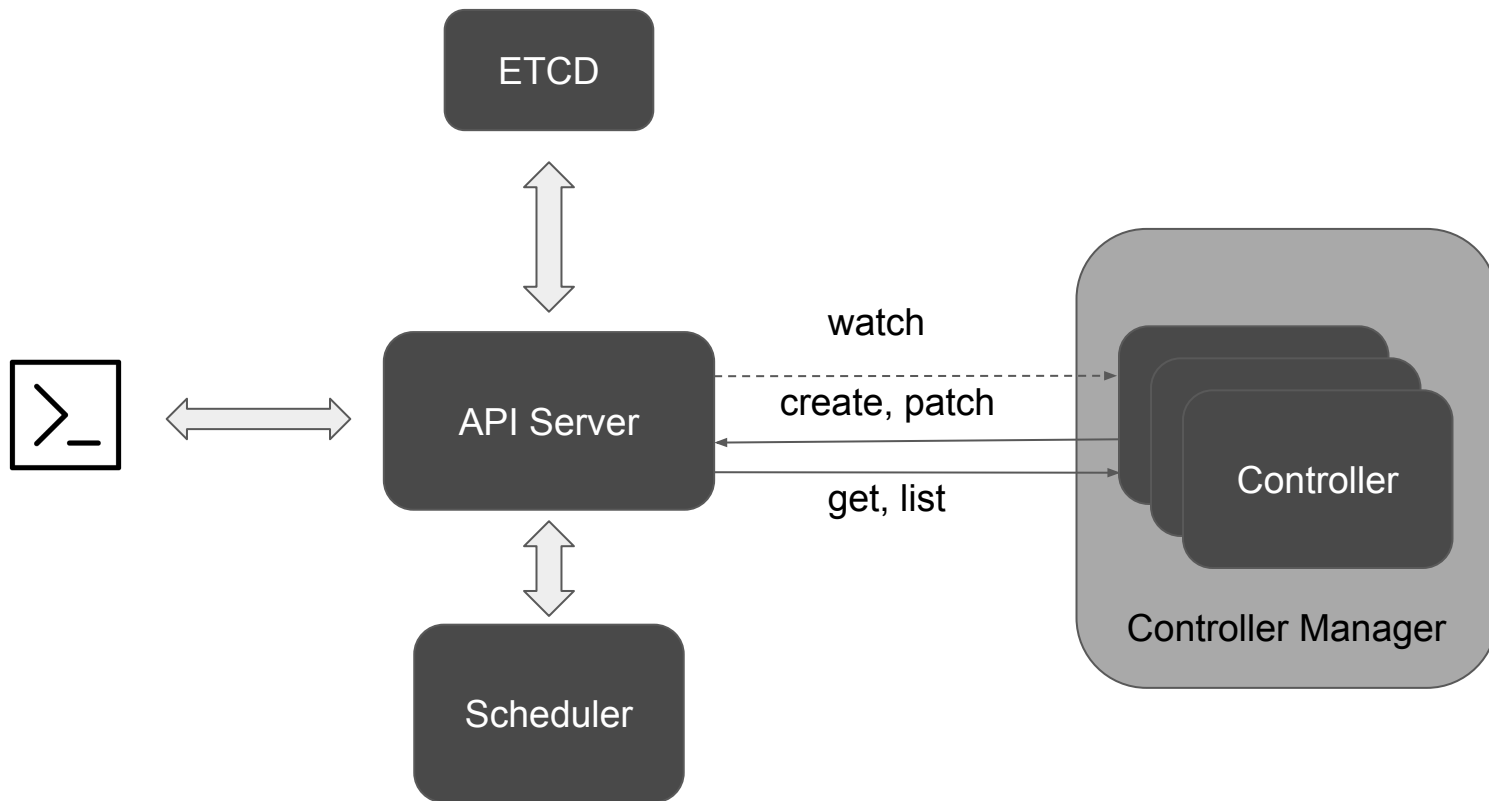
“

How does Kubernetes do the trick? What's behind these special resources?




- A controller ...
 - ... tracks a Kubernetes resource/object
 - ... makes sure that the object specification, desired state, matches the current state on cluster
 - ... interacts with the API server
 - ... acts in a control loop
- Built-in controllers
 - ReplicaSet
 - Deployment
 - ... the Kubernetes controller manager runs them

Kubernetes controllers





- **Observe**
 - Watch for resource/object creation or changes
- **Analyze**
 - Check that the resource/object desired state (“spec”) reflects the current state on the cluster
- **Act**
 - Makes the needed changes



ReplicaSet controller



```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-application
          image: quay.io/ppatierno/my-application:latest
```

my-replicaset-bf5zv

my-replicaset-1tf5a

my-replicaset-gb65f

ReplicaSet controller



```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod-1
  labels:
    app: my-app
spec:
  containers:
    - name: my-app
      image: quay.io/

apiVersion: v1
kind: Pod
metadata:
  name: my-pod-2
  labels:
    app: my-app
spec:
  containers:
    - name: my-application
      image: quay.io/ppatierno/my-application:latest
```

my-replicaset-bf5zv

my-replicaset-1tf5a

my-replicaset-gb65f

~~my-pod-1~~

~~my-pod-2~~

ReplicaSet controller



my-pod-1

my-pod-2

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod-1
  labels:
    app: my-app
spec:
  containers:
    - name: my-app
      image: quay.
      apiVersion: v1
      kind: Pod
      metadata:
        name: my-pod-2
        labels:
          app: my-app
      spec:
        containers:
          - name: my-application
            image: quay.io/ppatierno/my-application:latest
```


ReplicaSet controller



```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-application
          image: quay.io/ppatierno/my-application:latest
```

my-pod-1

my-pod-2

my-replicaset-65rt3




“

Humans build
complex applications
with these “bricks”

Human operating Kubernetes workloads



- Human operator knows ...
 - ... about a complex application or service
 - ... how it behaves internally
 - ... how to deploy it
 - ... how to react and fix issues with it






“

How to automate
operating applications?
The Operator pattern!

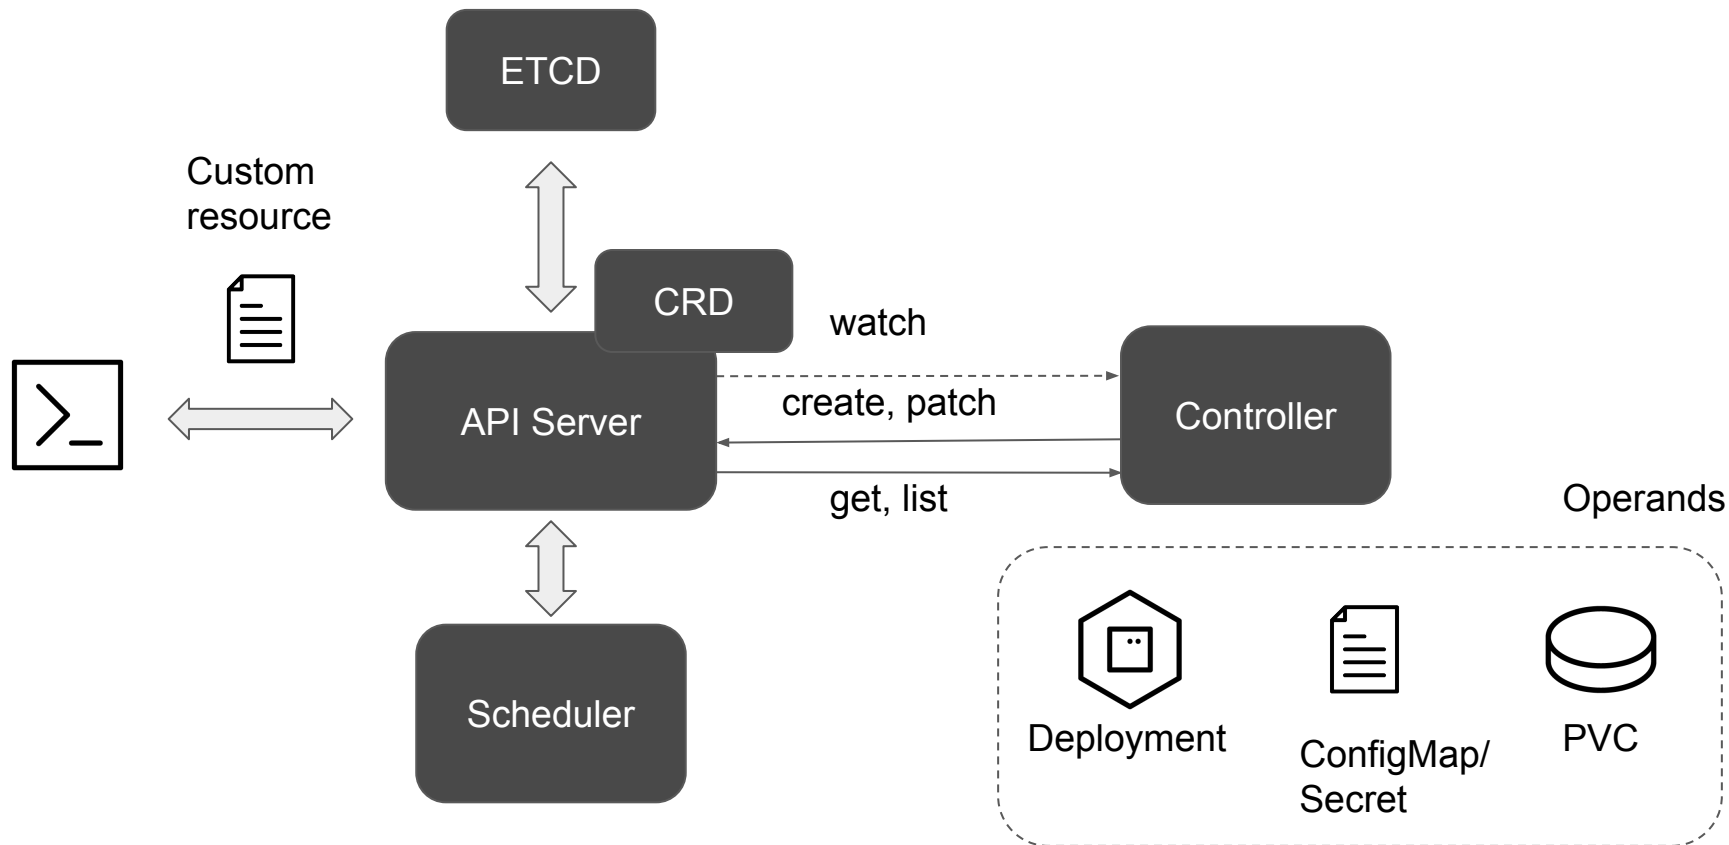
Operator



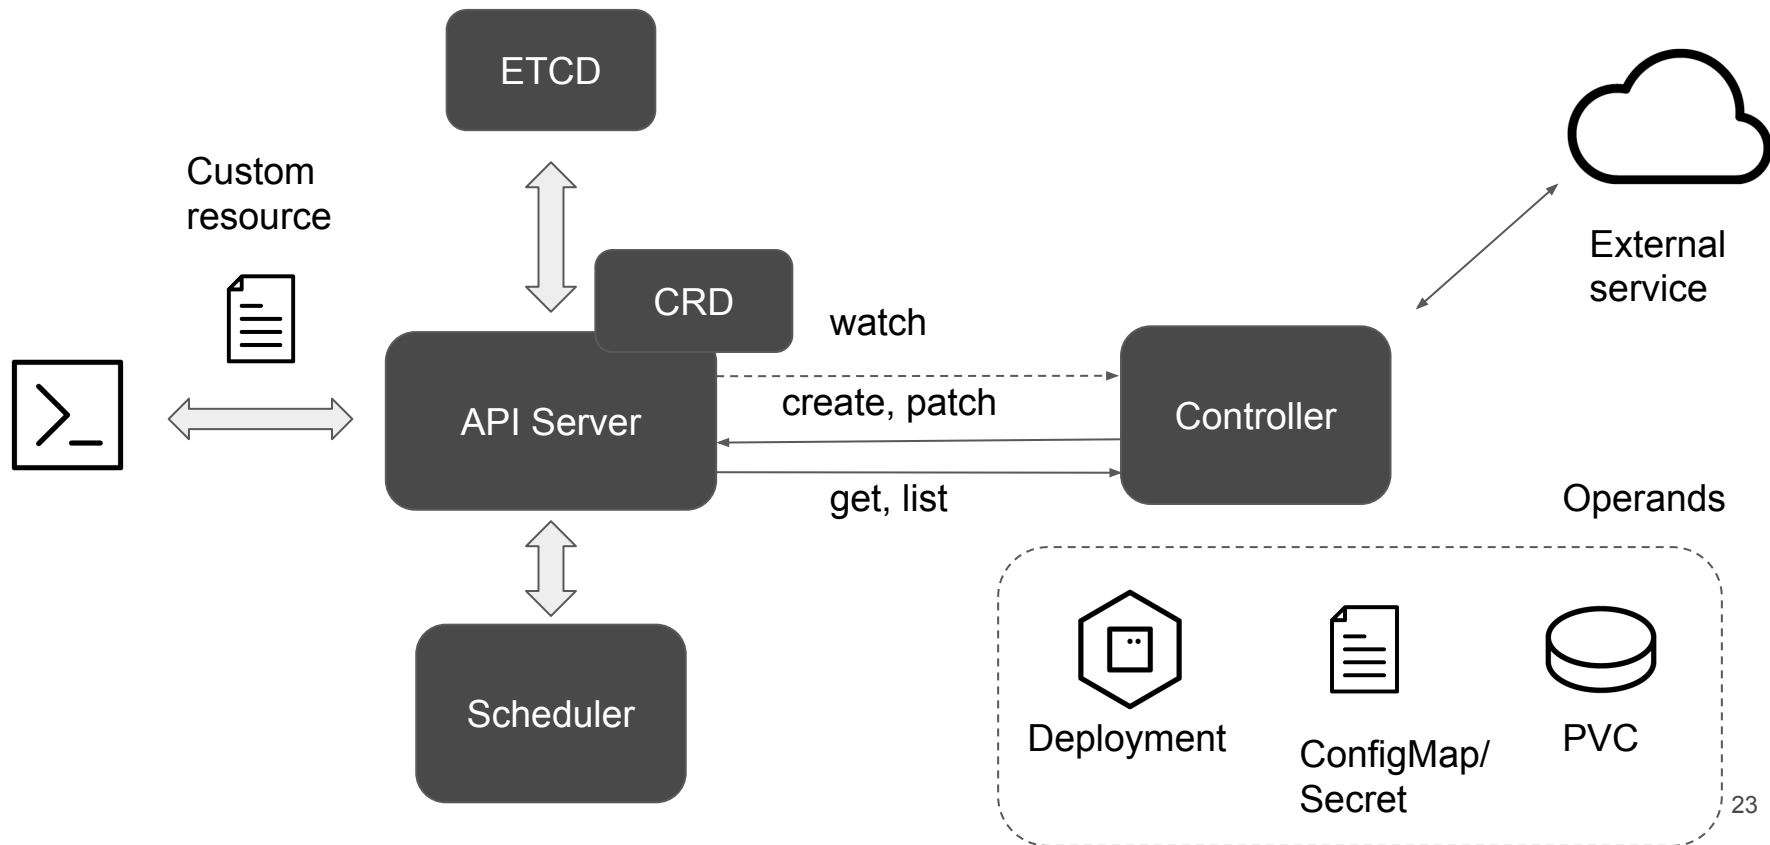
- It's yet another containerized application!
- Has the knowledge of a specific business domain
- Manage the application lifecycle
- Leverage CRDs (Custom Resource Definition) to extend API server
- Takes care of one (or more) custom resources/objects
 - By having a controller for each resource
 - Creating native Kubernetes resources ... aka "operands"
 - Leveraging built-in controllers via API server interaction



Operator



Operator



From the Custom Resource Definition ...






- It's a ... Kubernetes resource!
- Declare a new Kubernetes “kind”
 - group
 - versions
- Define the new “kind” structure using an OpenAPI schema
 - spec
 - status

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name: kafkas.kafka.strimzi.io
spec:
  group: kafka.strimzi.io
  names:
    kind: Kafka
    listKind: KafkaList
  ...
  versions:
  - name: v1beta2
    schema:
      openAPIV3Schema:
        type: object
        properties:
          spec:
            # spec definition with for
            # the custom resource
            ...
          status:
            # status definition reported back
            # in the custom resource
```


... to the Custom Resource



```
apiVersion: kafka.strimzi.io/v1beta2
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    version: 2.8.0
    replicas: 3
    listeners:
      - name: plain
        port: 9092
        type: internal
        tls: false
      - name: tls
        port: 9093
        type: internal
        tls: true
    config:
      log.message.format.version: "2.8"
      inter.broker.protocol.version: "2.8"
      ...
    storage:
      type: ephemeral
    zookeeper:
      replicas: 3
      storage:
        type: ephemeral
status:
  ...
  ...
```



StatefulSet



ConfigMap/
Secret



PVC



NetworkPolicy



- Each “internal” controller watches a corresponding custom resource
 - Create/update/delete the corresponding “operands” as native Kubernetes resources but ... could be other custom resources for other operators :-)
- Watch the “operands”
 - They can be touched only by operator controllers ... not humans
 - Reverts back any manual changes
 - “Owned by” the custom resource, leveraging Kubernetes garbage collection
- Can interact with external service
 - A custom resource could be related to handle a non-Kubernetes service (i.e. Azure Service Operator, ...)



- **Configuration**
 - Simplifies configuration of a complex application in a single and meaningful place
- **Installation and upgrade**
 - Executes installation but mostly is able to run a potential complex upgrade procedure
- **Security**
 - Allows to take care of security related stuff like encryption, access control, network policies
- **Scaling**
 - Enable scaling and maybe, depending on application, autoscaling



“

Why? What about
Helm Charts?



- “Package manager” for Kubernetes
 - Charts = template + values
- Rely on Kubernetes built-in resources (i.e. Deployment, ConfigMap, ...)
- Simplify to write YAMLS with parameters via templating
- Simplify day-1 operation, for deploying applications
- Key problems fixed
 - Deploy same application with different configuration
 - Deploy same application on different environments



- Control “life cycle” of Kubernetes workloads
 - Operator = CRD(s) + Controller(s)
- Extend the Kubernetes API with CRDs (Custom Resource Definitions)
- Simplify to write one (or a few) “custom resource” related YAMLs
- Acting since day-1 to day-2 operation from deployment to upgrades, through manage
- ... deployable via an Helm Charts :-)
 - Helm guarantees CRDs are installed before operator
 - Operator configurable via values or ConfigMap



“

You convinced me!
How to start?



- Operator framework
 - Toolkit to manage Kubernetes operator
 - SDK for writing operators in Golang
 - Operator Lifecycle Manager (OLM) to handle ... operators
 - Operator registry to provide operators to OLM
 - ... and much more
 - <https://sdk.operatorframework.io/>
- Java Operator SDK
 - Writing operators in Java
 - Support for Quarkus
 - <https://javaoperatorsdk.io/>



OPERATOR FRAMEWORK



[JAVA OPERATOR SDK]

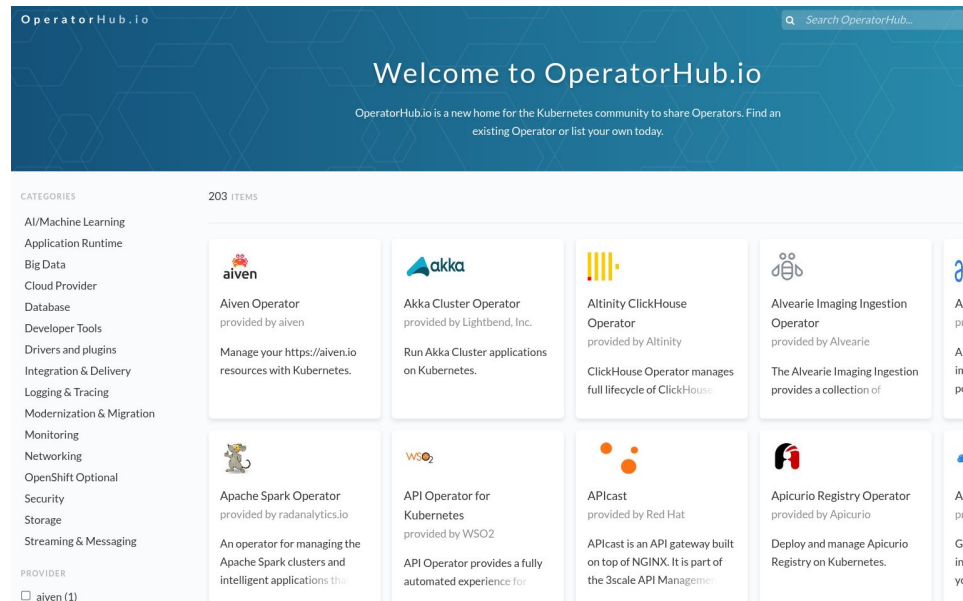
Operator Lifecycle Manager (OLM)



- An operator to rule them all!
 - An operator handling ... operators!
- Define a packaging format
- Expose a catalog
- Handle operator updates
- Provided out of the box with OpenShift (enterprise Kubernetes distribution by Red Hat) but installable on vanilla Kubernetes



- Home for Kubernetes operators
 - A lot of categories (Database, Streaming & messaging, Logging & Tracing, ...)
- Installation via Helm Charts or YAML files
- You can develop your own and provide to the community
- <https://operatorhub.io/>




Strimzi: The Apache Kafka operator



- Open source project licensed under Apache License 2.0
- Focuses on running Apache Kafka on Kubernetes
 - Container images for Apache Kafka, Apache ZooKeeper and other components
 - Operators for deploying, managing and configuring Kafka clusters
- Provides a Kubernetes-native experience
 - Not only Kafka clusters, but also users, topics and the rest of Kafka ecosystem
- CNCF sandbox project since September 2019
- <http://strimzi.io>



STRIMZI



Demo time



Thank you

Any questions?



ppatierno



@ppatierno



<https://www.linkedin.com/in/paolopatierno/>