

# Event Streaming and Message Queues: which one should you use?

Domenico Francesco Bruscino & Paolo Patierno



Linux Day Napoli 2022



# Who are we?



Senior Principal Software Engineer  
Apache Kafka and Strimzi  
**@ppatierno**



Senior Software Engineer  
Apache ActiveMQ Artemis  
and ArtemisCloud  
**@bruscinodf**

# Let's start from the beginning : Event vs Message

A thing that happens

A fact making the history

A record to retain for longer time

Replayable to rebuild status

A business process to run

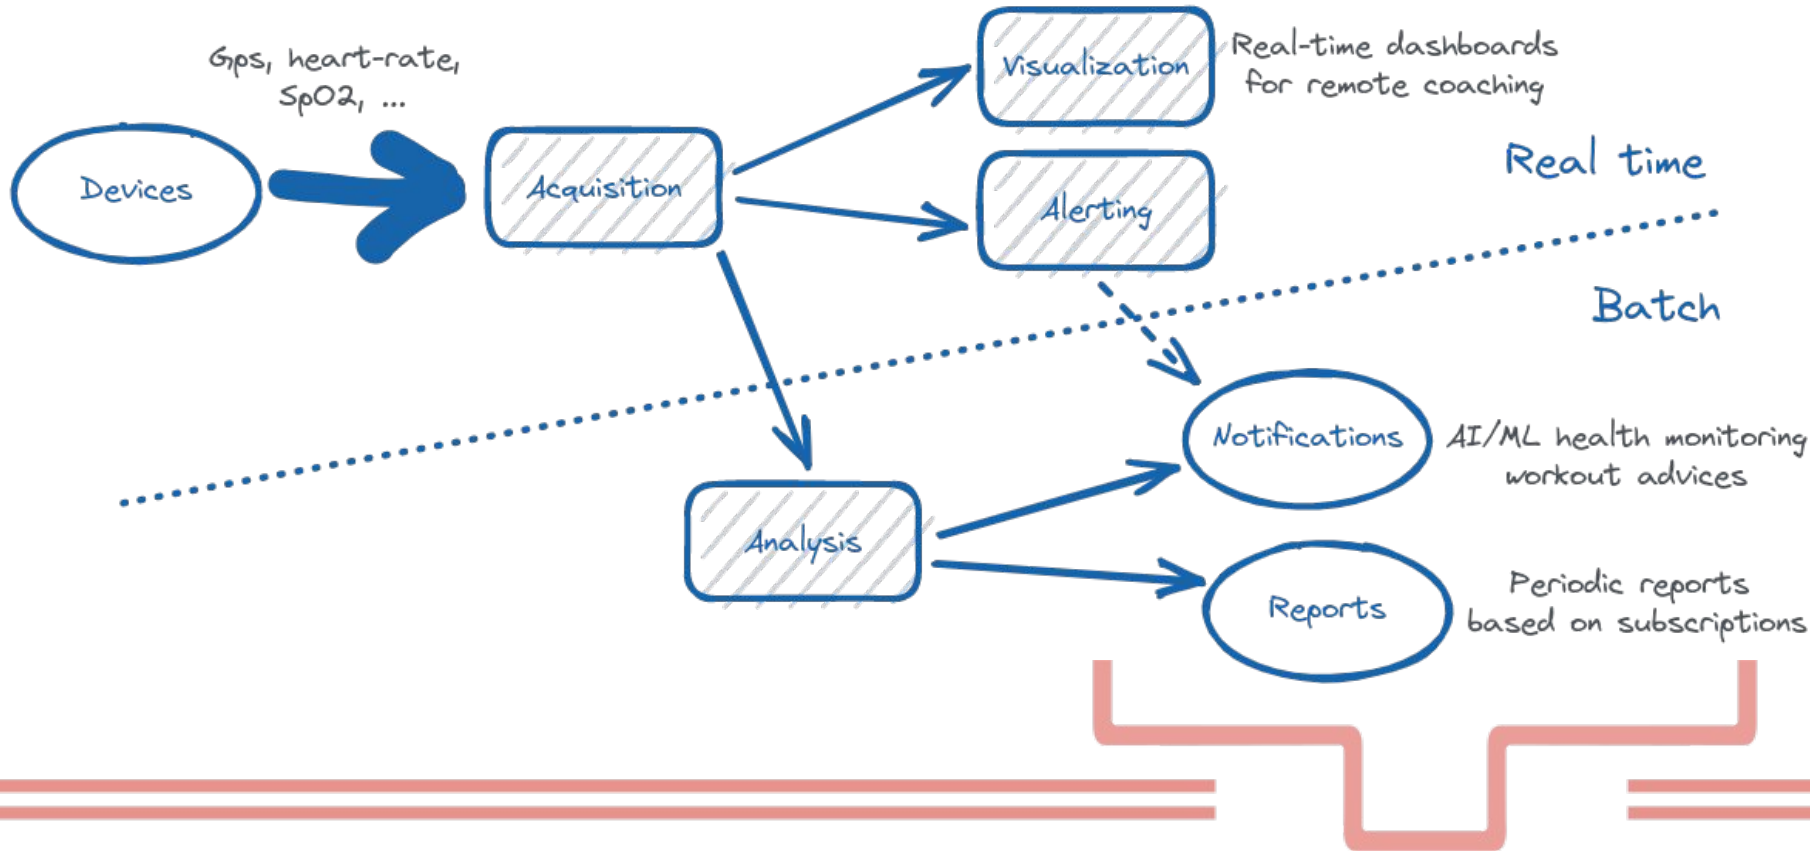
A task on a to-do list

A command to be executed


To be deleted when it's done

Fun fact ... Events **ARE** (anyway) Messages :-)


# Welcome to Strada



# Deep into the workflow



# Real-time vs Batch: challenges



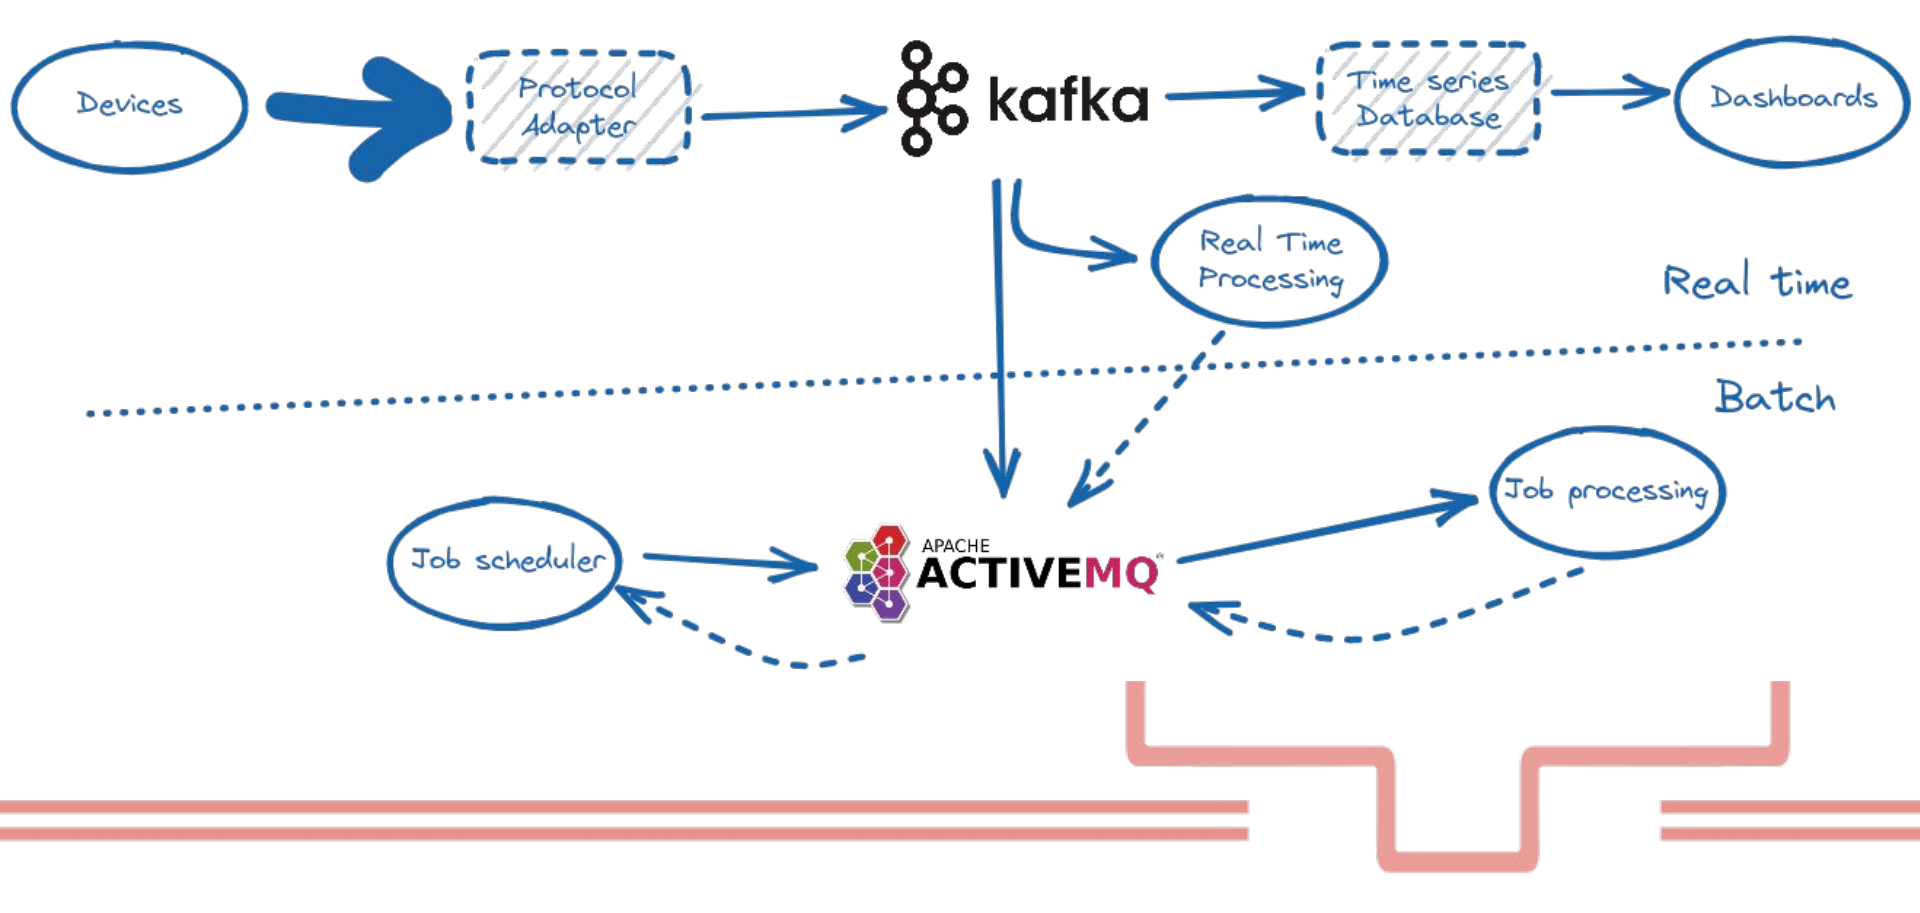
How do we solve this?

Is using one technology enough?

Should we use the right tool for different goals?



# Welcome to ... Apache Kafka & Apache ActiveMQ Artemis






# Apache Kafka

- Messaging system, data/event streaming platform ... a commit log
- Developed at LinkedIn back in 2010, open sourced in 2011
- Designed to be fast, scalable, durable and available
- Distributed by nature
- Data partitioning (sharding)
- High throughput / low latency
- Ability to handle huge number of consumers
- Dumb broker, smart client

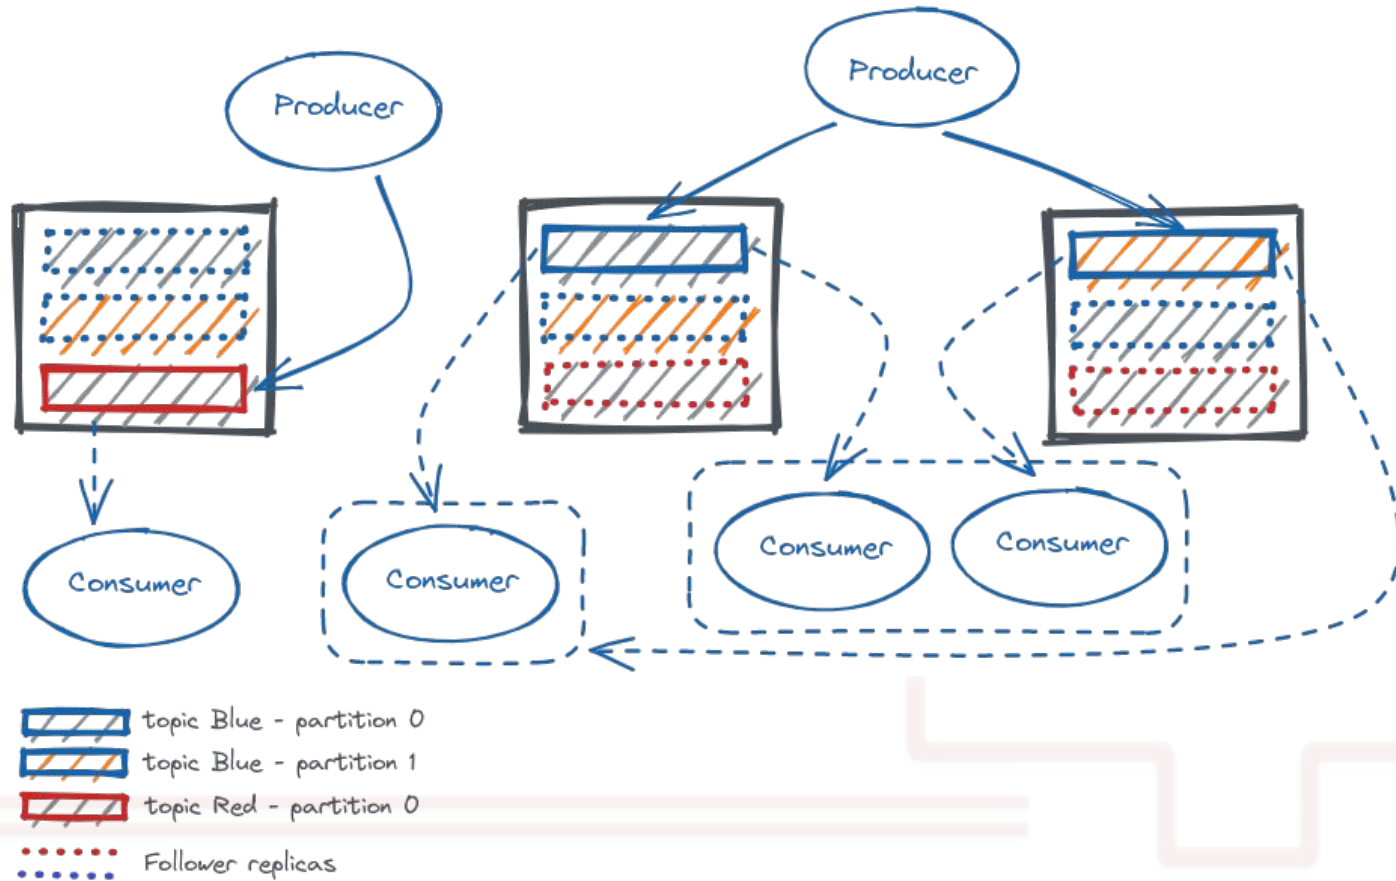


# Apache Kafka : the basics



- Consumer keeps track of the offset of the last read message

# Apache Kafka : brokers & replicas




# Apache Kafka ... good at

- Long-time durable storage
  - Messages stored and available for re-processing
- High throughput, Low latency
  - No complex process on the broker, just get and store the message
  - Consumers keep track of offset, Producers select destination and do batching
- Scalability
  - Topic partitions allow to scale consumers
- High Availability
  - Topic partitions are replicated across brokers




# Apache Kafka ... good at : scalability

- Partitioning is the key for scalability
- Partitions spread across consumers
- Multiple consumer group allows pub/sub




# Apache Kafka ... good at : high availability



\* actually possible from followers as well

# Apache Kafka ... good at : high availability



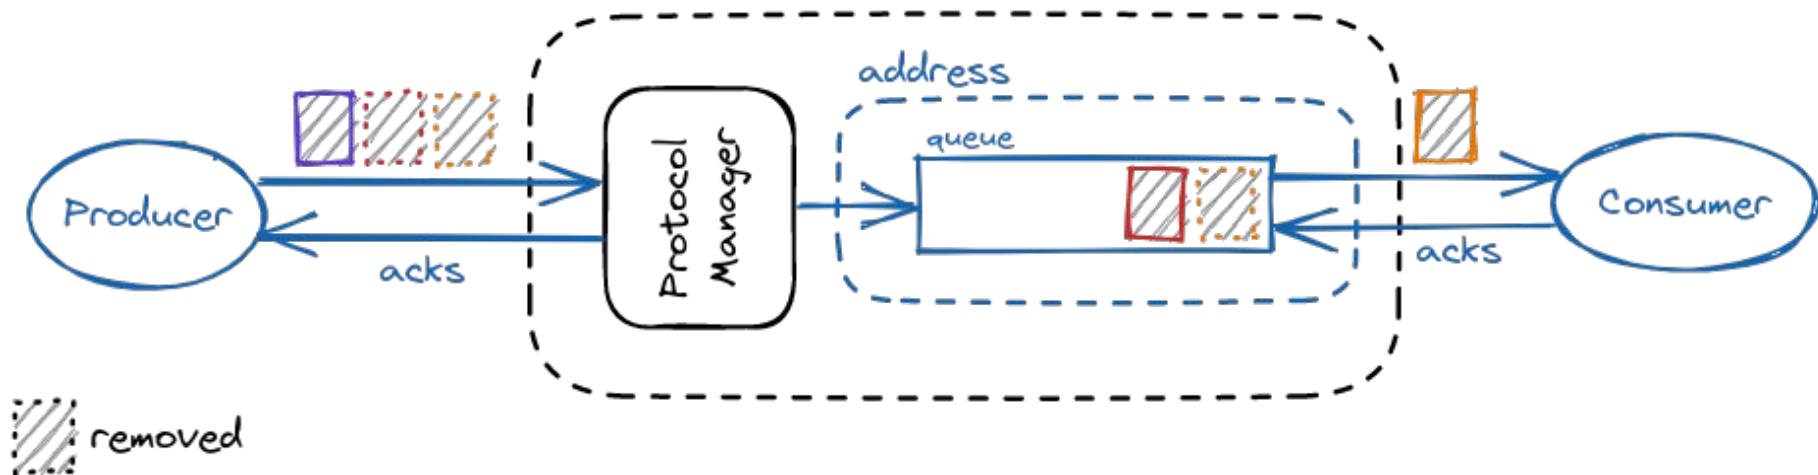
# Apache ActiveMQ Artemis

- Standards based (JMS 1.x/2.x/3.x)
- Multi protocol (AMQP, MQTT, STOMP, OpenWire)
- Powerful and flexible address model
- Cluster ready and highly available
- Trusted and used in many highly critical use cases
- Dumb client, smart broker



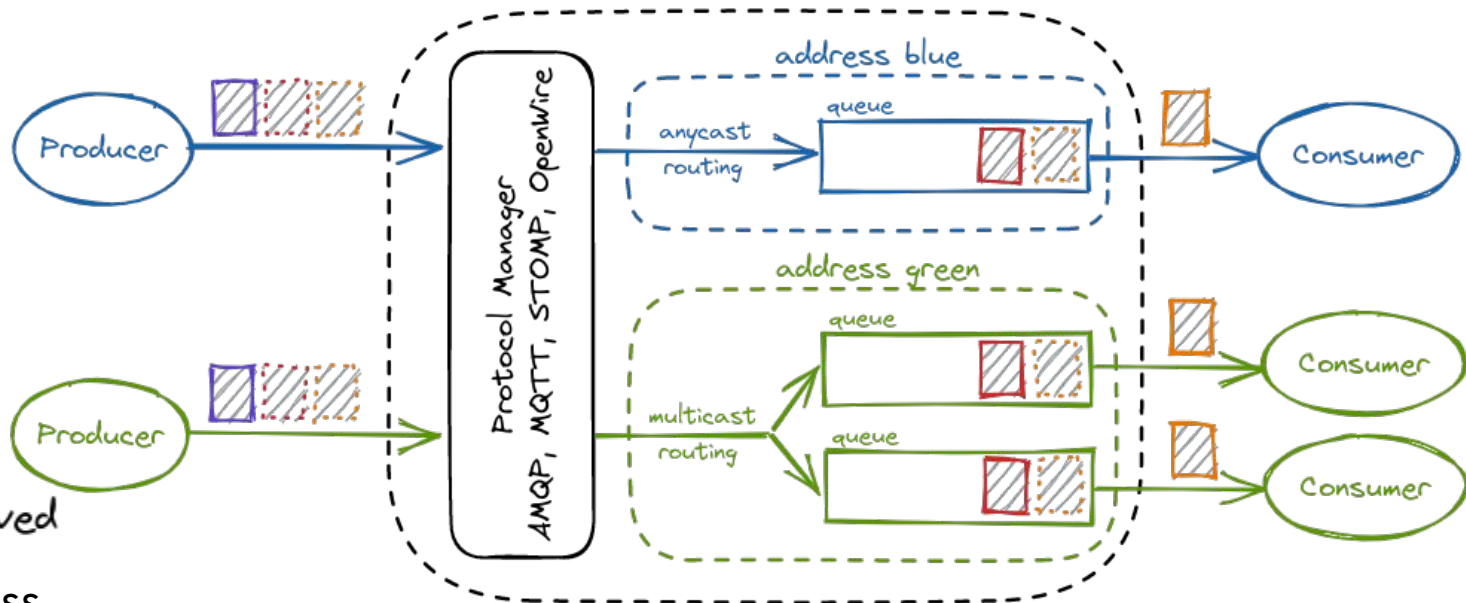


# Apache ActiveMQ Artemis : the basics




- Broker sends acknowledgments to confirm messages received by producers
- Broker removes messages from queues when it receives acknowledgments from consumers

# Apache ActiveMQ Artemis address model



- address
- queue
- routing type


# Apache ActiveMQ Artemis ... good at : message priority



Messages are enqueued in order of priority

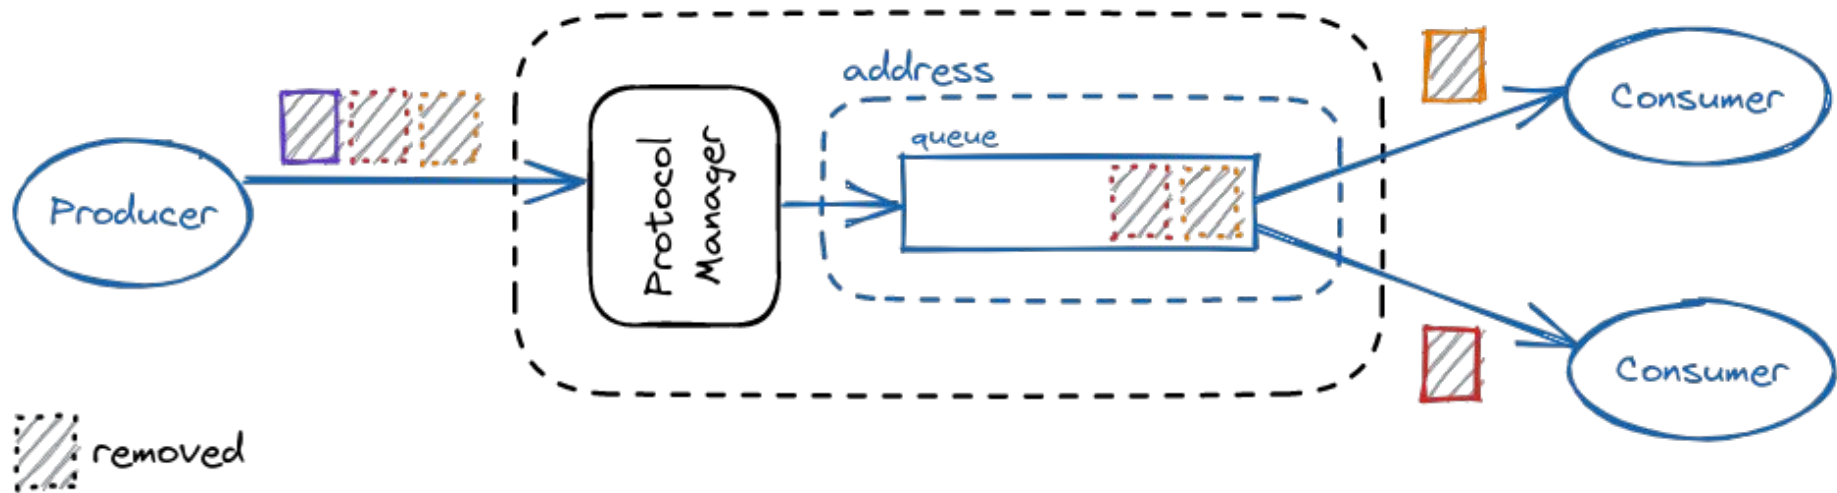


# Apache ActiveMQ Artemis ... good at : request/reply



Broker automatically create a temporary address and queue for the reply

# Apache ActiveMQ Artemis ... good at : competing consumers



- Message order is preserved
- Delivered messages are removed
- No required change to attach a new consumer

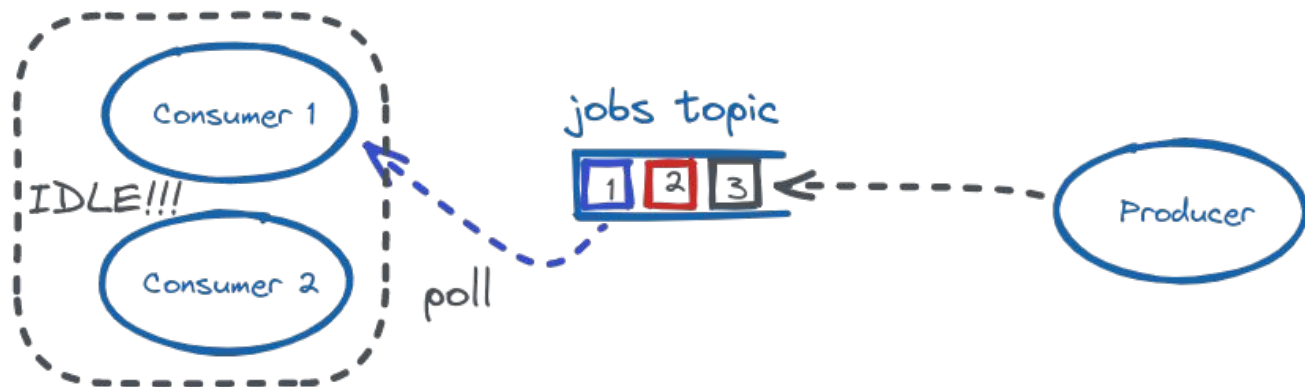
## Why these choices?

What's wrong with using Kafka for job processing?

Or with using Artemis ActiveMQ for event streaming?

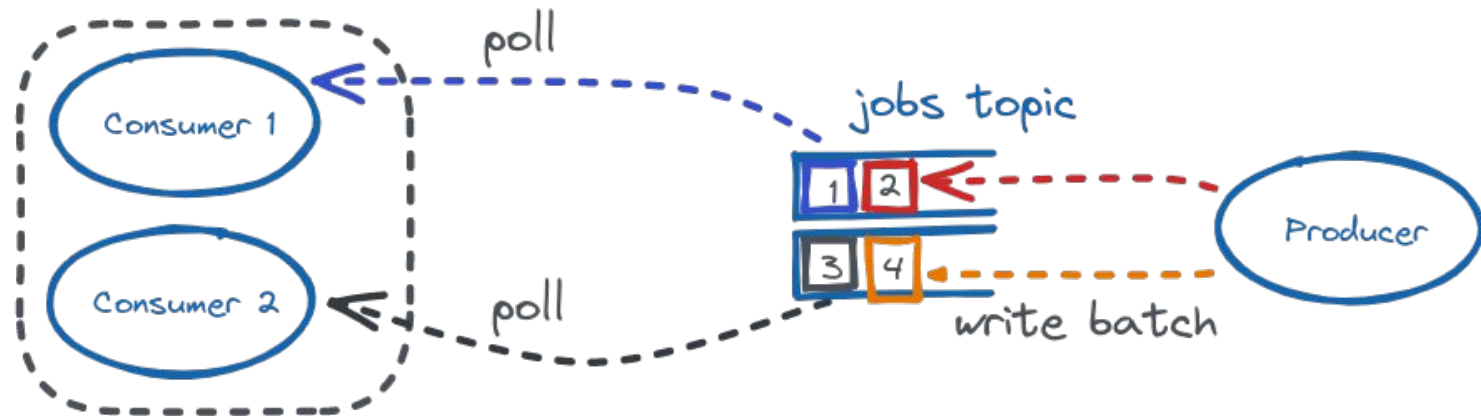


# Apache Kafka ... not good at : job ordering



- Order is per partition not per topic
- Only one consumer per partition to get each job and execute it
- To keep jobs ordering, we cannot scale
- Jobs get stuck if consumer is slow, or bad message (head of line blocking)

# Apache Kafka ... not good at : job ordering

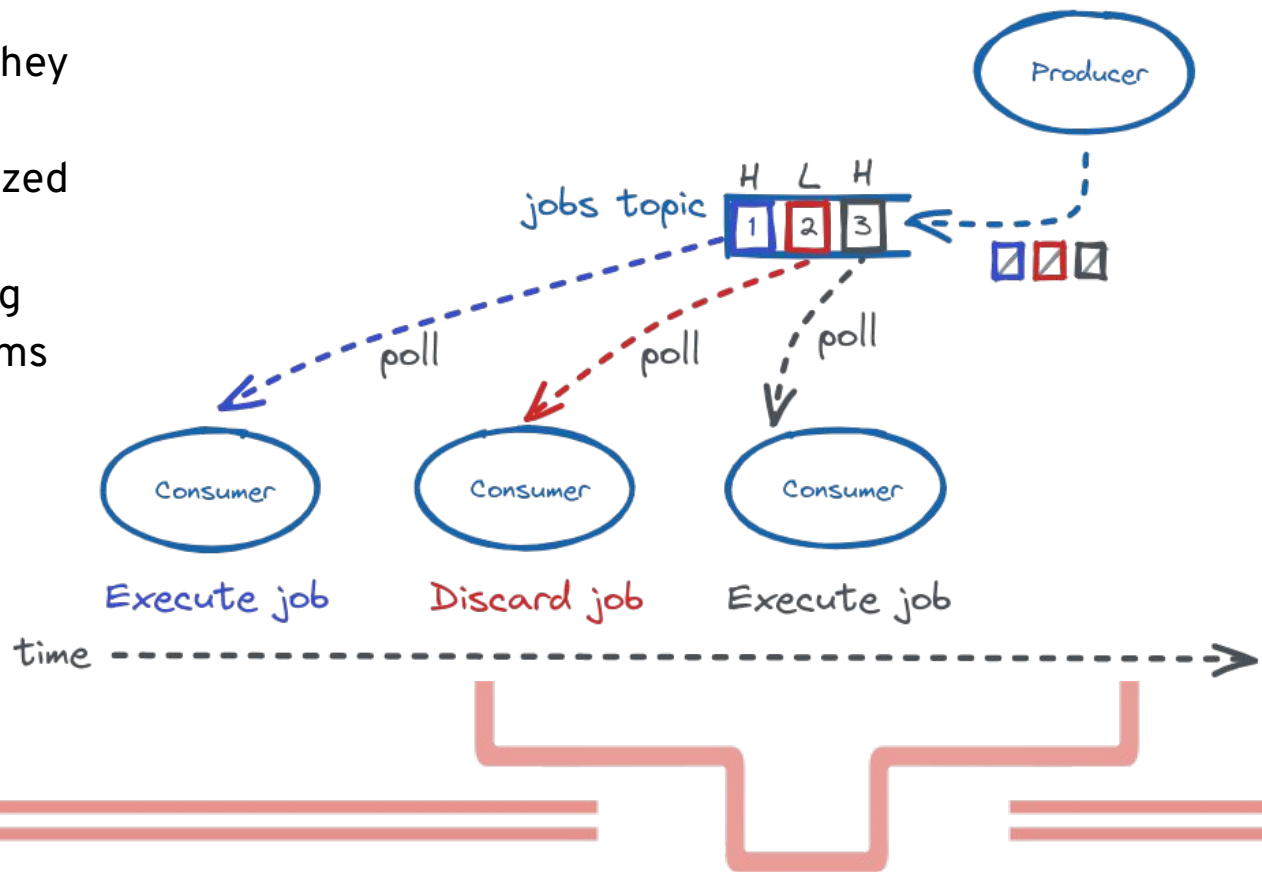


- Trying to scale consumers but ...
- ... Producer writes in “batch”
- ... More partitions can screw up the jobs order



# Apache Kafka ... not good at : message priority

- Messages are appended as they arrive
- Broker cannot return prioritized messages
- Need for an application doing the filtering (i.e. Kafka Streams API based)

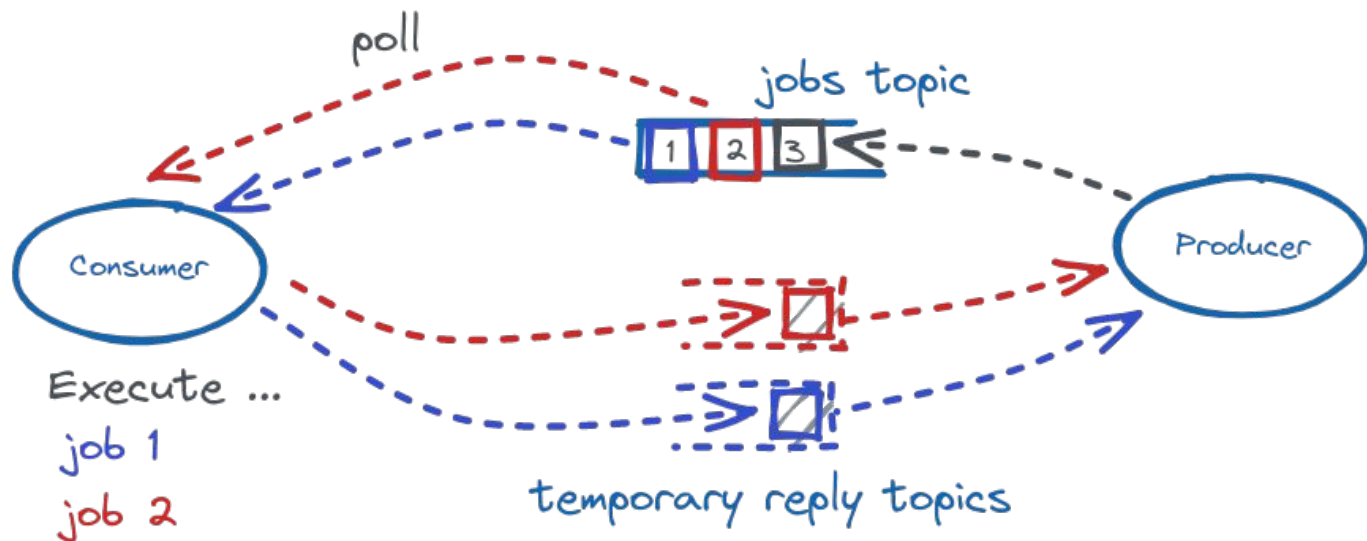


# Apache Kafka ... not good at : request/reply

- Request/Reply pattern
  - Thought for an event-driven architecture which doesn't cope well with a request/reply
  - Correlate response with request is not native
    - No correlation-id
    - No reply-to for creating a “temporary” topic for reply
- Onerous reply topic(s) handling




# Apache Kafka ... not good at : request/reply




- “Temporary” topic creation is onerous
- Disk allocation, replication set up, ...

# Apache Kafka ... not good at : request/reply




- Error prone, re-reading same acks
- Higher level logic needed

# Apache ActiveMQ Artemis ... not good at : horizontal scaling




Broker 1 needs to process messages before routing

# Apache ActiveMQ Artemis ... not good at : horizontal scaling



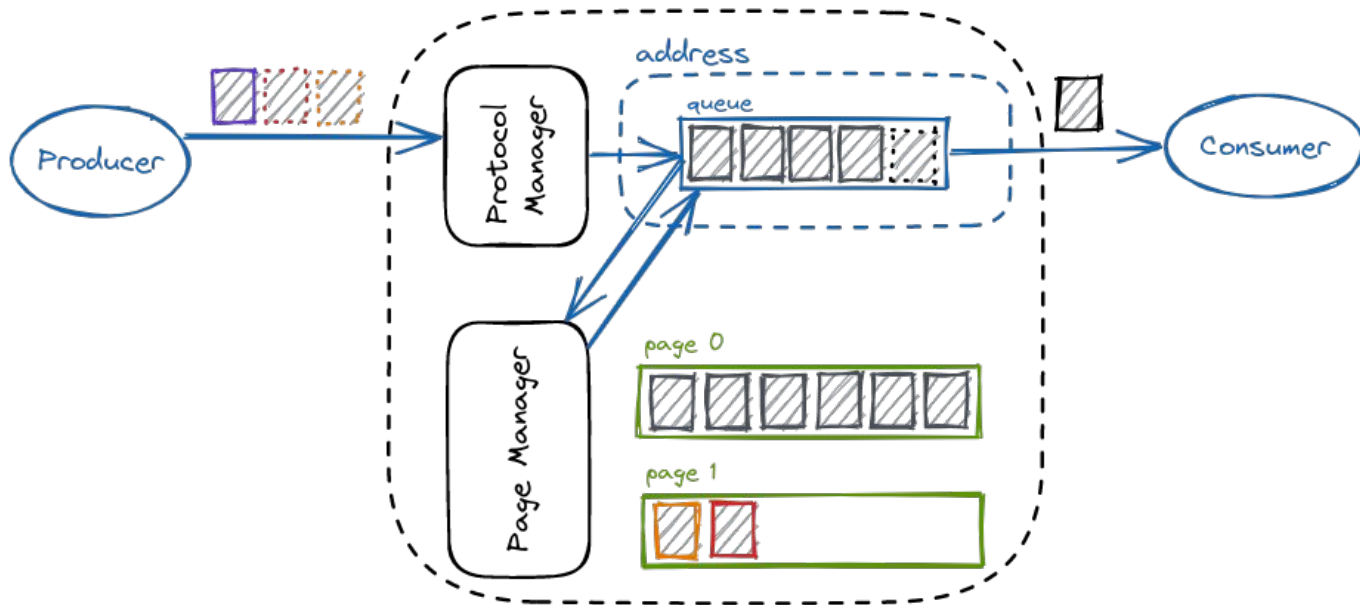
Broker 2 needs to redistribute messages

# Apache ActiveMQ Artemis ... not good at : horizontal scaling



Messages lose order after redistribution


# Apache ActiveMQ Artemis ... not good at : long-time storage



Fast producers and slow consumers cause paging

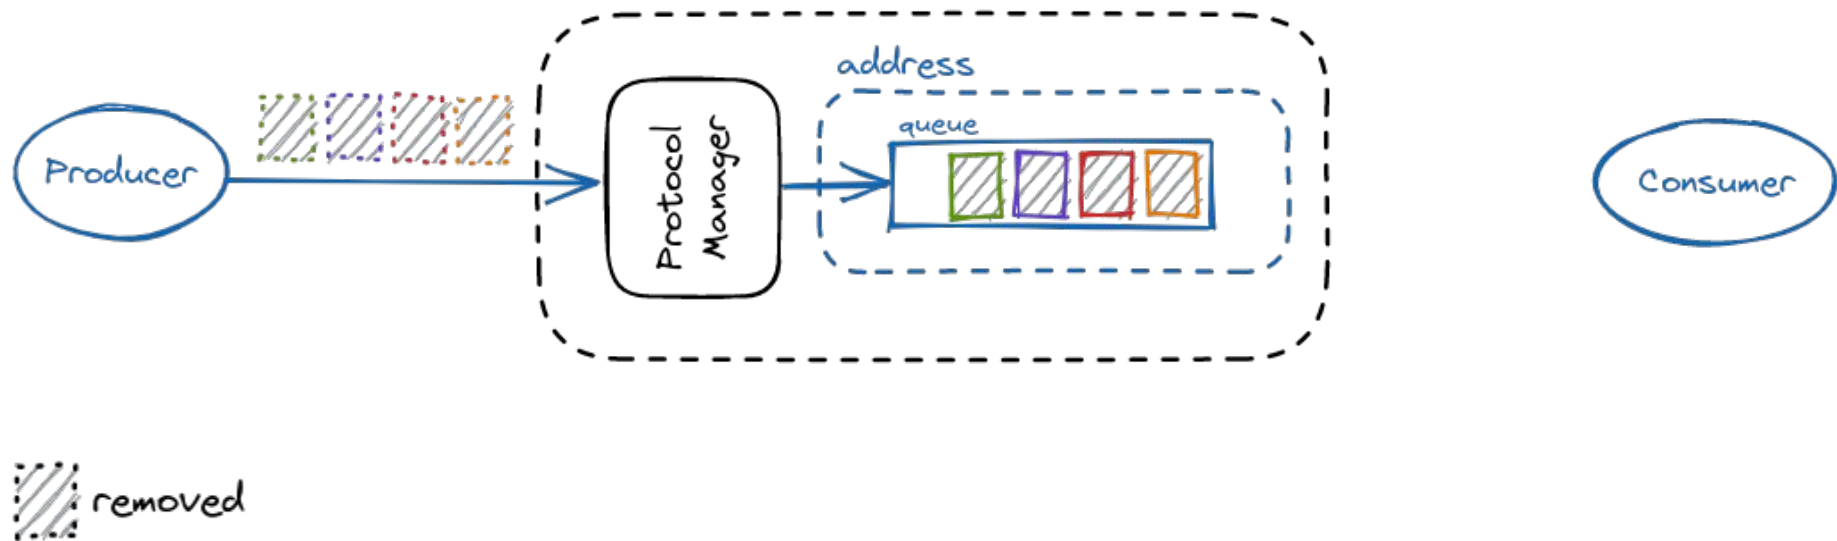


# Apache ActiveMQ Artemis ... not good at : long-time storage




- The queue pointer is managed by the broker
- Non destructive queue messages are not removed

# Apache ActiveMQ Artemis ... not good at : long-time storage



- The queue pointer is lost when the session is closed

# Apache ActiveMQ Artemis ... not good at : long-time storage



 removed

- The queue pointer starts from the beginning for new sessions



# Clemens Vasters (Microsoft) says ...

WM> World  
Congress  
2022

[]

//

Event Streaming is not "modern" and Queues  
are not "traditional"

Both are patterns of state-of-the art  
messaging infrastructures.



Photo: ...



Thanks!

