

Managing Kubernetes workloads

Extend the platform with operators

Paolo Patierno

Kate Stanley



Paolo Patierno

- ▶ Senior Principal Software Engineer at Red Hat
- ▶ CNCF Ambassador
- ▶ Strimzi maintainer
- ▶ Formula 1 & MotoGP addicted



Kate Stanley

- ▶ Principal Software Engineer at Red Hat
- ▶ Java Champion
- ▶ LinkedIn Learning Presenter
- ▶ Author

Kubernetes

Kubernetes

" A system for ... " " ... automating deployment ... "
" ... scaling ... " " ... management ... "
" ... of containerized applications ... "

Kubernetes

“ A system for ...” “ ... automating deployment ...”

“ ... scaling ...” “ ... management ...”

“ ... of containerized applications ...”

“ It’s like a Linux kernel ... but for distributed
systems”

“ ... automating deployment ... ”

Container
scheduling

Automated
rollout/rollback

Self healing

Batch
execution

“ ... scaling ... ”

Horizontal
scaling

Load balancing

“ ... management ... ”

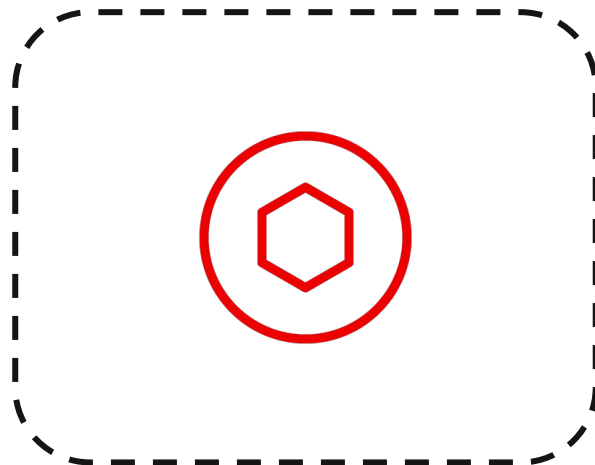
Service
discovery

Storage
orchestration

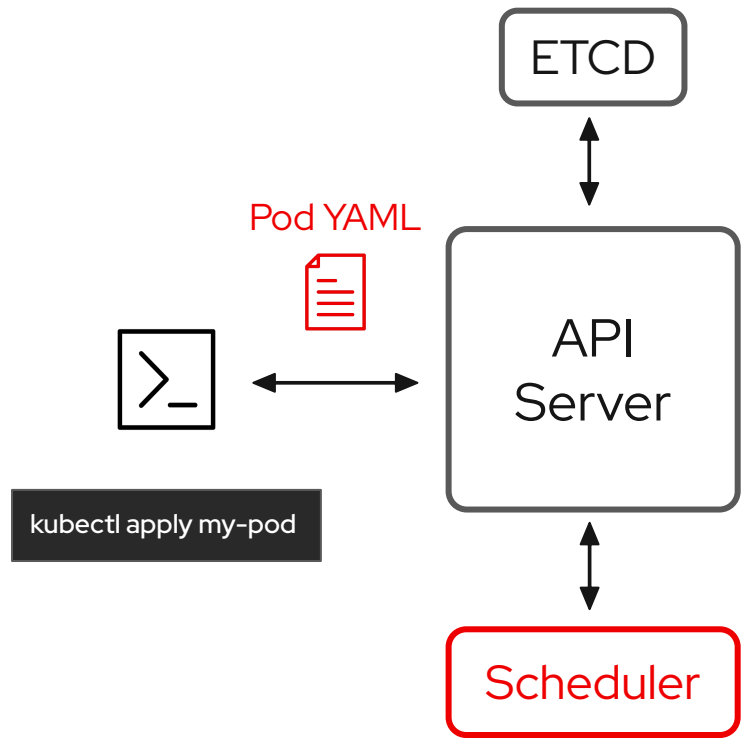
Secret &
configuration
management

It's declarative!

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.14.2
    ports:
    - containerPort: 80
```



Kubernetes
cluster





CONTAINERS ... PODS ...

...PODS EVERYWHERE



How does Kubernetes handle scaling, rollout, batch execution and so on?

```
apiVersion: v1
kind: Pod
# ...
```

```
apiVersion: apps/v1
kind: ReplicaSet
```

```
#
```

```
apiVersion: apps/v1
kind: Deployment
```

```
#
```

```
apiVersion: apps/v1
kind: StatefulSet
```

```
#
```

```
apiVersion: v1
kind: ConfigMap
```

```
# ...
```

⋮

```
apiVersion: v1
kind: Secret
# ...
```

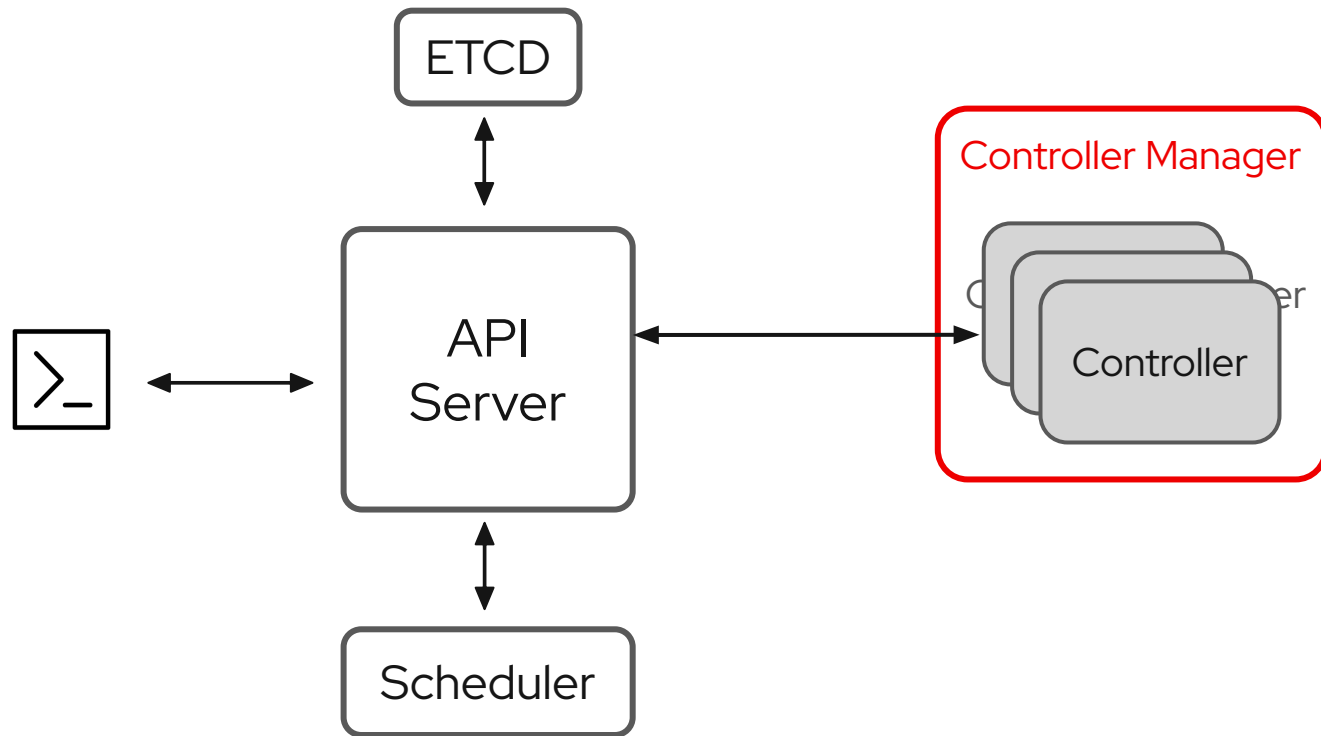


How does it work?

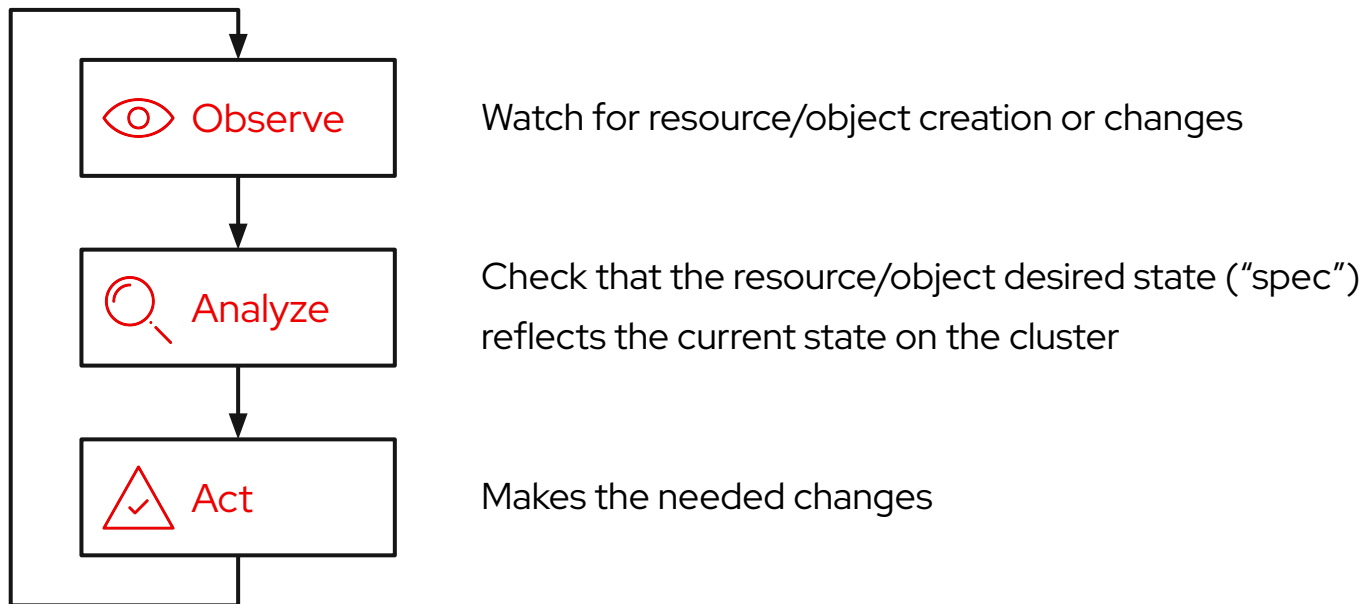
Let's use a controller!!!

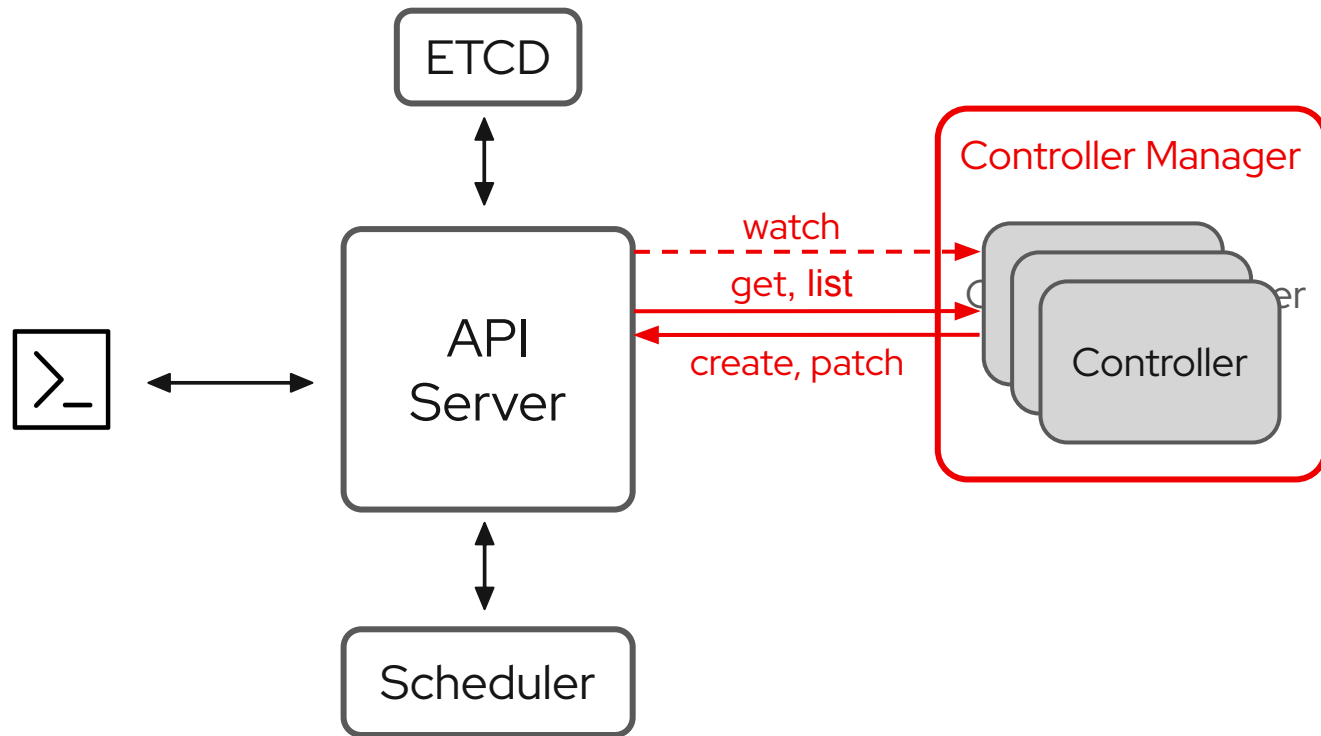


.... But not this one ;-)



Reconcile Loop







ReplicaSet
resource created

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset
spec:
  replicas: 2
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
      - name: my-application
        image: quay.io/devoxxuk/my-application:latest
```




2 replicas, spec

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset
spec:
  replicas: 2
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-application
          image: quay.io/devoxxuk/my-application:latest
```



Act

Create new pods

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset
spec:
  replicas: 2
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
      - name: my-application
        image: quay.io/devoxxuk/my-application:latest
```



```
apiVersion: v1
kind: Pod
metadata:
  name: my-replicaset-bf5zv
  labels:
    app: my-app
spec:
  containers:
  - name: my-application
    image: quay.io/devoxxuk/my-application:latest
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-replicaset-1tf5a
  labels:
    app: my-app
spec:
  # ...
```



What happens if the spec changes?



ReplicaSet resource updated

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset
spec:
  replicas: 2
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-application
          image: quay.io/devoxxuk/my-application:latest
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-replicaset-bf5zv
  labels:
    app: my-app
spec:
  # ...
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-replicaset-1tf5a
  labels:
    app: my-app
spec:
  # ...
```



ReplicaSet resource updated

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-application
          image: quay.io/devoxxuk/my-application:latest
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-replicaset-bf5zv
  labels:
    app: my-app
spec:
  # ...
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-replicaset-1tf5a
  labels:
    app: my-app
spec:
  # ...
```



Search for pods
matching
resources

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-application
          image: quay.io/devoxxuk/my-application:latest
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-replicaset-bf5zv
  labels:
    app: my-app
spec:
  # ...
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-replicaset-1tf5a
  labels:
    app: my-app
spec:
  # ...
```



Create new pod

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-application
          image: quay.io/devoxxuk/my-application:latest
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-replicaset-bf5zv
  labels:
    app: my-app
spec:
  # ...
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-replicaset-1tf5a
  labels:
    app: my-app
spec:
  # ...
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-replicaset-gb65f
  labels:
    app: my-app
spec:
  # ...
```



What happens if there are
resources already created?


```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod-1
  labels:
    app: my-app
spec:
  # ...
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod-2
  labels:
    app: my-app
spec:
  # ...
```



ReplicaSet resource created

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-application
          image: quay.io/devoxxuk/my-application:latest
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod-1
  labels:
    app: my-app
spec:
  # ...
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod-2
  labels:
    app: my-app
spec:
  # ...
```



Search for pods
matching resources

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-application
          image: quay.io/devoxxuk/my-application:latest
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod-1
  labels:
    app: my-app
spec:
  # ...
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod-2
  labels:
    app: my-app
spec:
  # ...
```



Take ownership
of existing pods

Create new pod

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-application
          image: quay.io/devoxxuk/my-application:latest
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod-1
  labels:
    app: my-app
spec:
  # ...
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod-2
  labels:
    app: my-app
spec:
  # ...
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-replicaset-gb65f
  labels:
    app: my-app
spec:
  # ...
```



Is it really that simple?

```
apiVersion:  
apps/v1  
kind: ReplicaSet  
# ...  
spec:  
  replicas: 2  
# ...
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: nginx  
spec:  
  # ...
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: nginx  
spec:  
  # ...
```

```
apiVersion: apps/v1
kind: Deployment
# ...
spec:
  replicas: 2
  strategy:
    type: RollingUpdate
# ...
```

```
apiVersion: apps/v1
kind: Deployment
# ...
spec:
  replicas: 2
  strategy:
    type: RollingUpdate
# ...
```

v1

```
apiVersion: apps/v1
kind: ReplicaSet
# ...
spec:
  replicas: 2
# ...
```

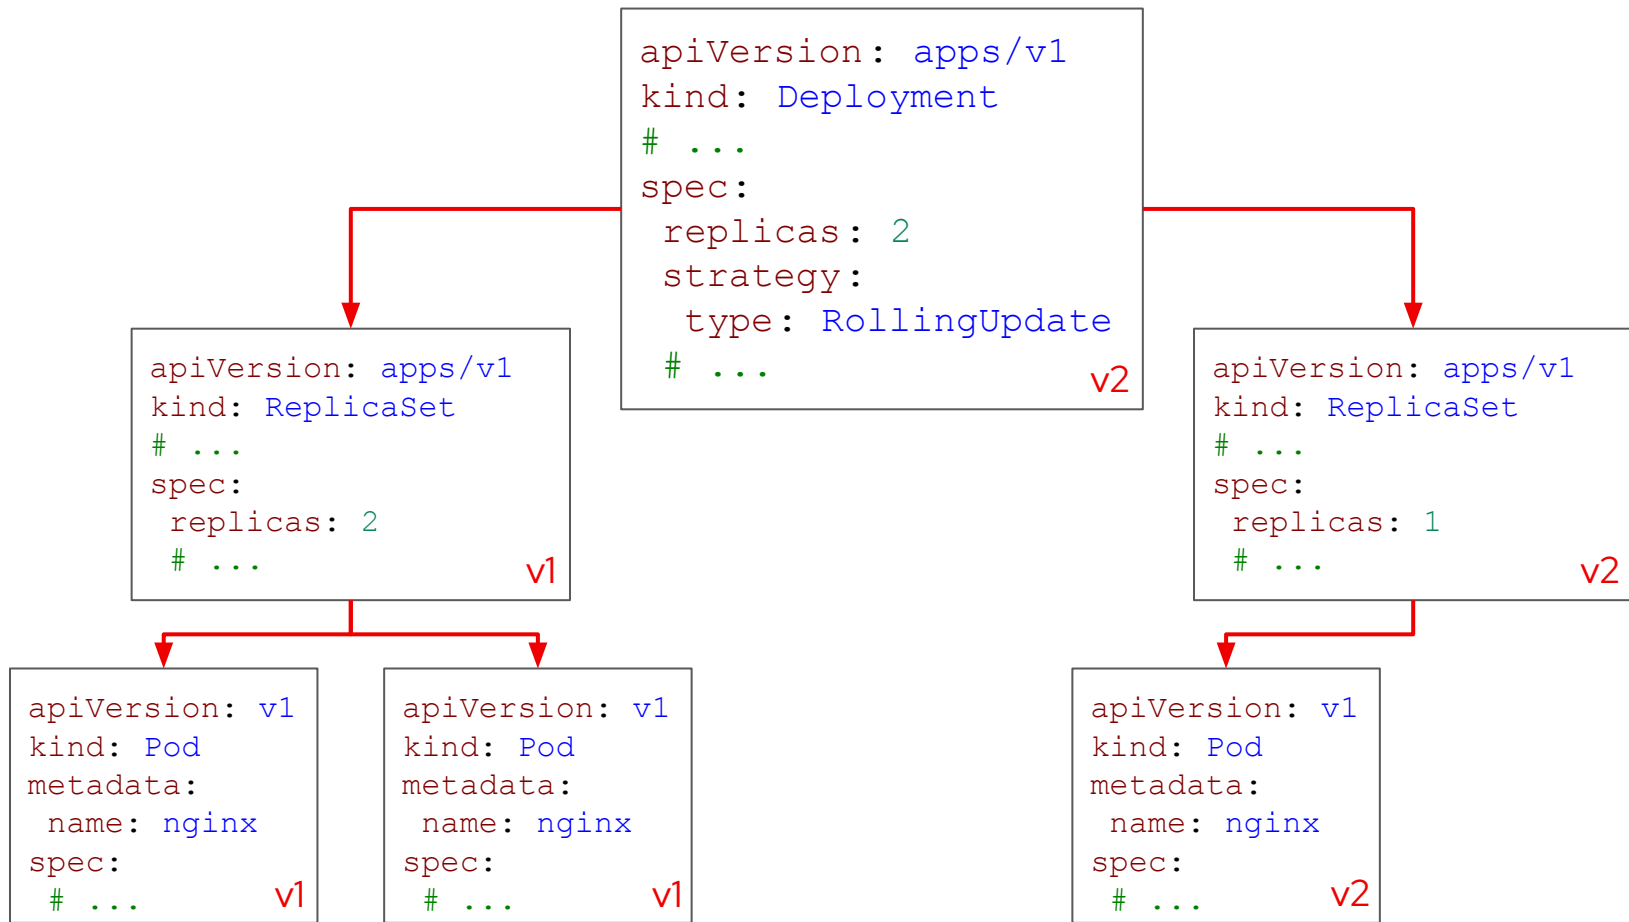
v1

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
# ...
```

v1

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
# ...
```

v1



```
apiVersion: apps/v1
kind: Deployment
# ...
spec:
  replicas: 2
  strategy:
    type: RollingUpdate
# ...
```

v2

```
apiVersion: apps/v1
kind: ReplicaSet
# ...
spec:
  replicas: 1
# ...
```

v1

```
apiVersion: apps/v1
kind: ReplicaSet
# ...
spec:
  replicas: 1
# ...
```

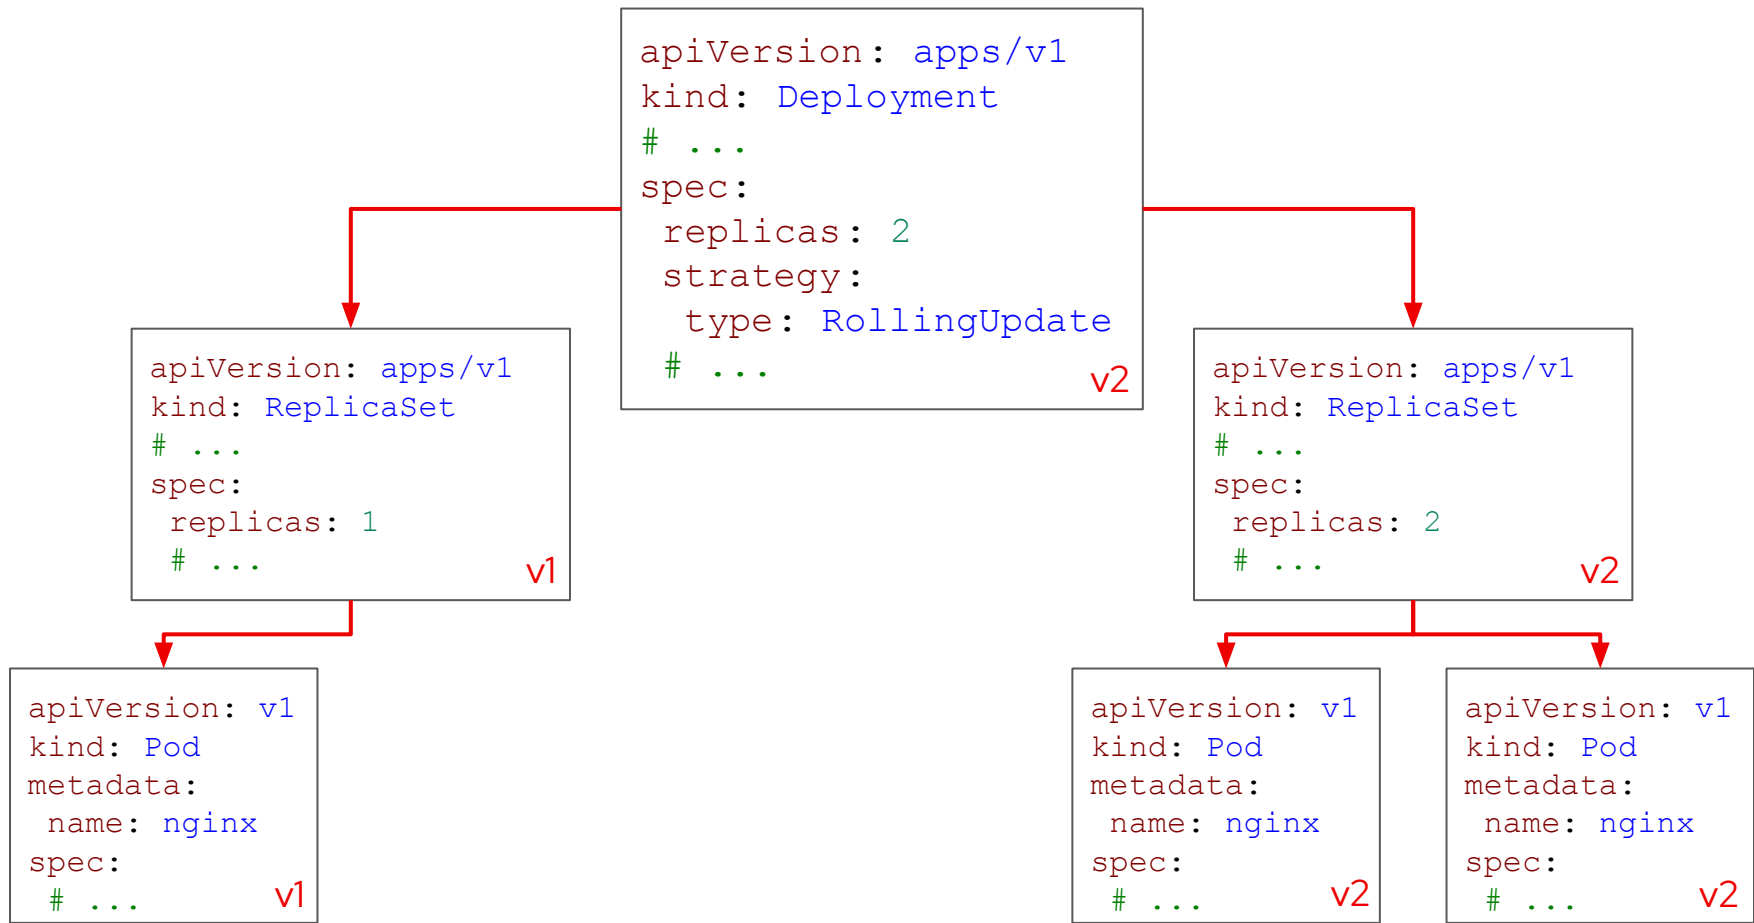
v2

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
# ...
```

v1

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
# ...
```

v2



```
apiVersion: apps/v1
kind: Deployment
# ...
spec:
  replicas: 2
  strategy:
    type: RollingUpdate
# ...
```

v2

```
apiVersion: apps/v1
kind: ReplicaSet
# ...
spec:
  replicas: 2
# ...
```

v2

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
# ...
```

v2

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
# ...
```

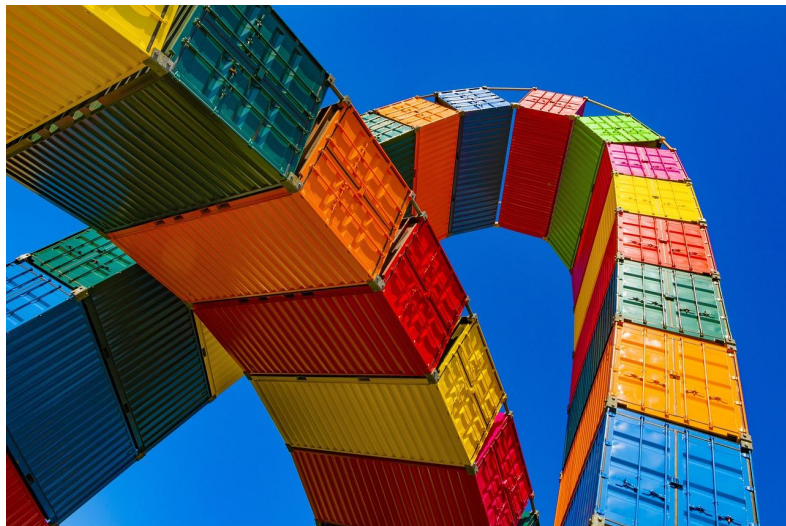
v2

What have we learnt?

- ▶ User manages highest level YAML
- ▶ Controllers perform reconcile loop
 - 👁 Observe
 - 🔍 Analyze
 - ⚙ Act



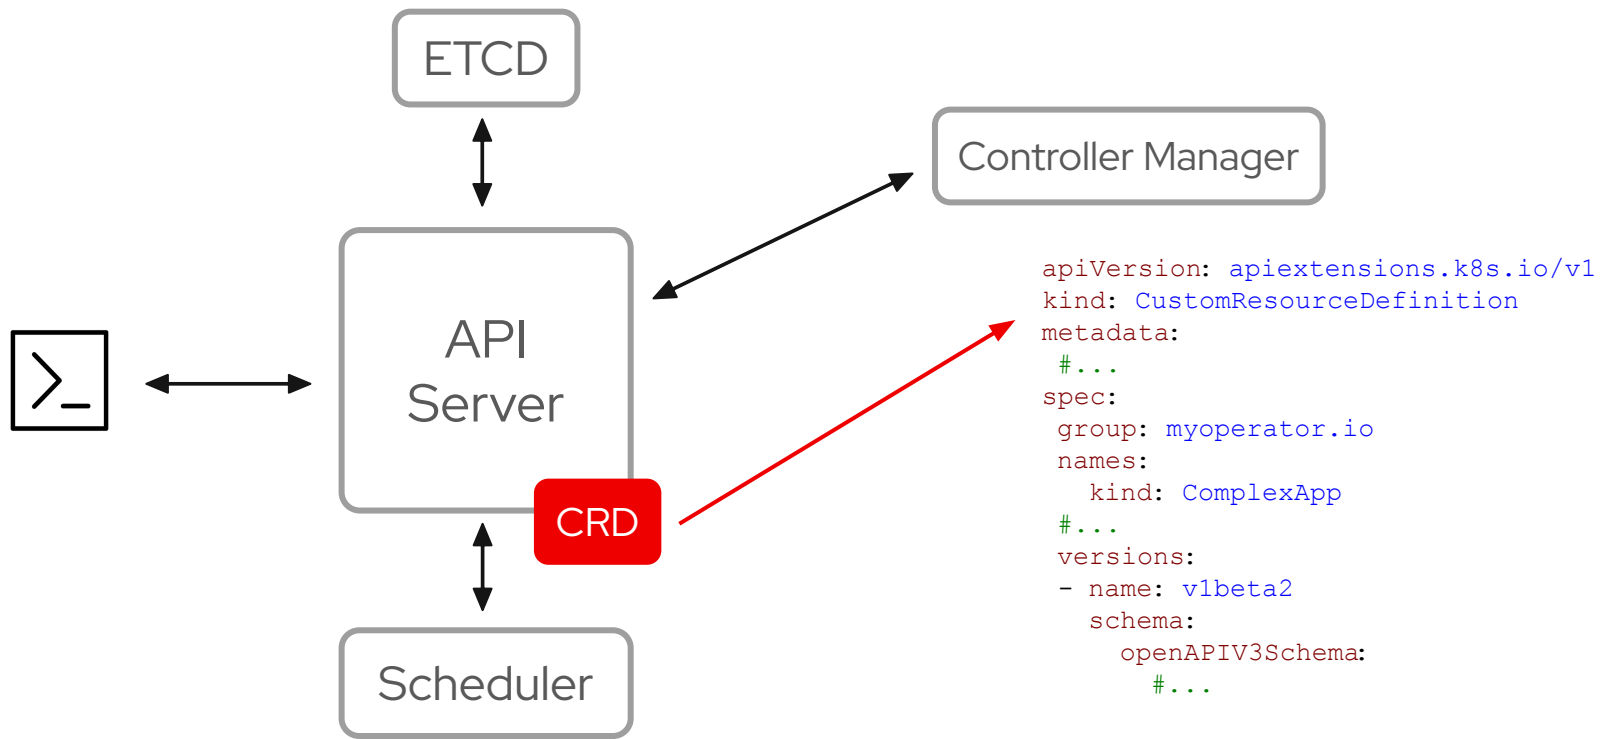
How to automate operating complex
“containerized” applications?





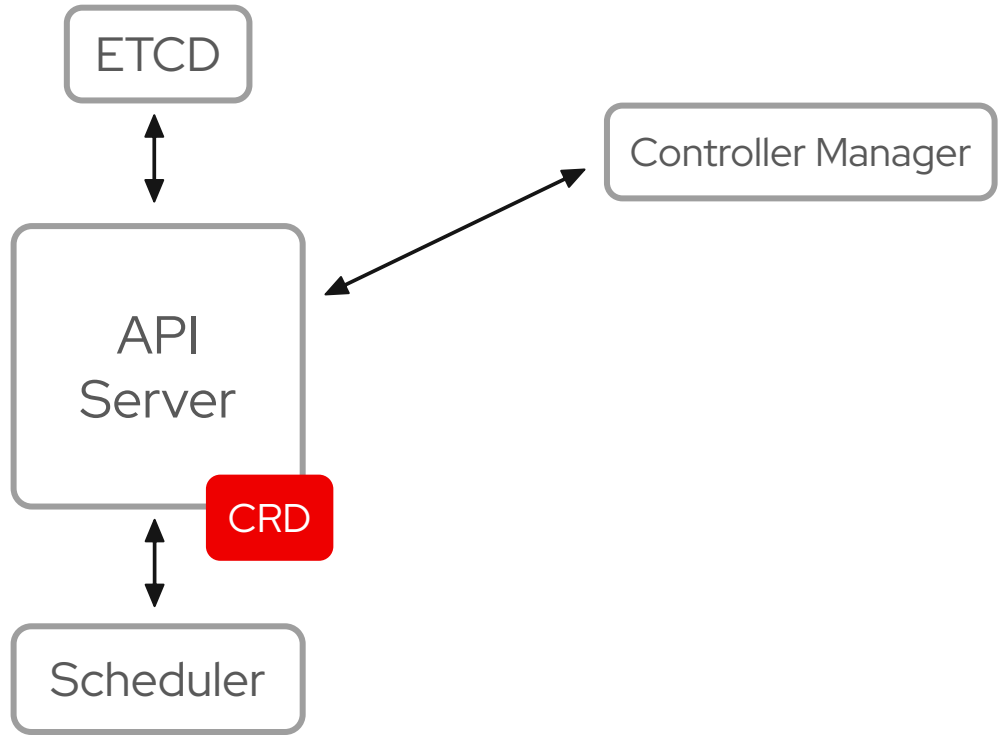
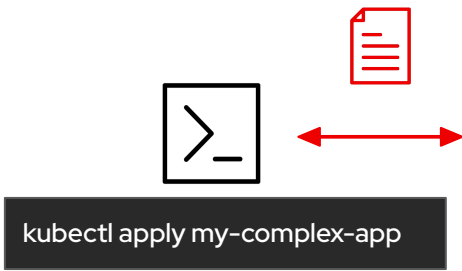
The Operator pattern!

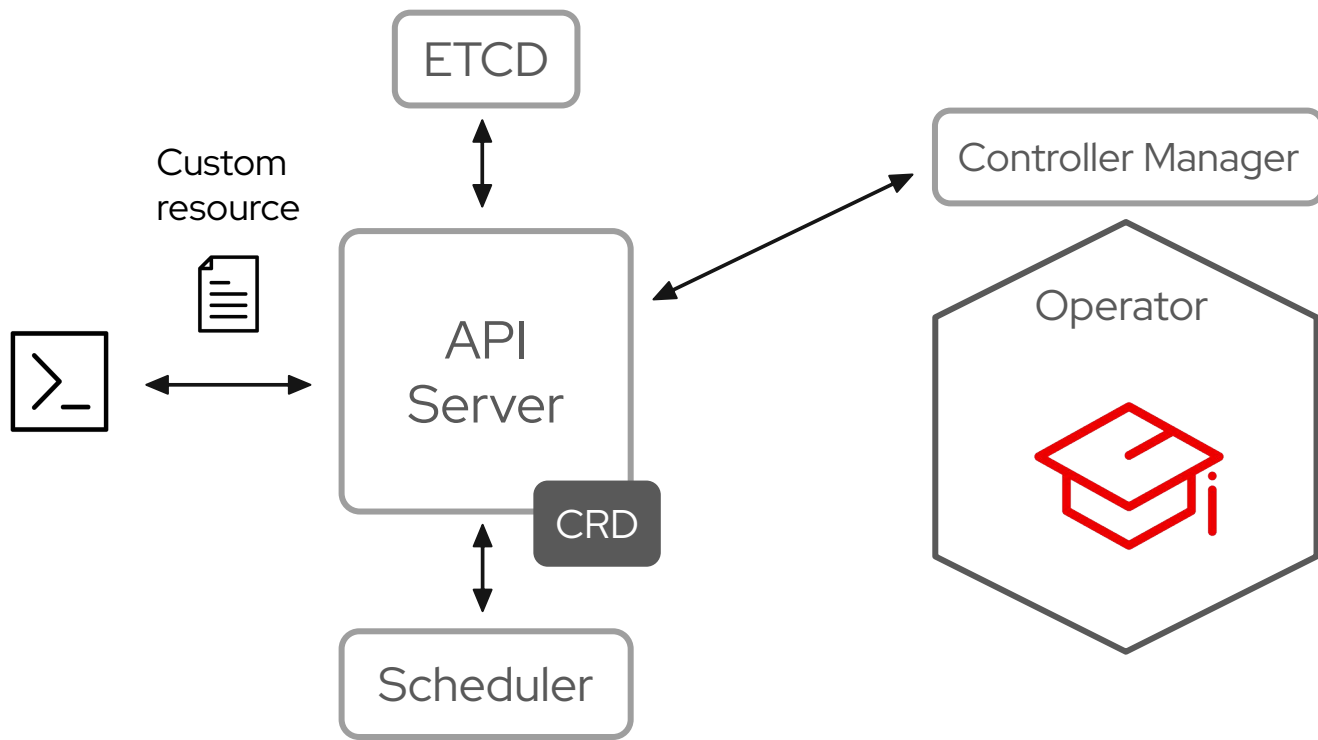


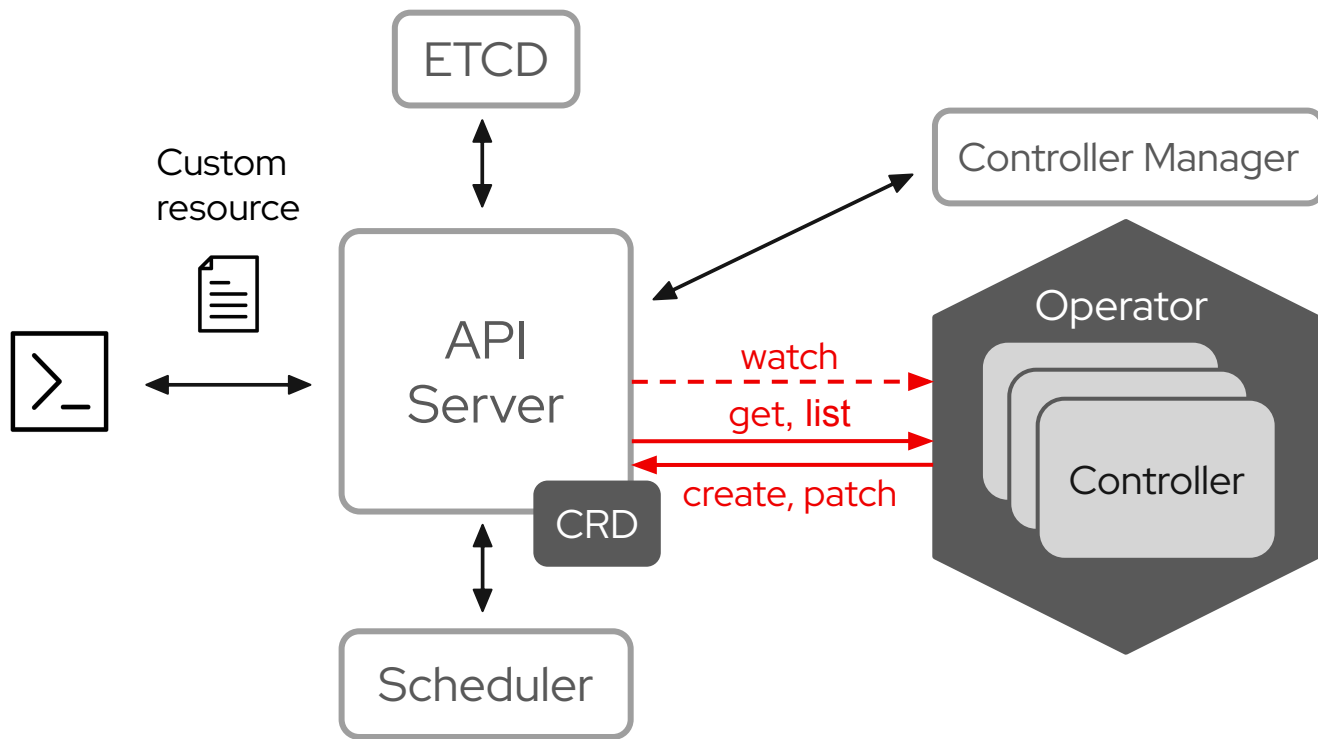


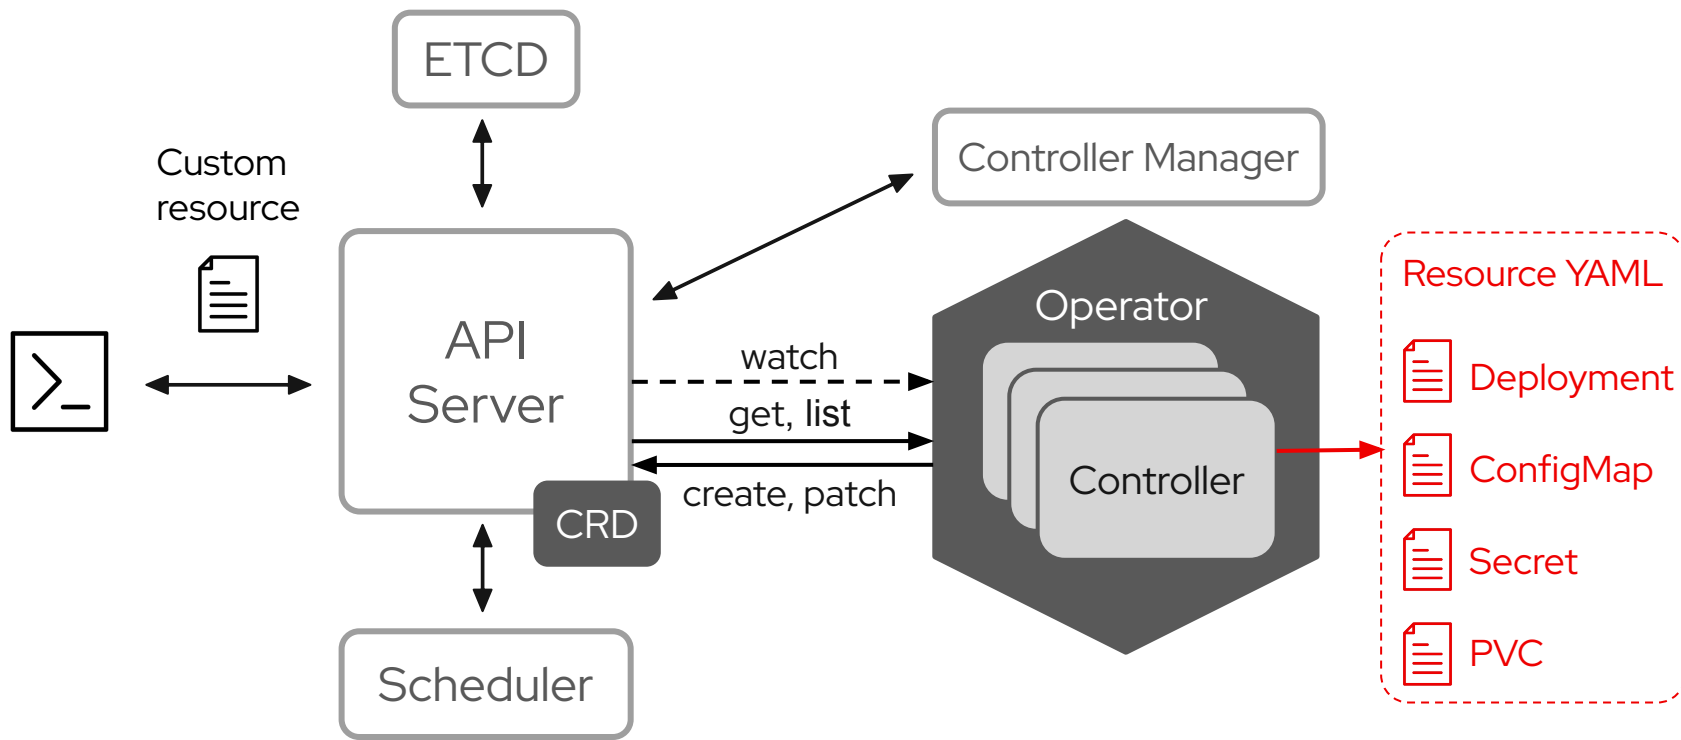

```
apiVersion: myoperator.io/v1
kind: ComplexApp
metadata:
  name: my-complex-app
spec:
  # ...
status:
  # ...
```

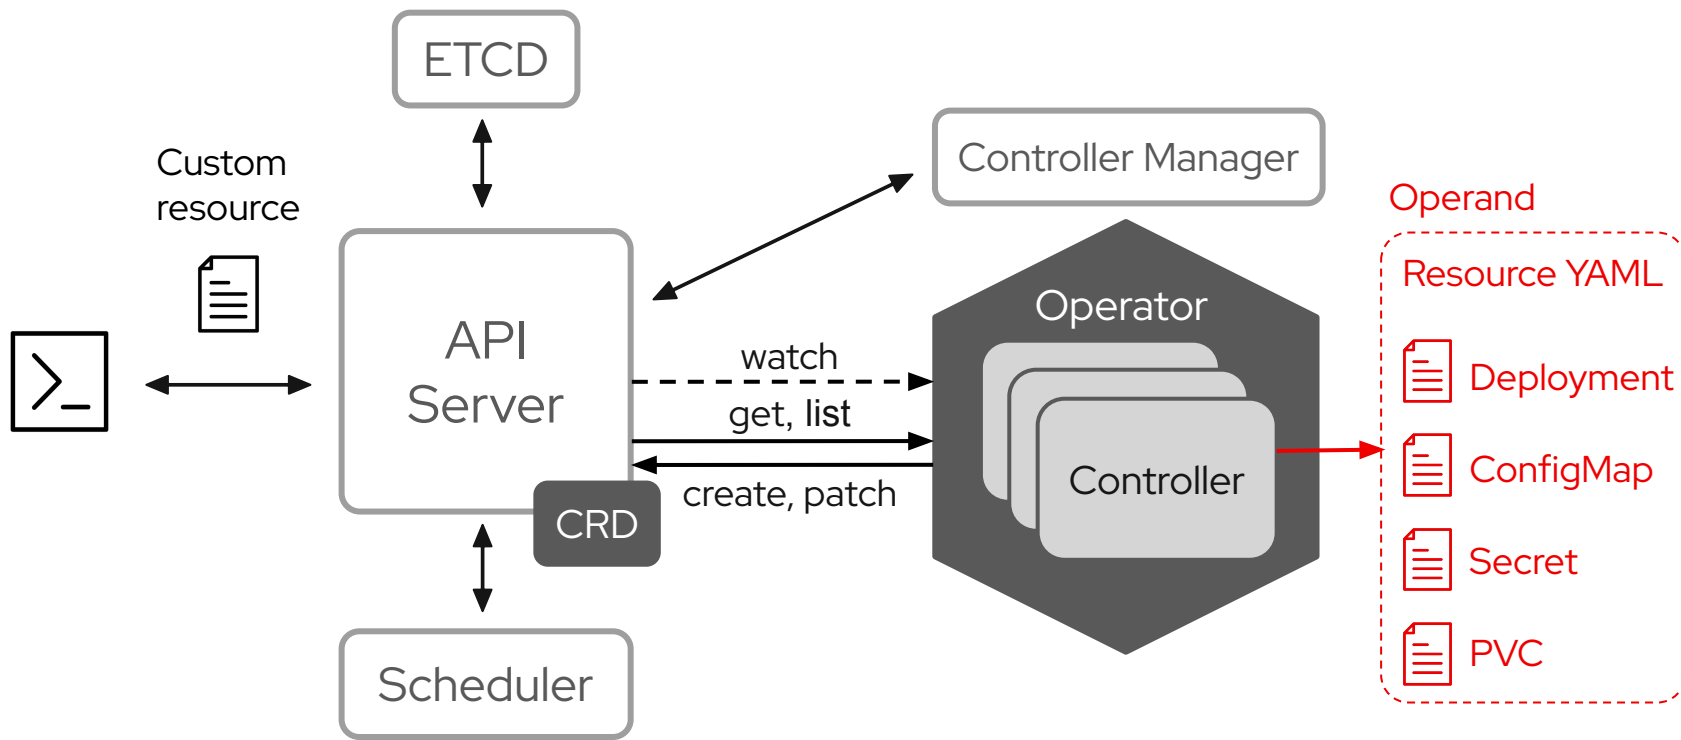
Custom resource

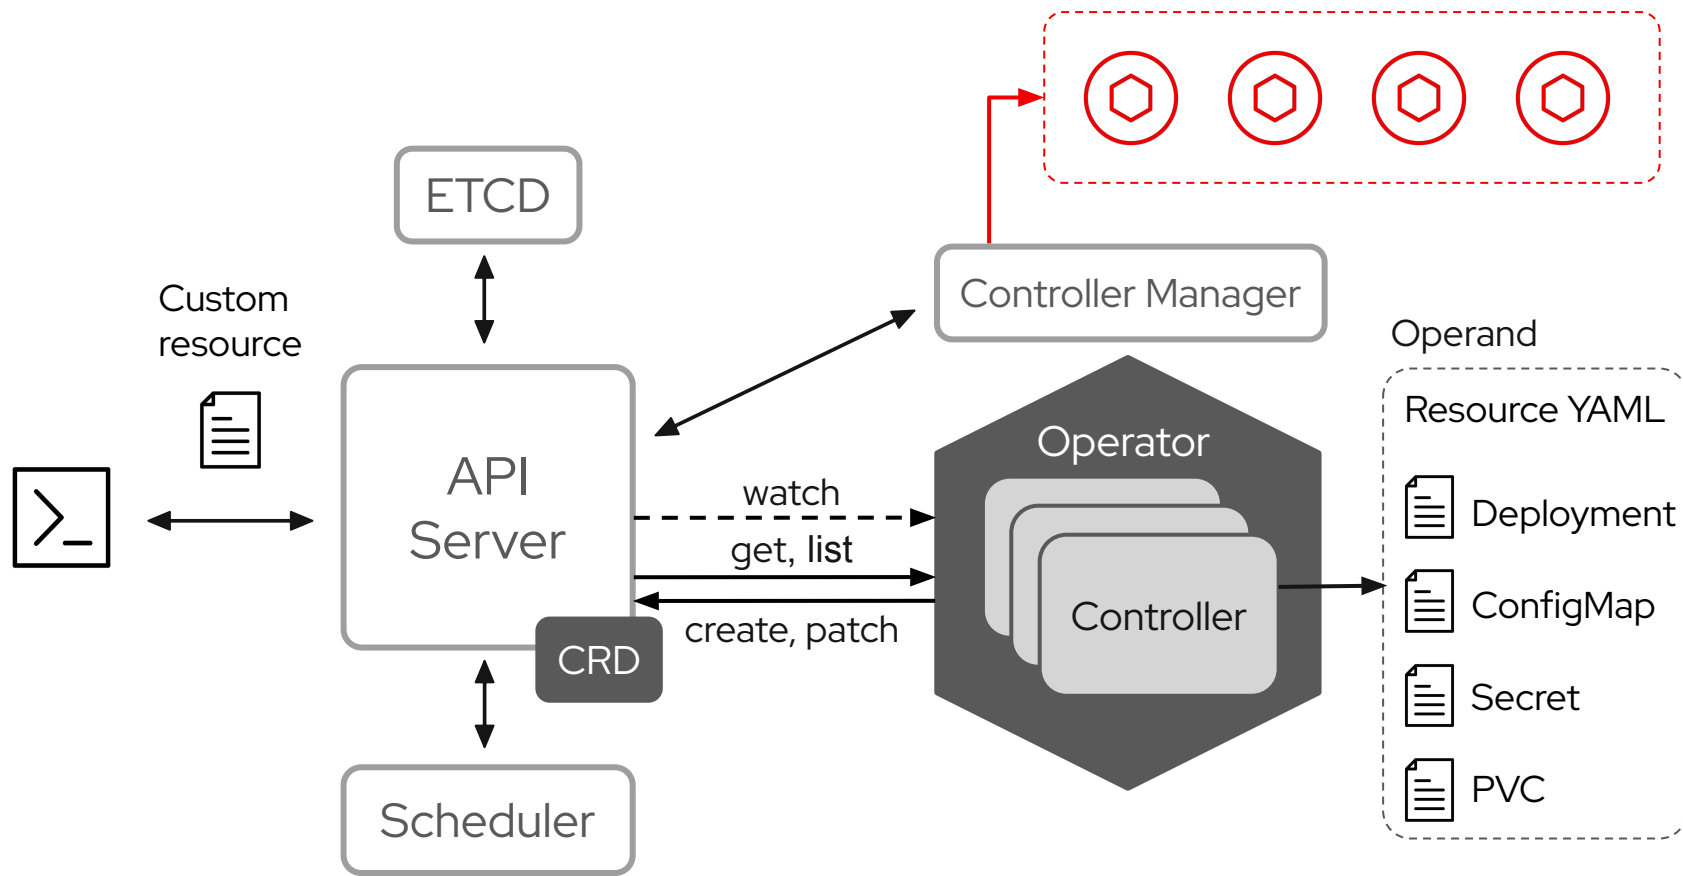


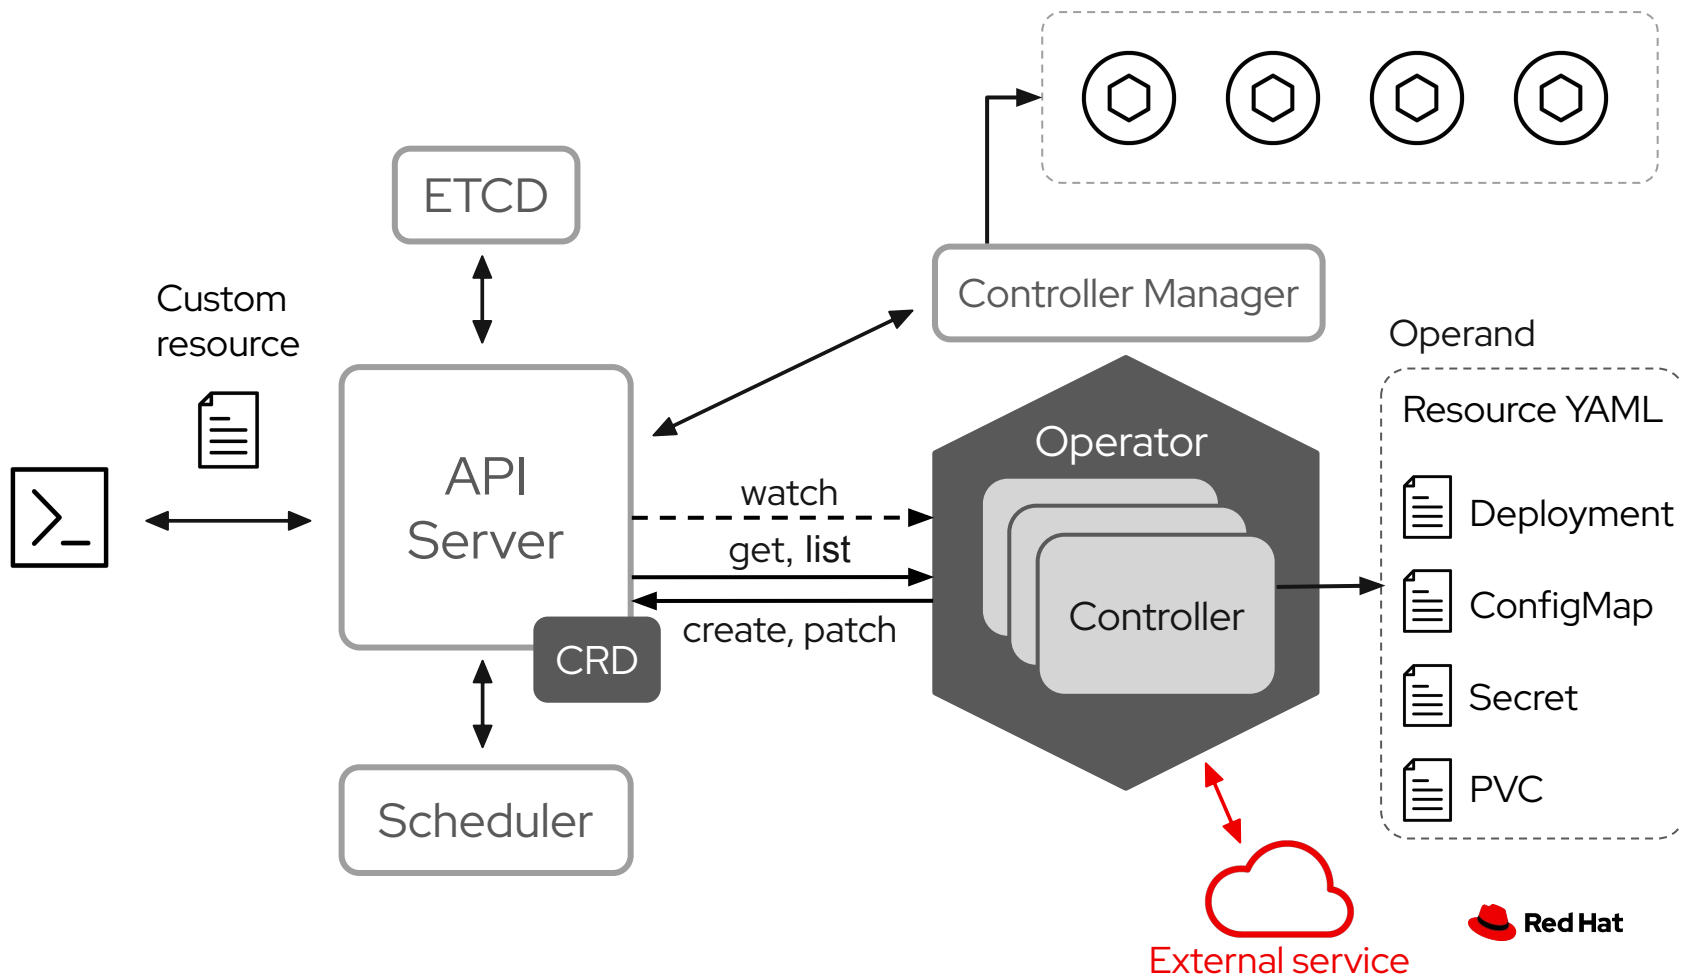














What if I want to deploy Apache Kafka on Kubernetes?

Apache Kafka

"A publish/subscribe messaging system ..."

"A data & event streaming platform ..."

Apache Kafka

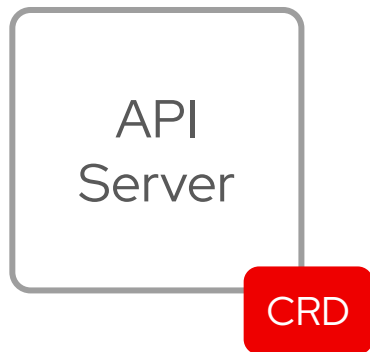
"A publish/subscribe messaging system ..."

"A data & event streaming platform ..."

with its own complexity to ...

- ▶ deploy and run
- ▶ apply configuration changes
- ▶ upgrade to a new release
- ▶ secure

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name: kafkas.kafka.strimzi.io
spec:
  group: kafka.strimzi.io
  names:
    kind: Kafka
    listKind: KafkaList
  #...
  versions:
  - name: v1beta2
    schema:
      openAPIV3Schema:
        type: object
        properties:
          spec:
            # spec definition for the custom resource
            kafka:
              #...
          status:
            # status definition reported back
            # in the custom resource
```



```
GET /apis/kafka.strimzi.io/v1beta2/kafkas/
kubectl get kafka
```

```

apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name: kafkas.kafka.strimzi.io
spec:
  group: kafka.strimzi.io
  names:
    kind: Kafka
    listKind: KafkaList
    #...
  versions:
    - name: v1beta2
      schema:

```

```

        openAPIV3Schema:
          type: object
          properties:
            spec:
              # spec definition for the custom resource
              kafka:
                #...

```

```

            status:
              # status definition reported back
              # in the custom resource

```

```

apiVersion: kafka.strimzi.io/v1beta2
kind: Kafka
metadata:
  name: my-cluster
spec:
  kafka:
    version: 3.7.0
    replicas: 3
    listeners:
      - name: plain
        port: 9092
        type: internal
        tls: false
    # ...
    config:
      default.replication.factor: 3
      min.insync.replicas: 2
    # ...
    storage:
      type: ephemeral
    # ...
  status:
    clusterId: UUq-xVw5TdW8GVBXappe4g
    conditions:
      - lastTransitionTime: "2024-04-09T12:51:02"
        status: "True"
        type: Ready
    kafkaMetadataState: Kraft
    kafkaMetadataVersion: 3.7-IV4
    # ...

```

Strimzi

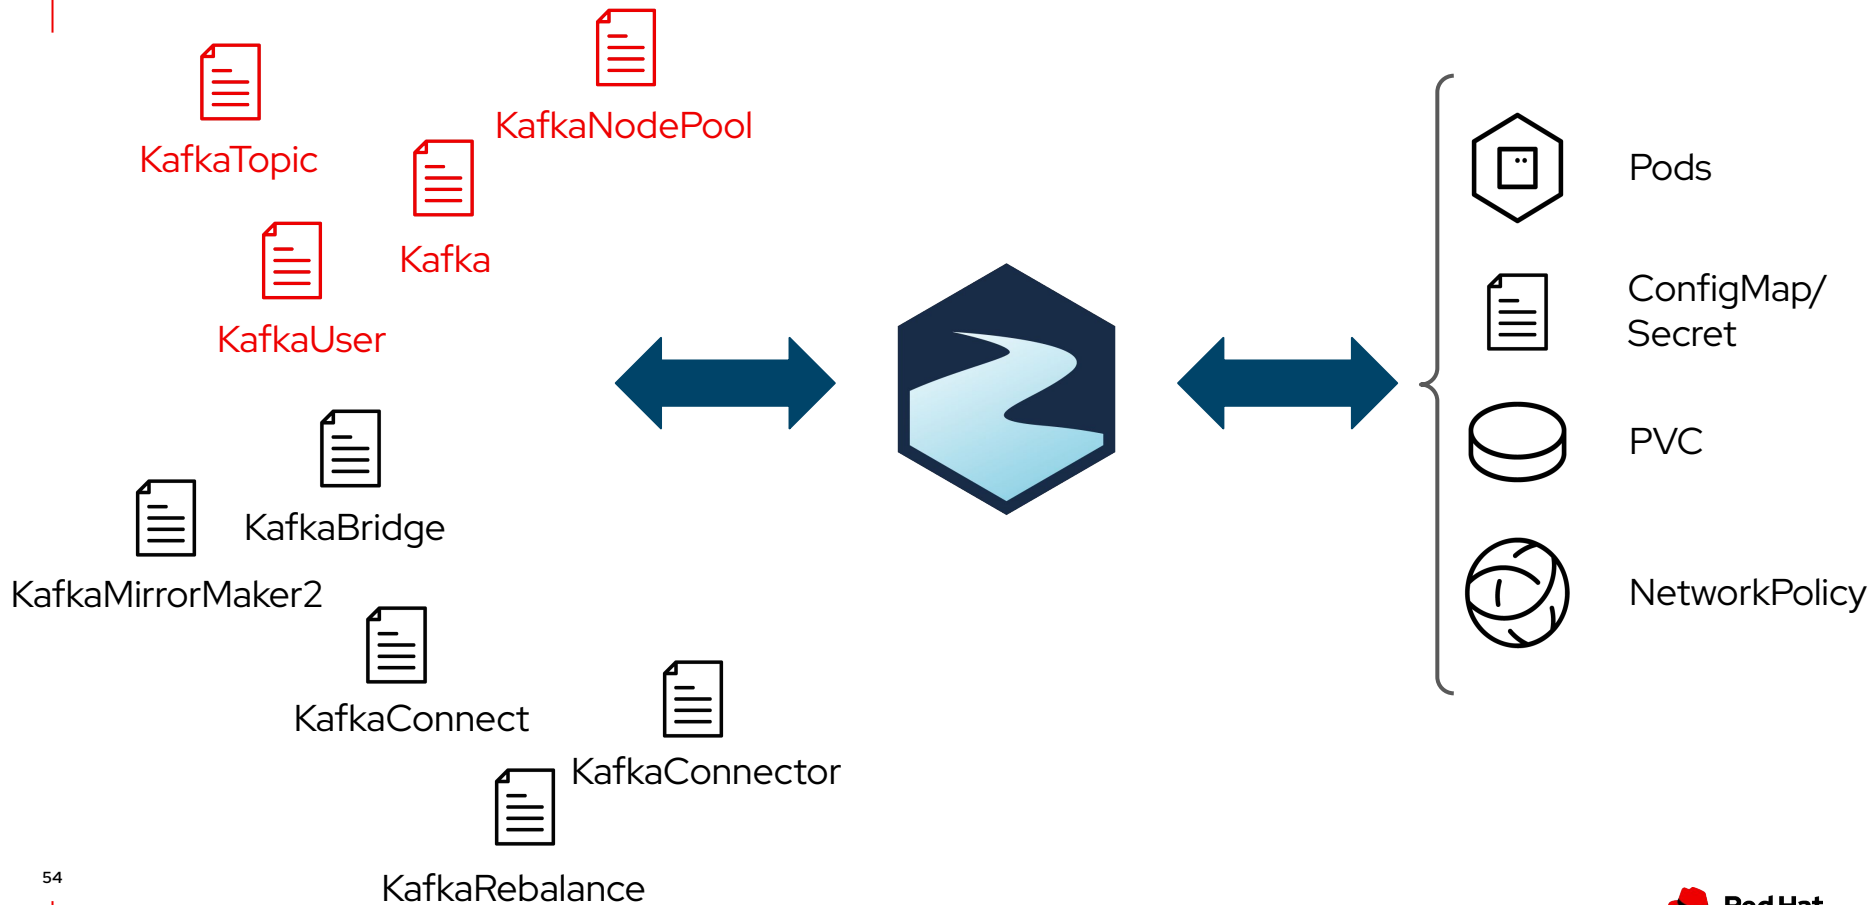
Open Source project (Apache License 2.0)

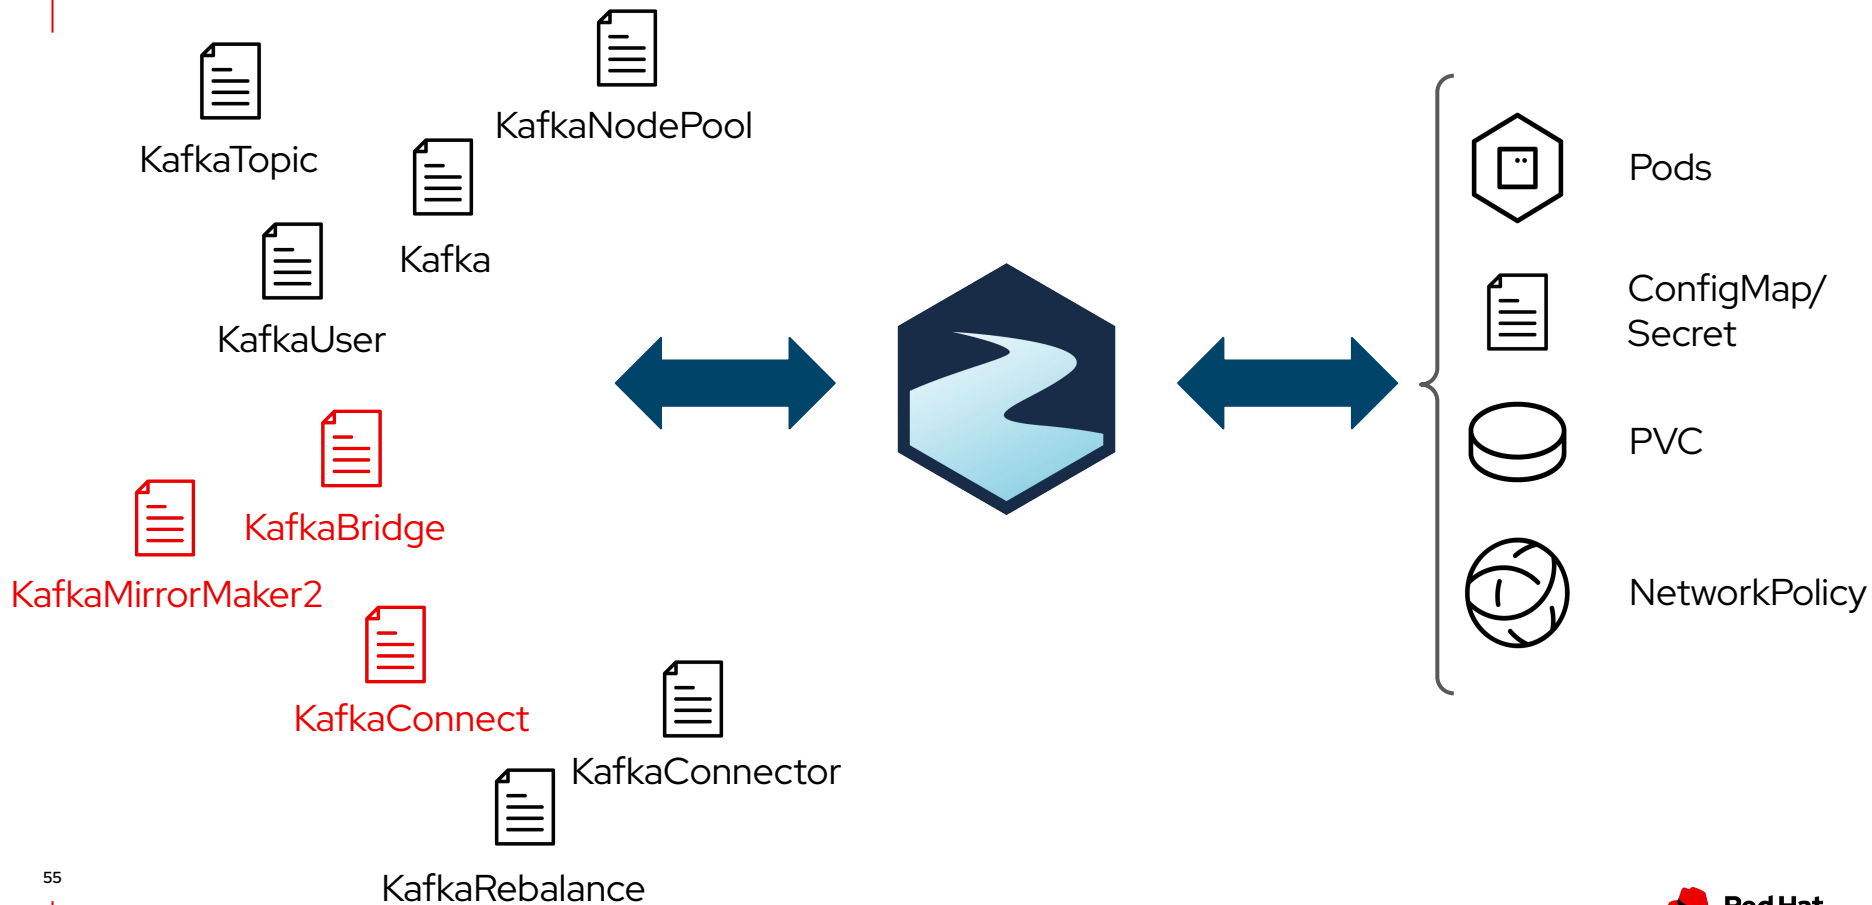
Focuses on Apache Kafka on Kubernetes

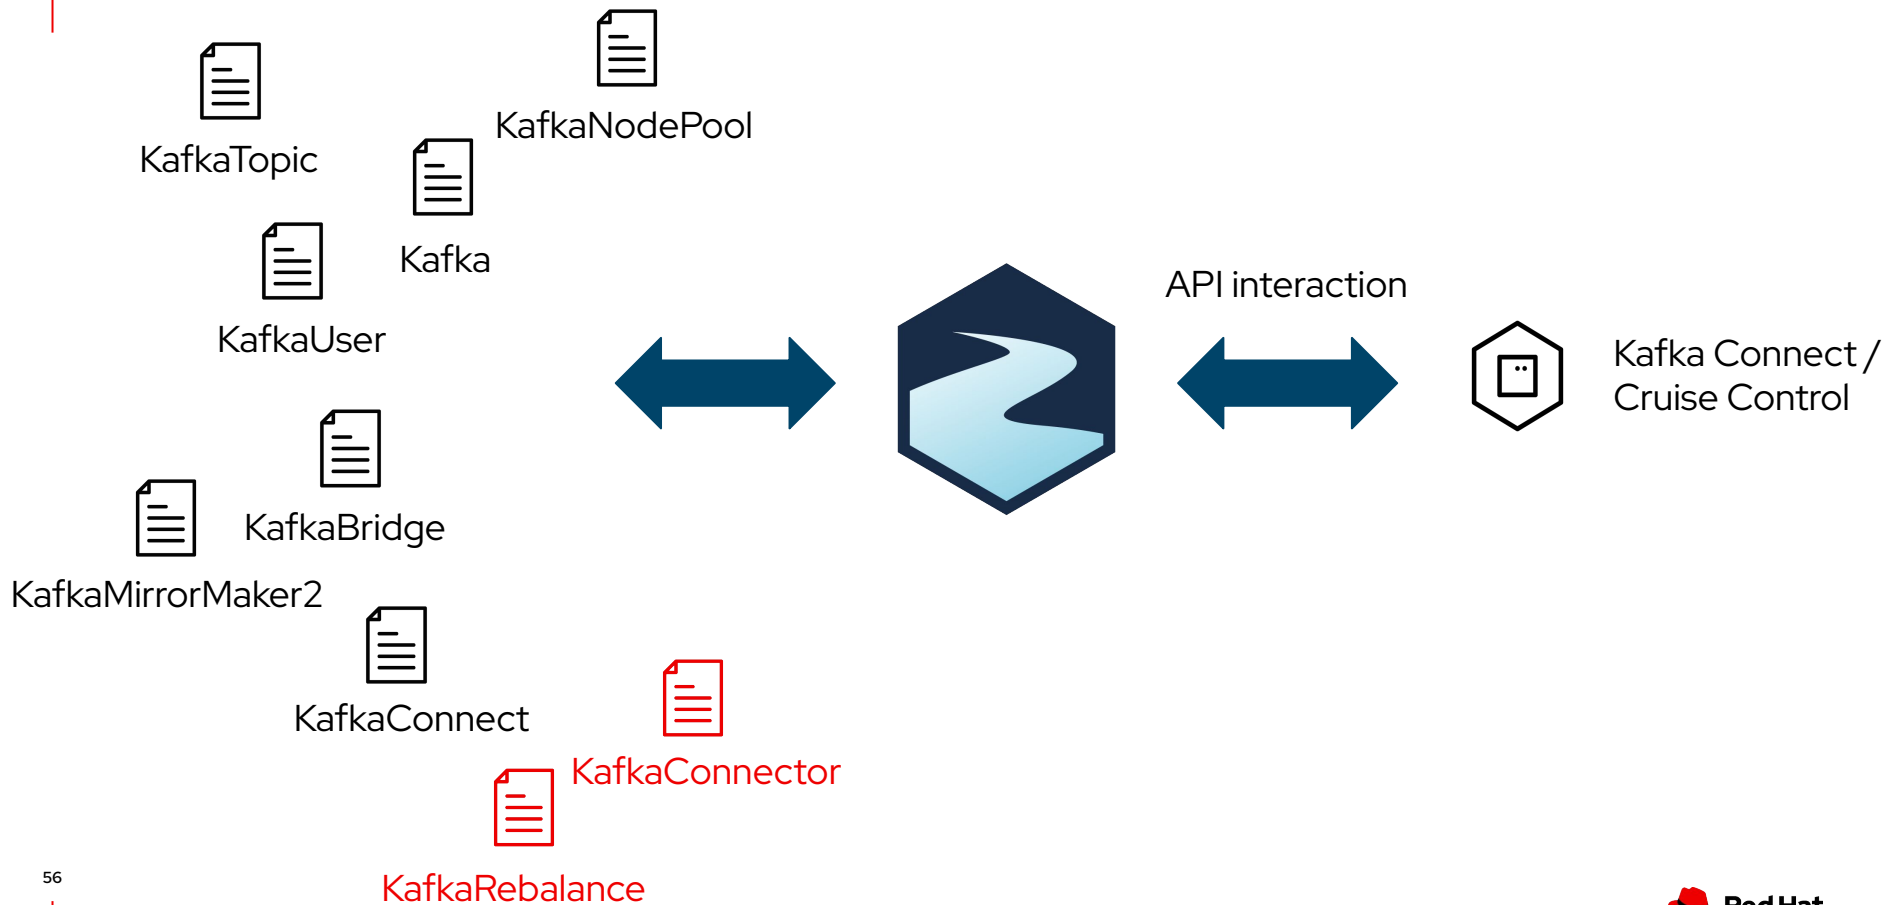
CNCF Incubating Project



strimzi.io







Demo time!



What are the alternatives?



Helm

Open Source project (Apache License 2.0)

Package manager for Kubernetes

CNCF Graduated Project



Helm vs Operator

Helm

- ▶ Relies on Kubernetes built-in resources

Operator

- ▶ Extends the Kubernetes API with CRDs

Helm vs Operator

Helm

- ▶ Relies on Kubernetes built-in resources
- ▶ Many YAMLs with customization via templating

Operator

- ▶ Extends the Kubernetes API with CRDs
- ▶ One (or a few) “custom resource” YAMLs

Helm vs Operator

Helm

- ▶ Relies on Kubernetes built-in resources
- ▶ Many YAMLs with customization via templating
- ▶ Ideal for day-1 operation (deploying)

Operator

- ▶ Extends the Kubernetes API with CRDs
- ▶ One (or a few) “custom resource” YAMLs
- ▶ Useful for day-1 and day-2 operations (upgrading, scaling)

Helm vs Operator

Helm

- ▶ Relies on Kubernetes built-in resources
- ▶ Many YAMLs with customization via templating
- ▶ Ideal for day-1 operation (deploying)

Operator

- ▶ Extends the Kubernetes API with CRDs
- ▶ One (or a few) “custom resource” YAMLs
- ▶ Useful for day-1 and day-2 operations (upgrading, scaling)
- ▶ Deployable via Helm charts!



You've convinced me!
How do I start?

Operator Framework



Open Source project (Apache License 2.0)

Toolkit to manage Kubernetes operators

- Build - Operator SDK

- Manage - OLM

- Discover - OperatorHub.io

CNCF Incubating Project

OperatorHub.io

OperatorHub.io

Search OperatorHub...

Contribute

Welcome to OperatorHub.io

OperatorHub.io is a new home for the Kubernetes community to share Operators. Find an existing Operator or list your own today.

CATEGORIES

352 ITEMS

VIEW SORT A-Z

AI/Machine Learning

Application Runtime

Big Data

Cloud Provider

Database


Developer Tools

Drivers and plugins

Integration & Delivery


Logging & Tracing

Modernization & Migration




Aerospike Kubernetes Operator
provided by Aerospike

The Aerospike Kubernetes Operator automates the




Airflow Helm Operator
provided by opdev

An experimental operator that installs Apache Airflow.




Aiven Operator
provided by aiven

Manage your https://aiven.io resources with Kubernetes.



Akka Cluster Operator
provided by Lightbend, Inc.

Run Akka Cluster applications on Kubernetes.



Altinity Operator for ClickHouse
provided by Altinity

ClickHouse Operator manages full lifecycle of ClickHouse



[JAVA OPERATOR SDK]

javaoperatorsdk.io

Java Operator SDK

Open Source project (Apache License 2.0)

Framework for writing Java based operators

Built on top of the Fabric8 Client

Thank you

Operator Pattern - <https://kubernetes.io/docs/concepts/extend-kubernetes/operator/>

Strimzi - strimzi.io

Apache Kafka at Red Hat - <https://developers.redhat.com/topics/kafka-kubernetes>



github.com/ppatierno

github.com/katheris