CLOUD NATIVE
COMPUTING FOUNDATION

PORTO
KCD
GROWING CLOUD NATIVE TOGETHER

Managing Kubernetes workloads: extend the platform with operators

Paolo Patierno,

Senior Principal Software Engineer @ Red Hat

Who am I?

@ppatierno

- Senior Principal Software Engineer @ Red Hat
  - Messaging & data streaming
- CNCF Ambassador
- Strimzi maintainer
- Running, swimming, Formula 1 & MotoGP addicted

We all know about Kubernetes … right? :-)

Kubernetes is … declarative!

```
apiVersion: v1
kind: Pod
metadata:
 name: nginx
spec:
 containers:
 - name: nginx
   image: nginx:1.14.2
   ports:
   - containerPort: 80
```

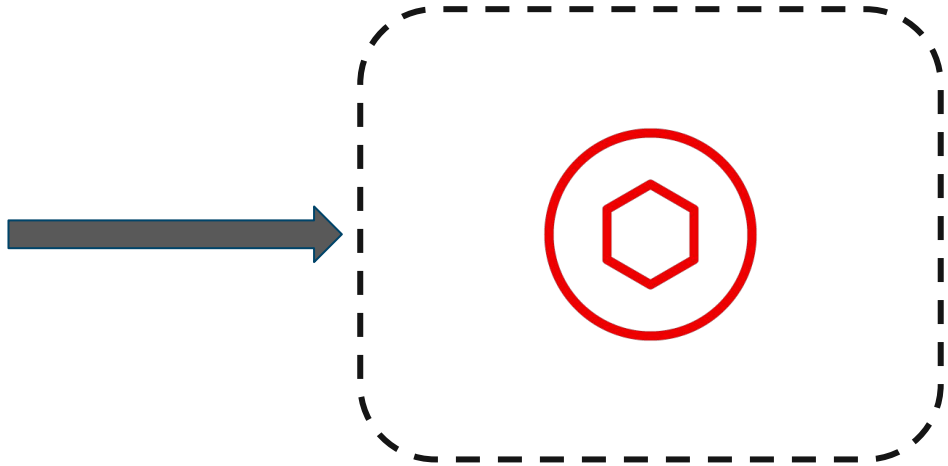

Kubernetes
cluster

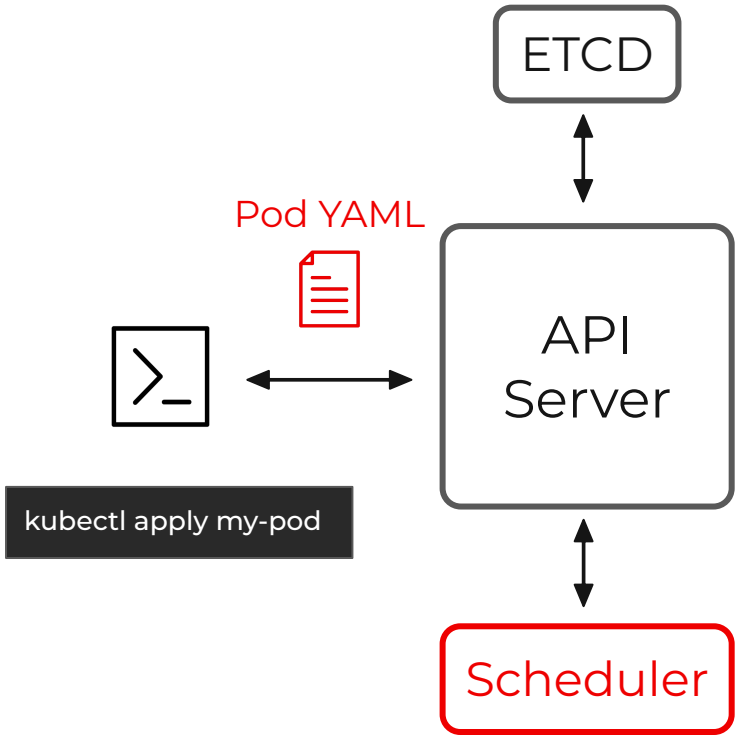

How does Kubernetes handle scaling, rollout, batch execution and so on?

```yaml
apiVersion: v1
kind: Pod
# ...
```

```yaml
apiVersion: apps/v1
kind: ReplicaSet
# ...
```

```yaml
apiVersion: apps/v1
kind: Deployment
# ...
```

```yaml
apiVersion: apps/v1
kind: StatefulSet
# ...
```

⋮

```yaml
apiVersion: batch/v1
kind: Job
# ...
```


How does it work?
Let’s use a controller!!!

…. But not this one ;-)

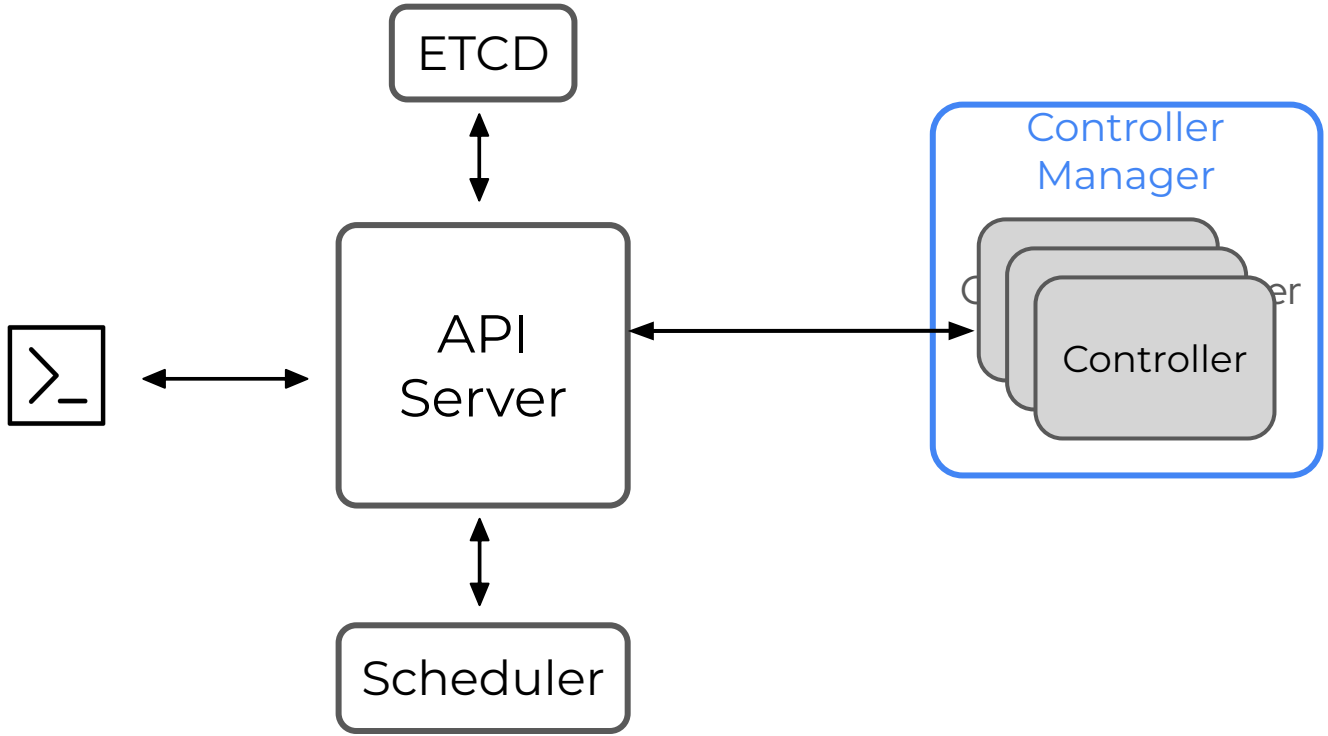

Reconcile Loop

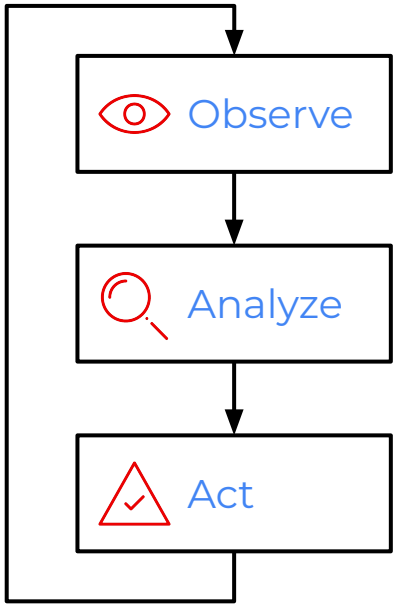

Watch for resource/object creation or changes

Check that the resource/object desired state (“spec”) reflects the current state on the cluster

Makes the needed changes

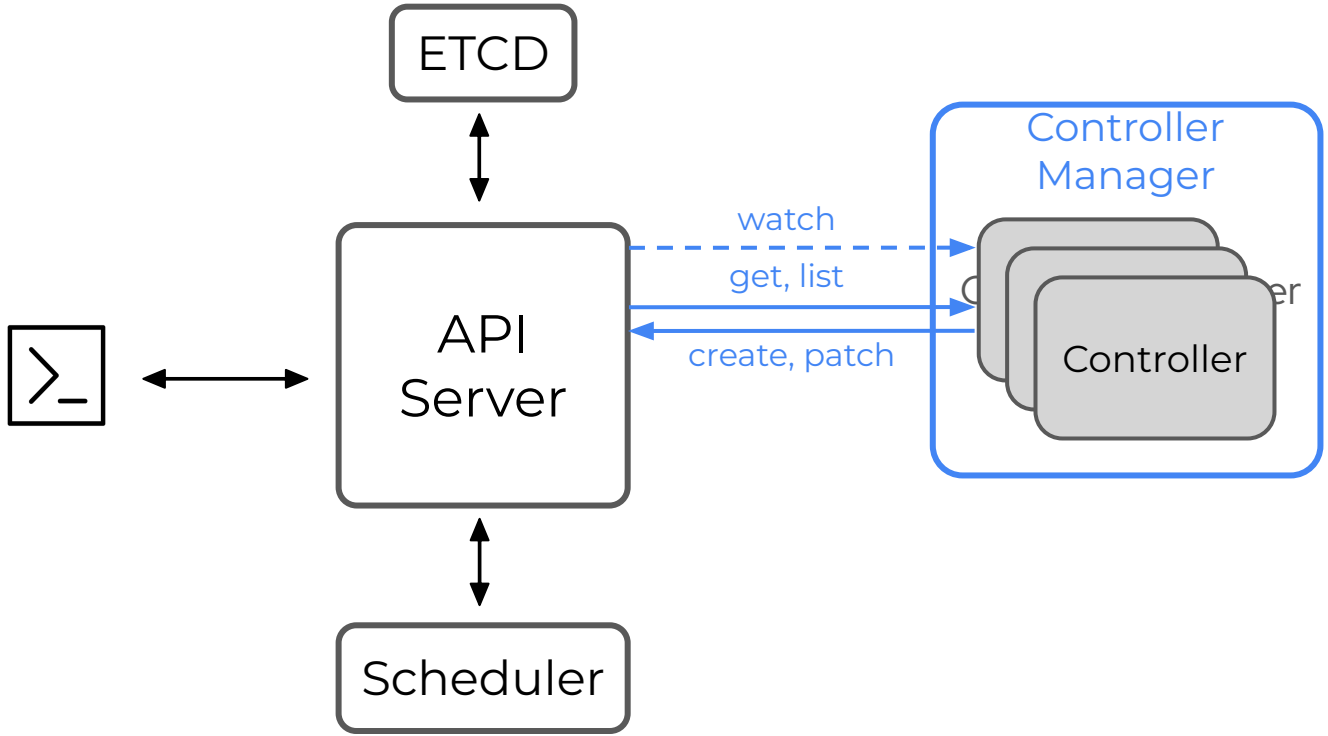

ReplicaSet
resource created

```yaml
apiVersion: apps/v1
kind: ReplicaSet
metadata:
 name: my-replicaset
spec:
 replicas: 2
 selector:
   matchLabels:
     app: my-app
 template:
   metadata:
     labels:
       app: my-app
   spec:
     containers:
     - name: my-application
       image: quay.io/devoxxuk/my-application:latest
```


2 replicas, spec

```yaml
apiVersion: apps/v1
kind: ReplicaSet
metadata:
 name: my-replicaset
spec:
 replicas: 2
 selector:
   matchLabels:
     app: my-app
 template:
   metadata:
     labels:
       app: my-app
   spec:
     containers:
     - name: my-application
       image: quay.io/devoxxuk/my-application:latest
```

Create new pods

```yaml
apiVersion: apps/v1
kind: ReplicaSet
metadata:
 name: my-replicaset
spec:
 replicas: 2
 selector:
   matchLabels:
     app: my-app
 template:
   metadata:
     labels:
       app: my-app
   spec:
     containers:
     - name: my-application
       image: quay.io/devoxxuk/my-application:latest
```

```yaml
apiVersion: v1
kind: Pod
metadata:
 name: my-replicaset-bf5zv
 labels:
   app: my-app
spec:
 containers:
 - name: my-application
   image: quay.io/devoxxuk…
```

```yaml
apiVersion: v1
kind: Pod
metadata:
 name: my-replicast-1tf5a
 labels:
   app: my-app
spec:
 # ...
```

What happens if the spec changes?

ReplicaSet resource updated

```yaml
apiVersion: apps/v1
kind: ReplicaSet
metadata:
 name: my-replicaset
spec:
 replicas: 2
 selector:
   matchLabels:
     app: my-app
 template:
   metadata:
     labels:
       app: my-app
   spec:
     containers:
     - name: my-application
       image: quay.io/devoxxuk/my-application:latest
```

```yaml
apiVersion: v1
kind: Pod
metadata:
 name: my-replicaset-bf5zv
 labels:
   app: my-app
spec:
 # ...
```

```yaml
apiVersion: v1
kind: Pod
metadata:
 name: my-replicaset-1tf5a
 labels:
   app: my-app
spec:
 # ...
```


```yaml
apiVersion: apps/v1
kind: ReplicaSet
metadata:
 name: my-replicaset
spec:
 replicas: 3
 selector:
   matchLabels:
     app: my-app
 template:
   metadata:
     labels:
       app: my-app
   spec:
     containers:
     - name: my-application
       image: quay.io/devoxxuk/my-application:latest
```

ReplicaSet resource updated

```yaml
apiVersion: v1
kind: Pod
metadata:
 name: my-replicaset-bf5zv
 labels:
   app: my-app
spec:
 # ...
```

```yaml
apiVersion: v1
kind: Pod
metadata:
 name: my-replicaset-1tf5a
 labels:
   app: my-app
spec:
 # ...
```

Search for pods matching resources

```yaml
apiVersion: apps/v1
kind: ReplicaSet
metadata:
 name: my-replicaset
spec:
 replicas: 3
 selector:
   matchLabels:
     app: my-app
 template:
   metadata:
     labels:
       app: my-app
   spec:
     containers:
     - name: my-application
       image: quay.io/devoxxuk/my-application:latest
```

```yaml
apiVersion: v1
kind: Pod
metadata:
 name: my-replicaset-bf5zv
 labels:
   app: my-app
spec:
 # ...
```

```yaml
apiVersion: v1
kind: Pod
metadata:
 name: my-replicaset-1tf5a
 labels:
   app: my-app
spec:
 # ...
```


Create new pod

```yaml
apiVersion: apps/v1
kind: ReplicaSet
metadata:
 name: my-replicaset
spec:
 replicas: 3
 selector:
   matchLabels:
     app: my-app
 template:
   metadata:
     labels:
       app: my-app
   spec:
     containers:
     - name: my-application
       image: quay.io/devoxxuk/my-application:latest
```

```yaml
apiVersion: v1
kind: Pod
metadata:
 name: my-replicaset-bf5zv
 labels:
   app: my-app
spec:
 # ...
```

```yaml
apiVersion: v1
kind: Pod
metadata:
 name: my-replicaset-1tf5a
 labels:
   app: my-app
spec:
 # ...
```

```yaml
apiVersion: v1
kind: Pod
metadata:
 name: my-replicaset-gb65f
 labels:
   app: my-app
spec:
 # ...
```


How to automate operating complex “containerized” applications?

The Operator pattern!

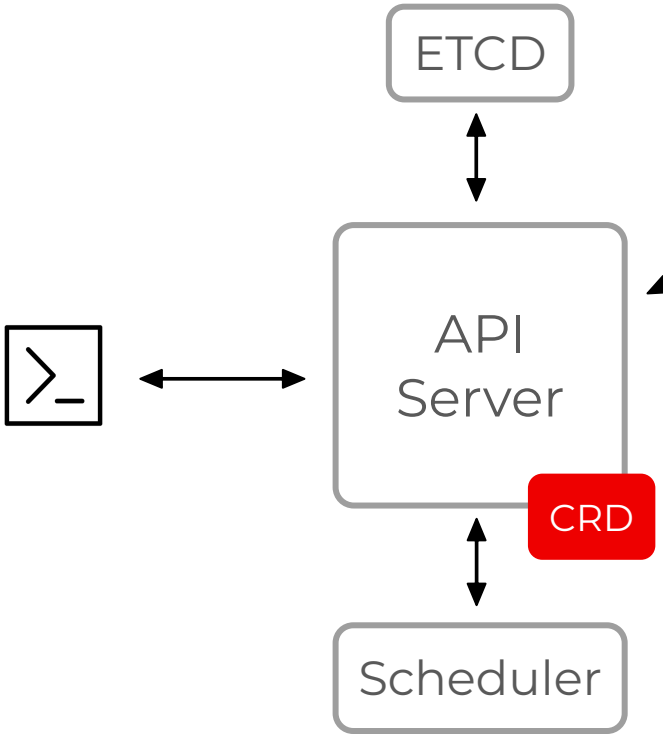


```yaml
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
 #...
spec:
 group: myoperator.io
 names:
   kind: ComplexApp
 #...
 versions:
 - name: v1beta2
   schema:
     openAPIV3Schema:
       #...
```

```yaml
apiVersion: myoperator.io/v1
kind: ComplexApp
metadata:
 name: my-complex-app
spec:
 # ...
status:
 # ...
```

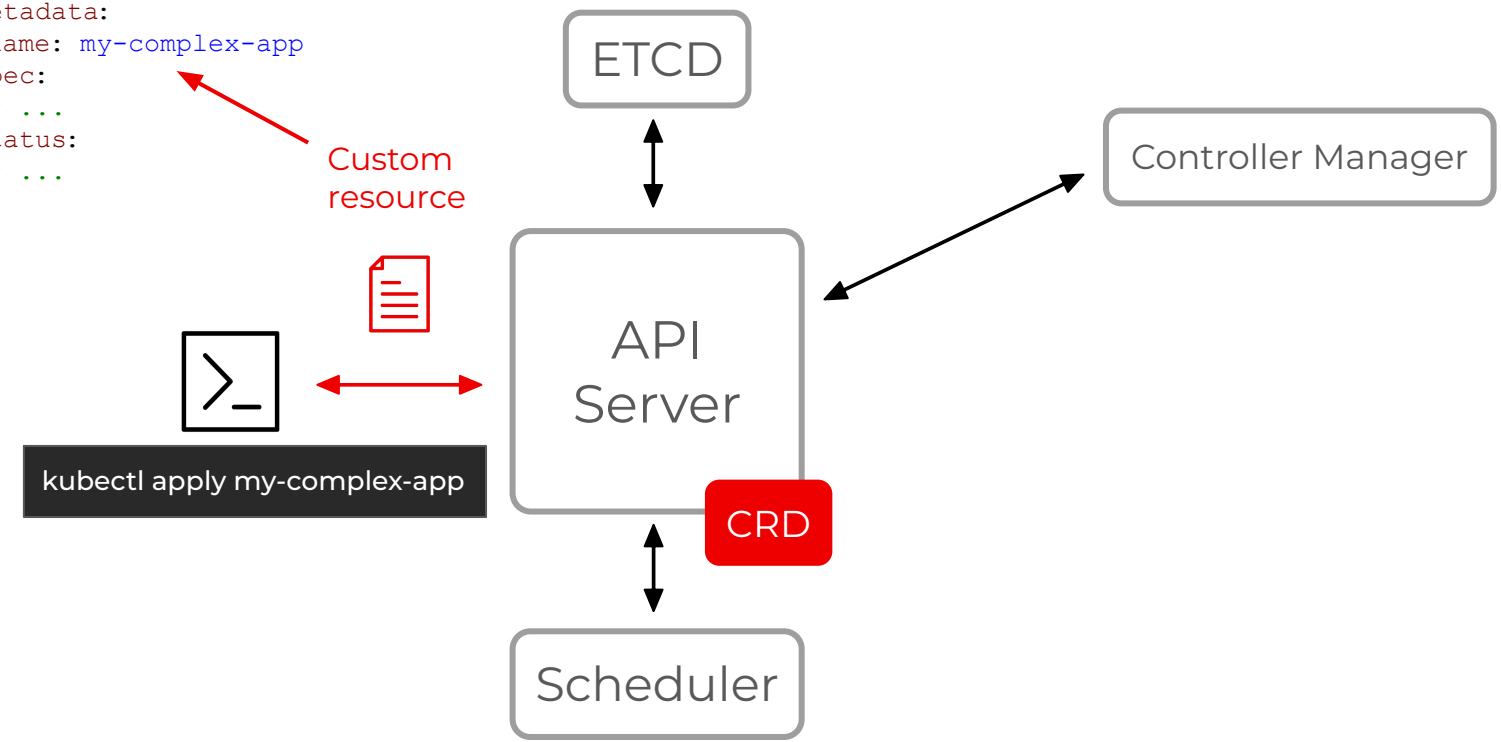

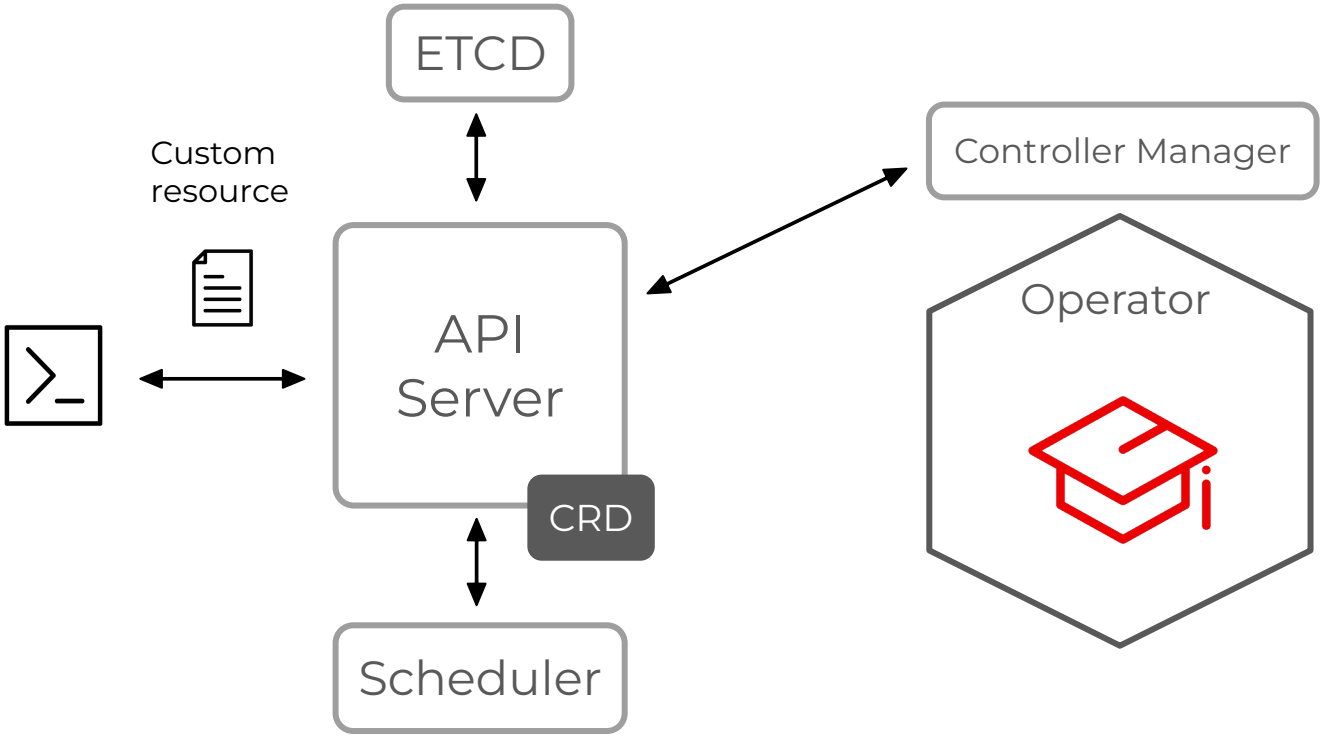

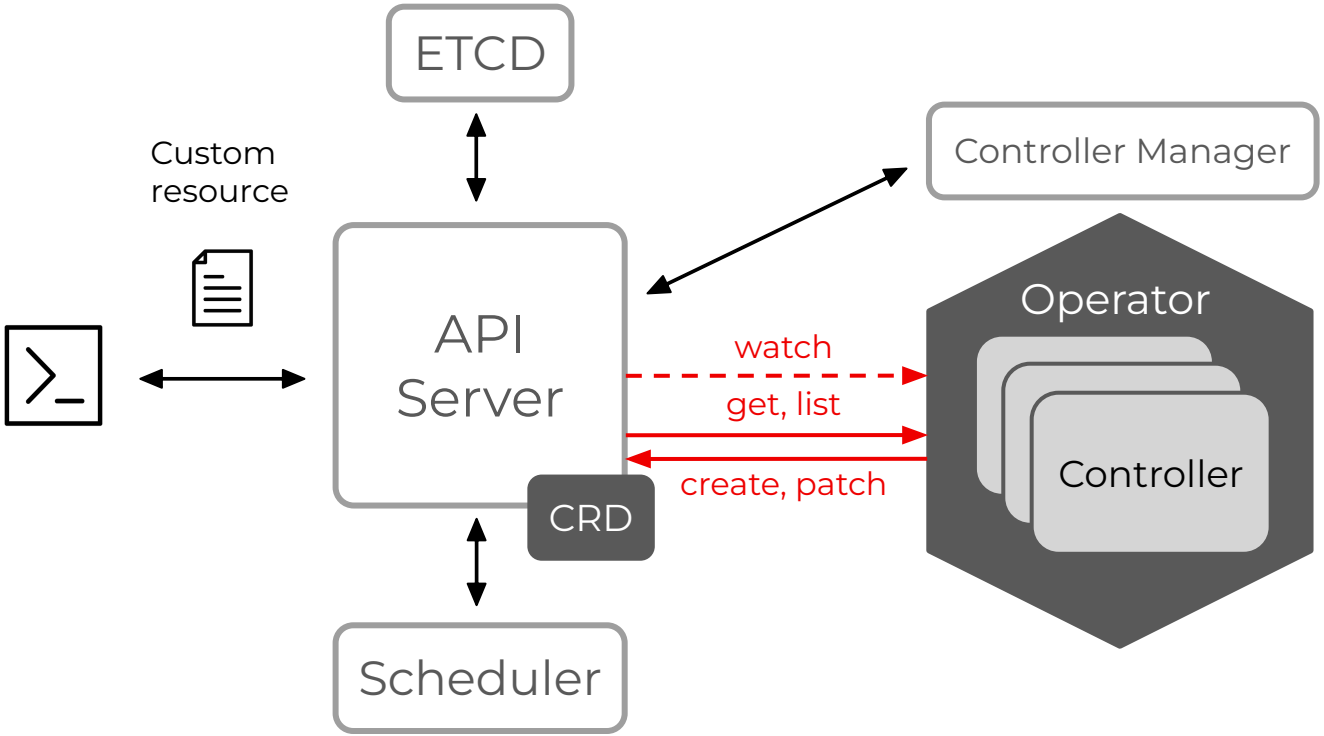

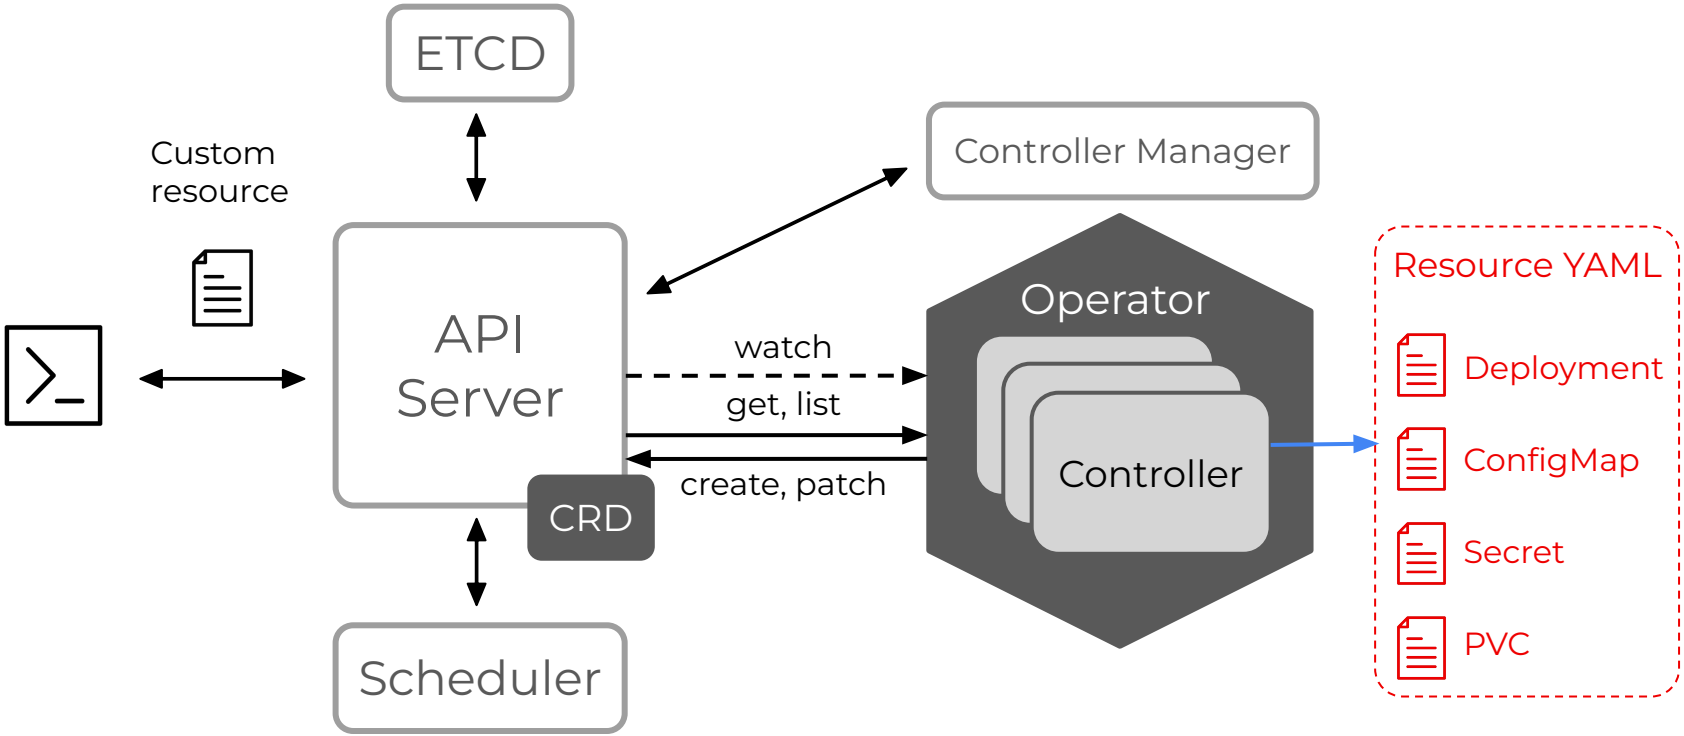

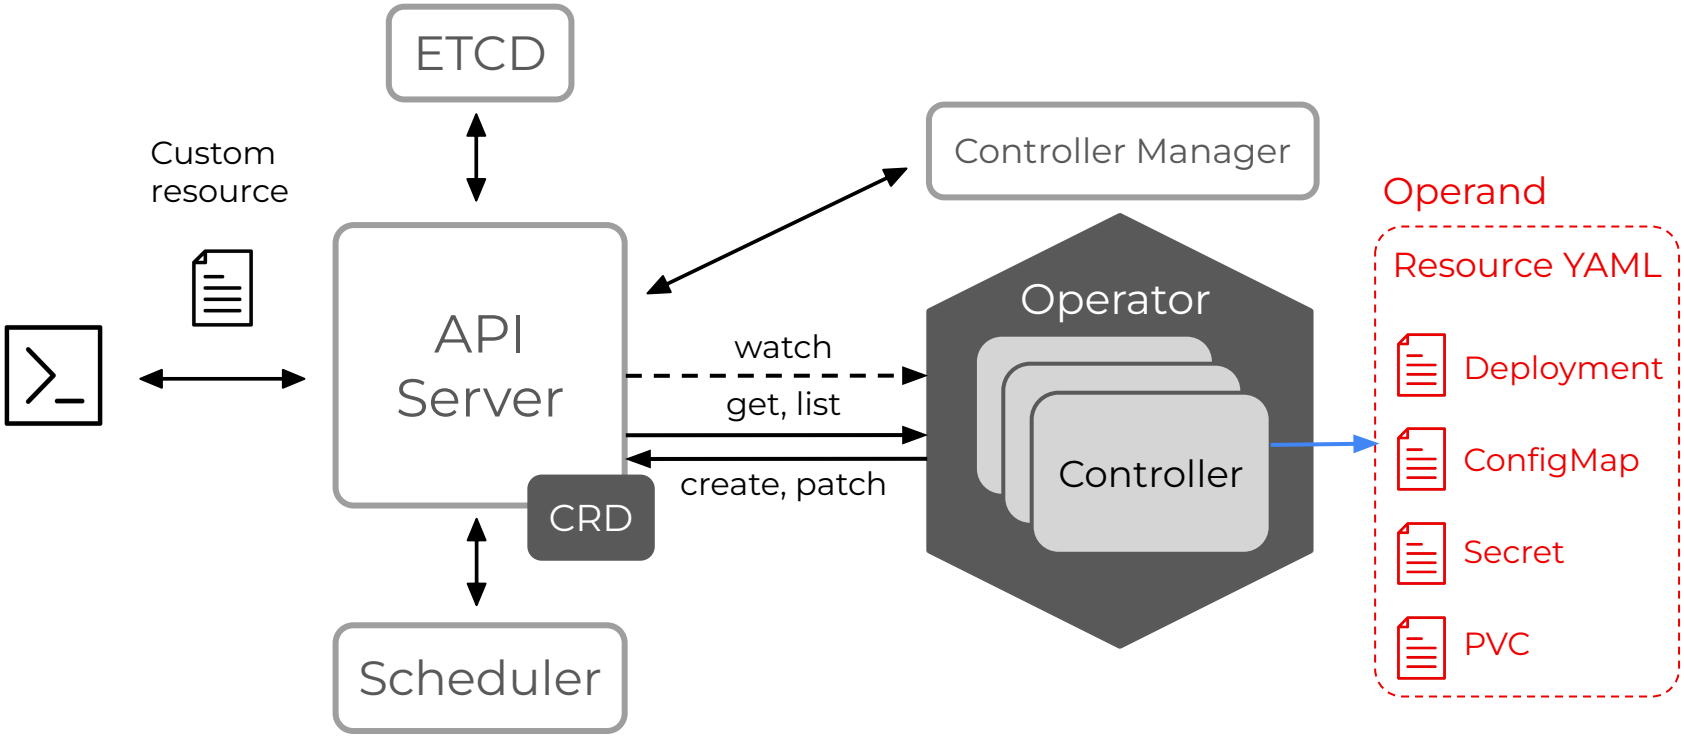

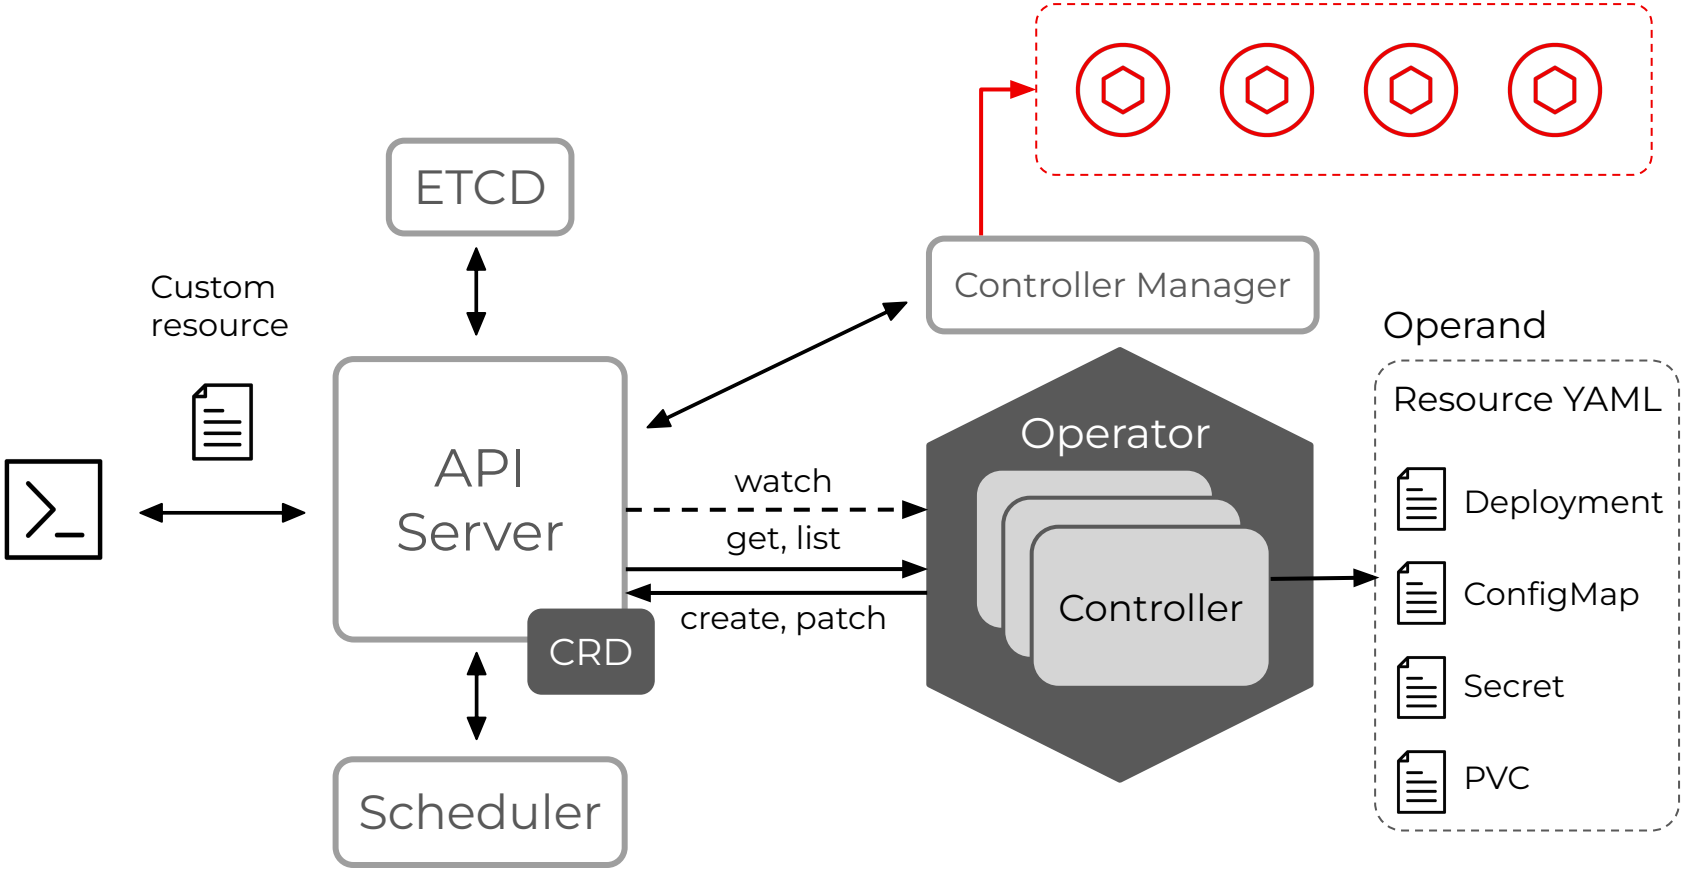

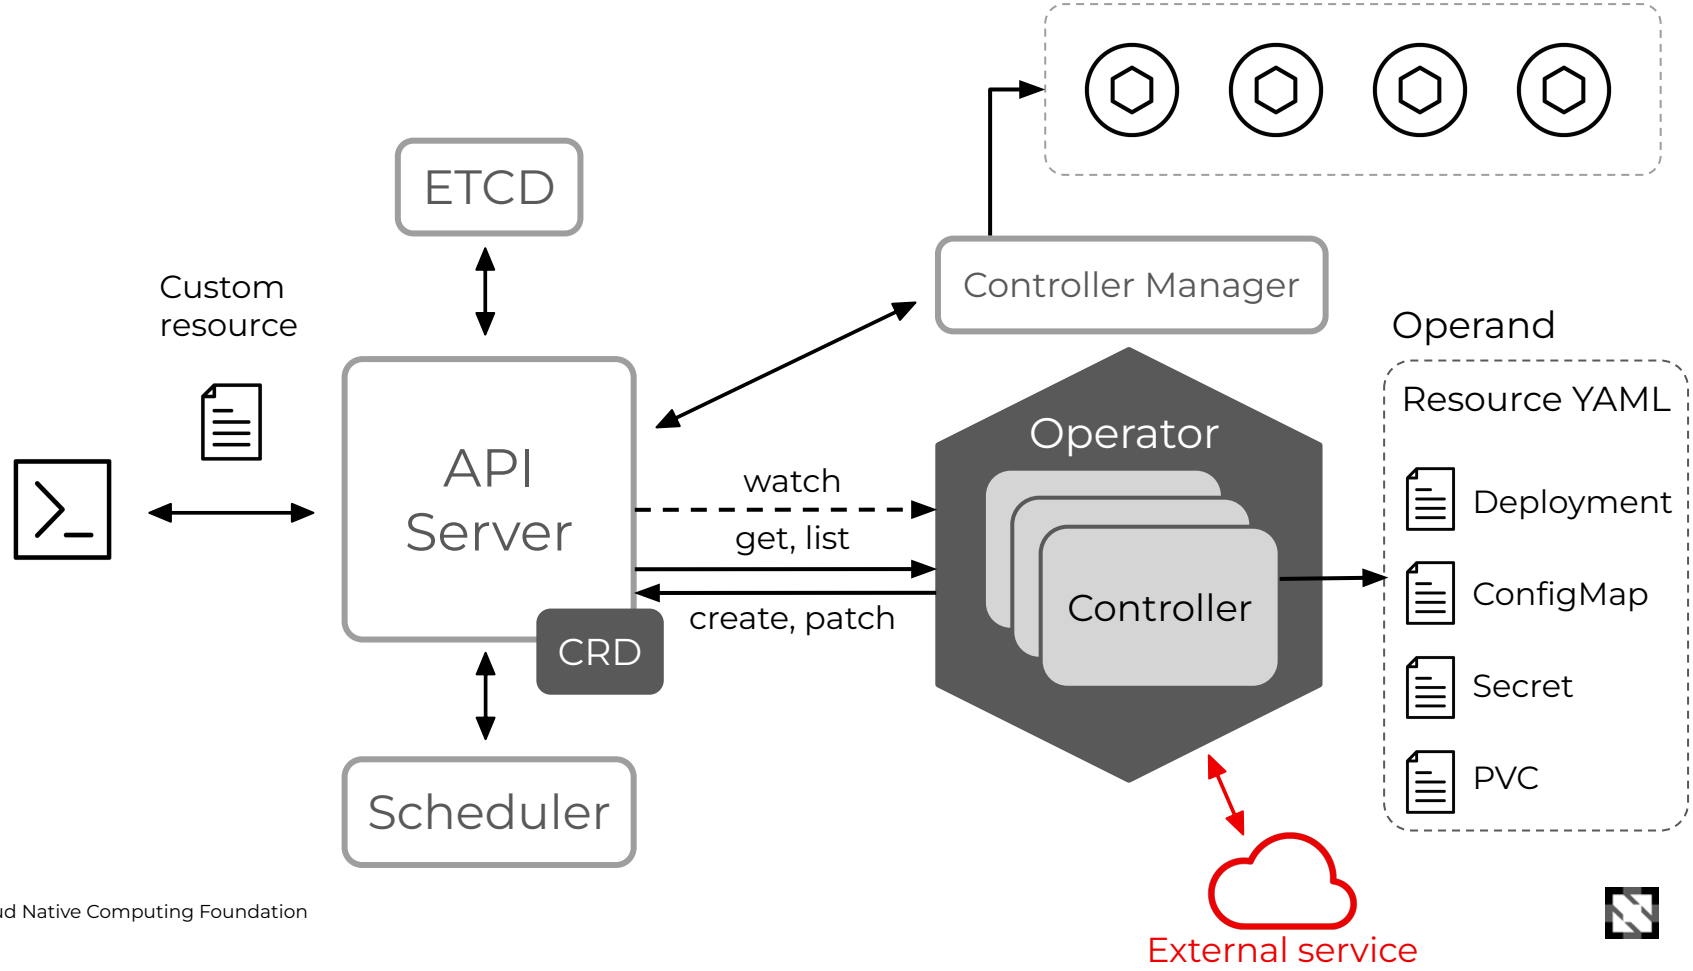

What if I want to deploy Apache Kafka on Kubernetes?

```yaml
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
 name: kafkas.kafka.strimzi.io
spec:
 group: kafka.strimzi.io
 names:
   kind: Kafka
   listKind: KafkaList
 #...
 versions:
 - name: v1beta2
   schema:
     openAPIV3Schema:
       type: object
       properties:
         spec:
           # spec definition for the custom resource
           kafka:
             #...
         status:
           # status definition reported back
           # in the custom resource
```

Strimzi

strimzi.io

Open Source project (Apache License 2.0)

Focuses on Apache Kafka on Kubernetes

CNCF Incubating Project

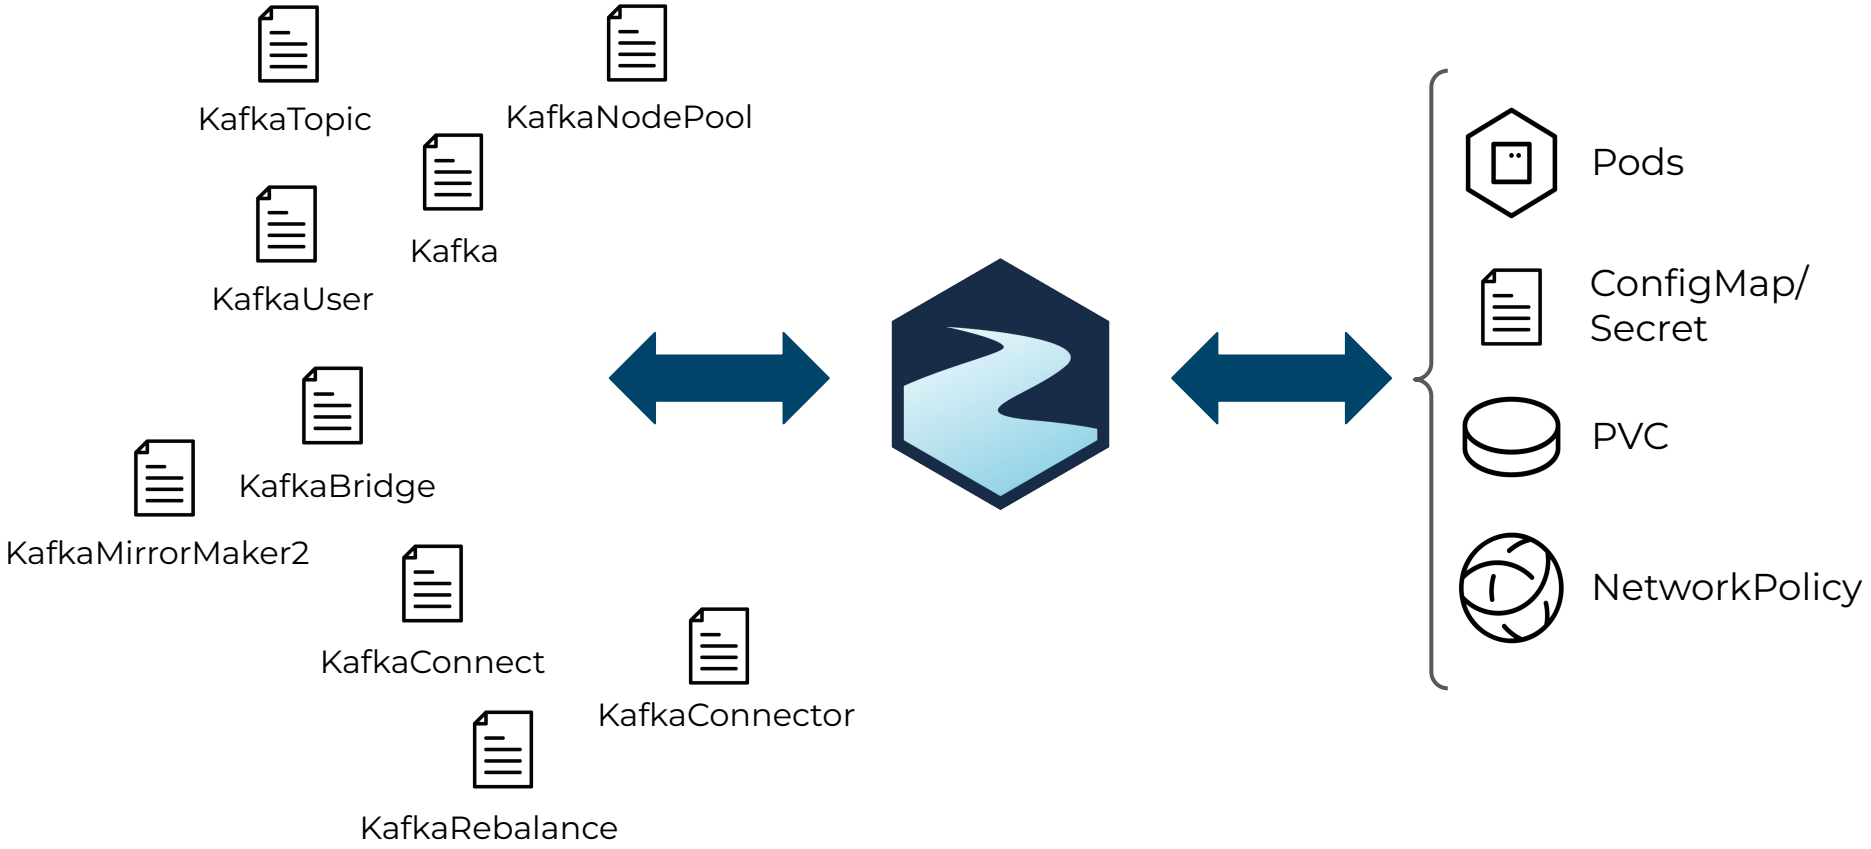

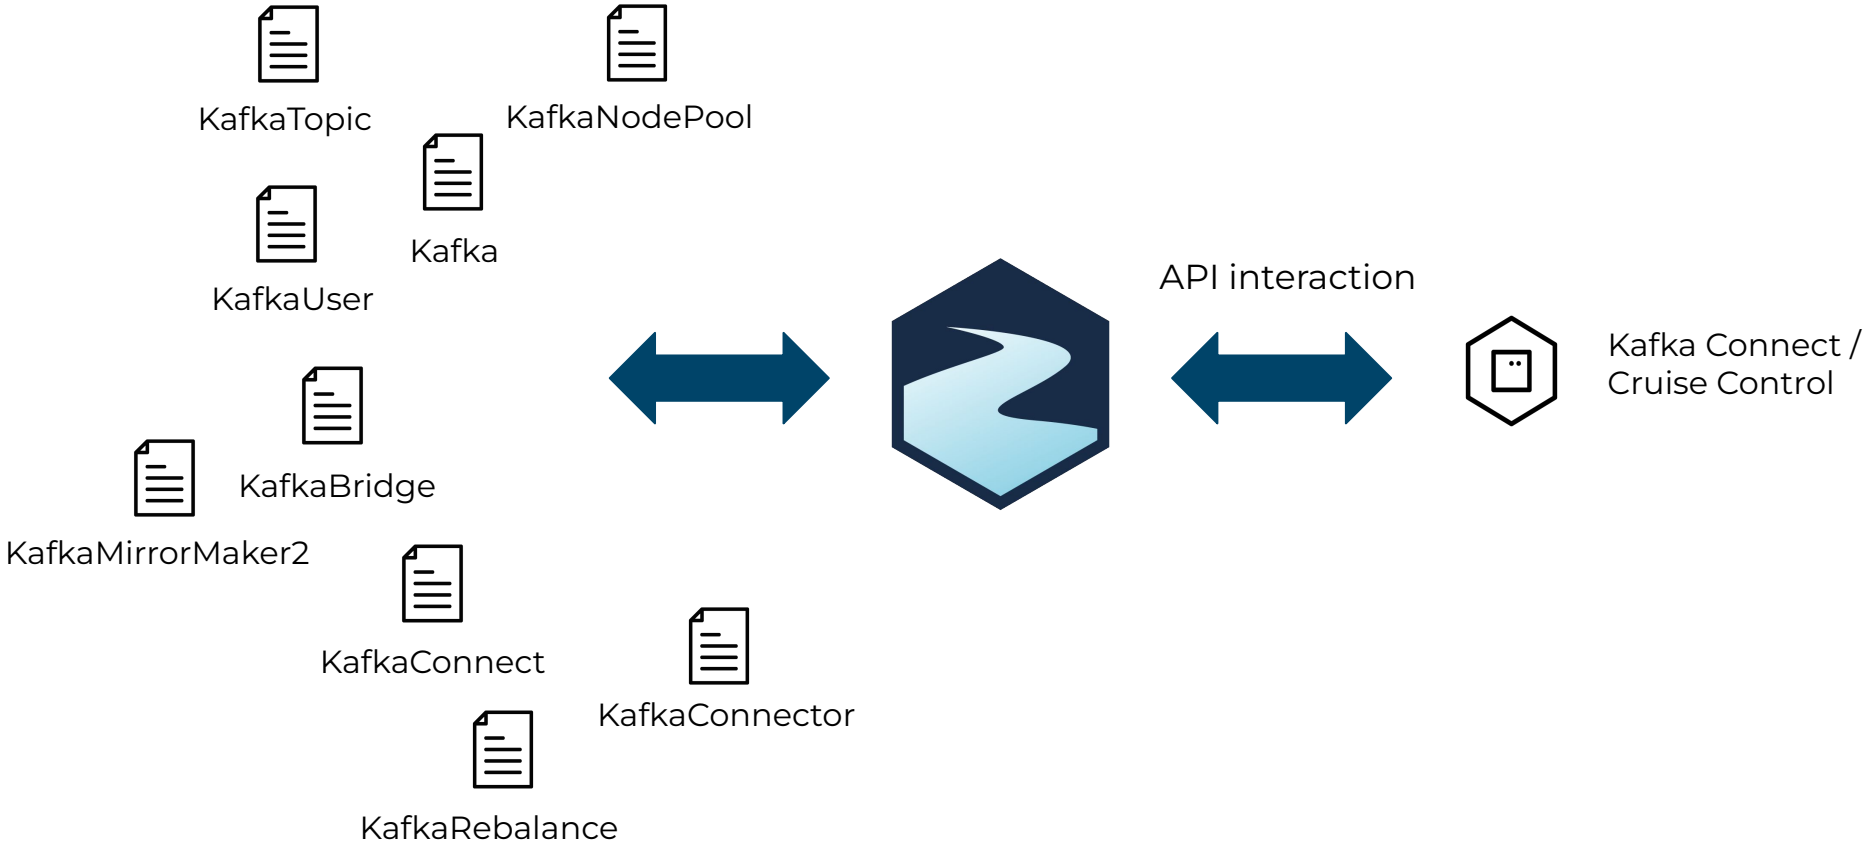

Demo time!

Helm

- Relies on Kubernetes built-in resources
- Many YAMLs with customization via templating
- Ideal for day-1 operation (deploying)

Operator

- Extends the Kubernetes API with CRDs
- One (or a few) “custom resource” YAMLs
- Useful for day-1 and day-2 operations (upgrading, scaling)
- Deployable via Helm charts!

Where to start?

- Operator SDK ... for writing operators in Go
- Java Operator SKD ... for writing operators in Java
- OperatorHub.io ... provides an operators catalog

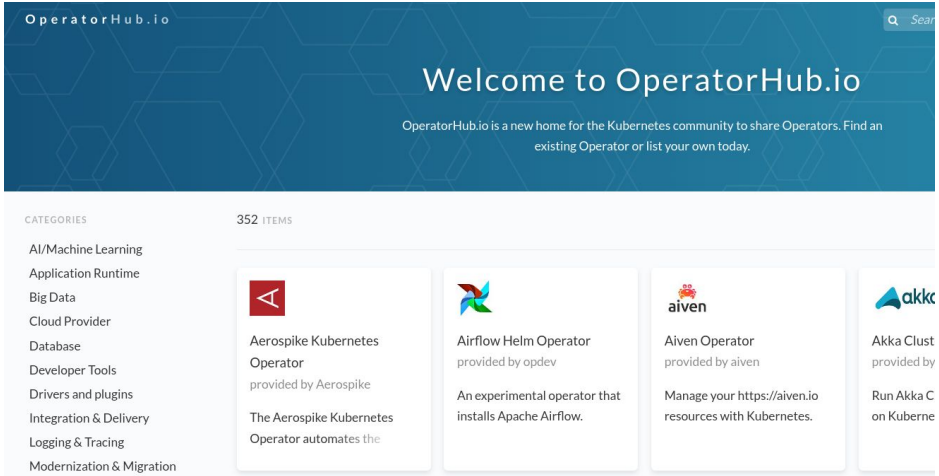

Thank you!